

PART III: THE FOUNDATIONS OF DEEP LEARNING

THE BASIC BUILDING BLOCKS OF CNNs

Discrete Convolution

1D:
[Spectra]

$$f(x) * g(x) = \sum_{k=-\infty}^{k=+\infty} f(k).g(k - x)$$

2D:
[Images]

$$f(x, y) * g(x, y) = \sum_{k=-\infty}^{k=+\infty} \sum_{l=-\infty}^{l=+\infty} f(k, l).g(x - k, y - l)$$

DISCRETE CONVOLUTION

1D:
[Spectra]

$$f(x) * g(x) = \sum_{k=-\infty}^{k=+\infty} f(k).g(k - x)$$

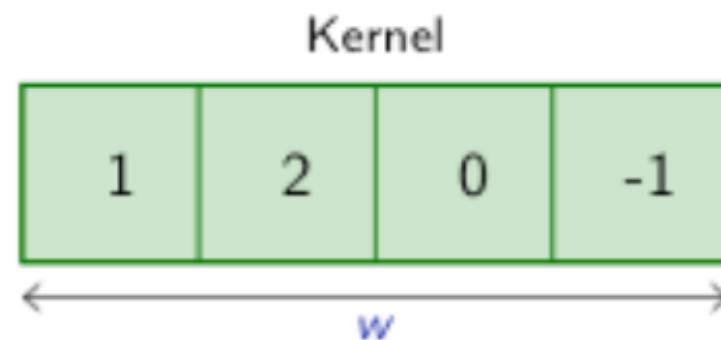
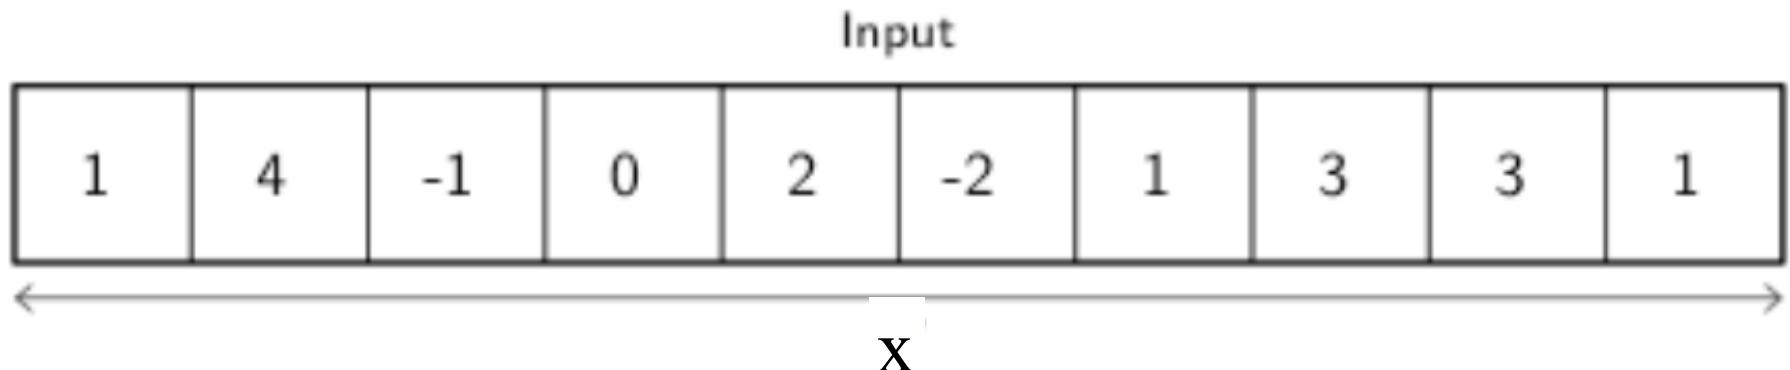
2D:
[Images]

$$f(x, y) * g(x, y) = \sum_{k=-\infty}^{k=+\infty} \sum_{l=-\infty}^{l=+\infty} f(k, l).g(x - k, y - l)$$

CONVOLUTION KERNEL

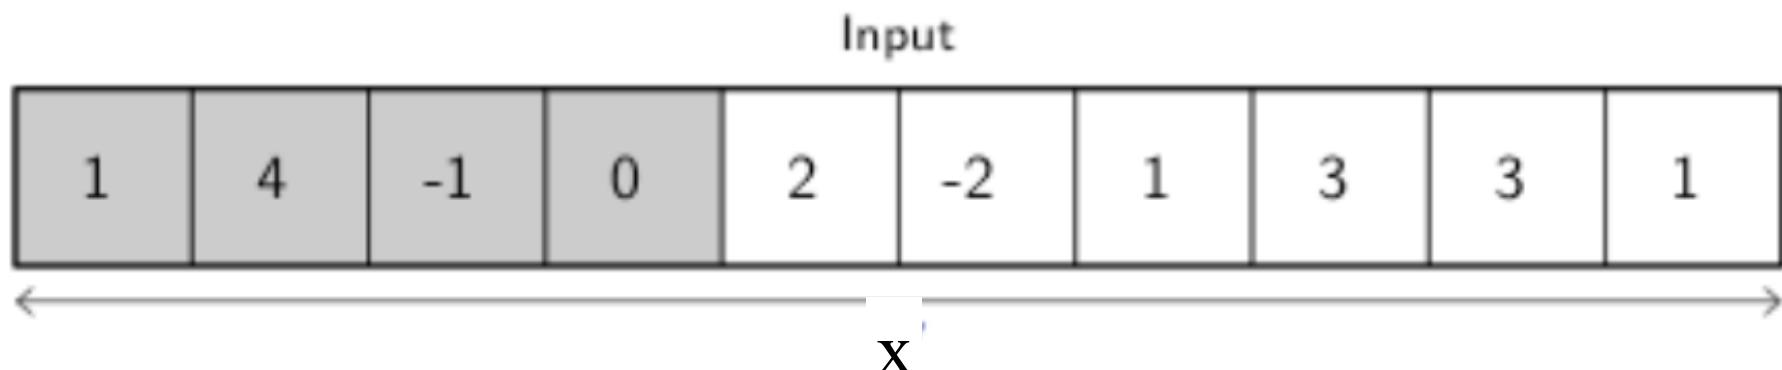
INPUT DATA

1-D CONVOLUTION



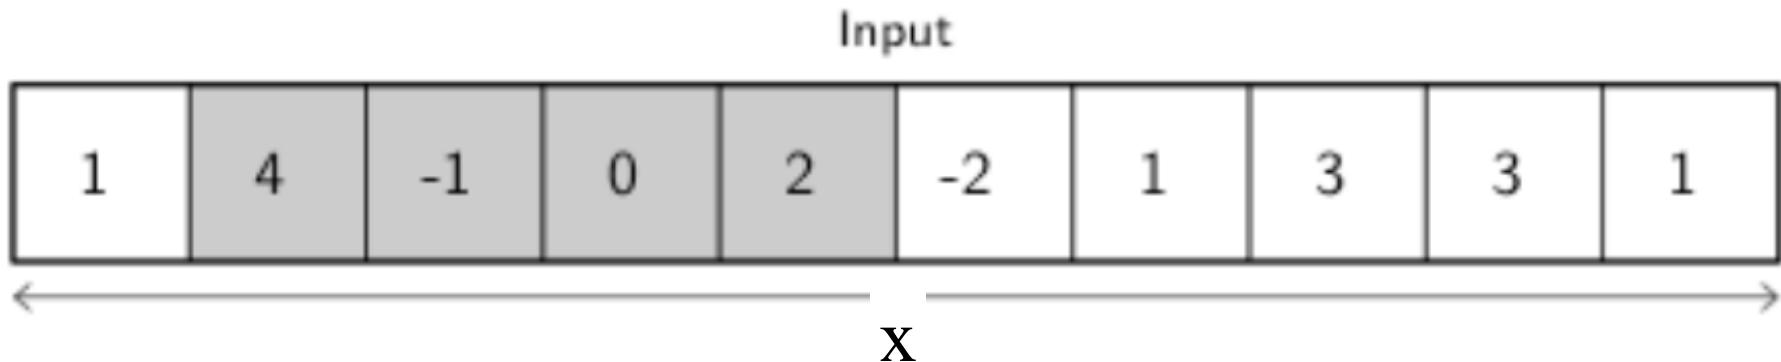
credit

1-D CONVOLUTION



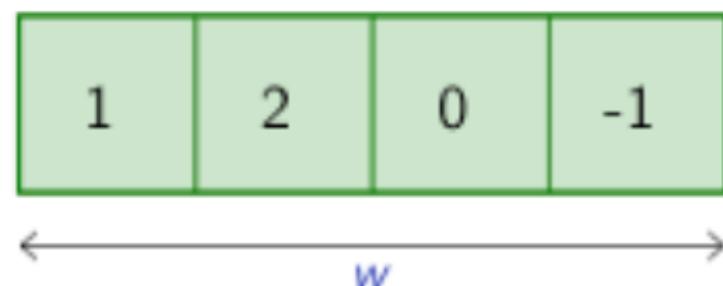
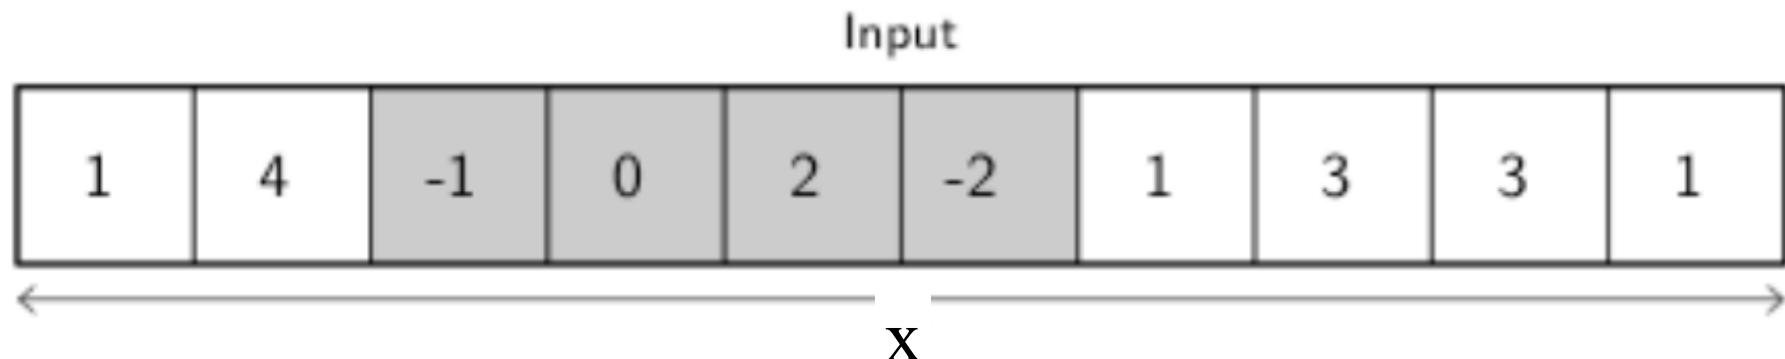
credit

1-D CONVOLUTION



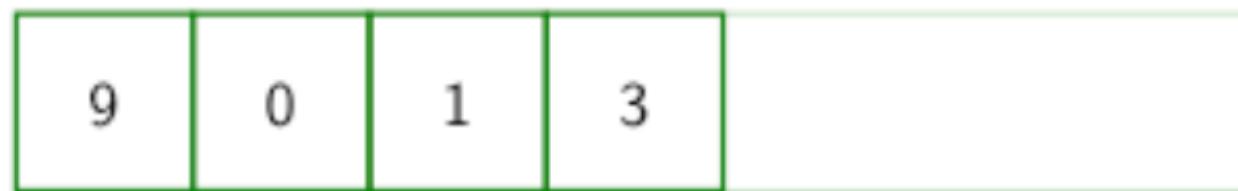
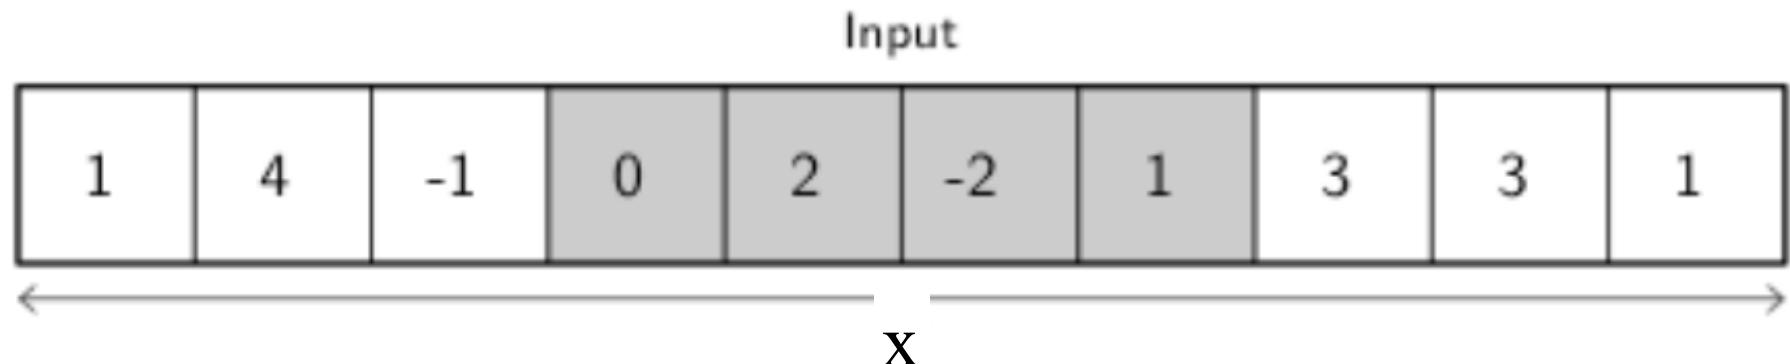
credit

1-D CONVOLUTION



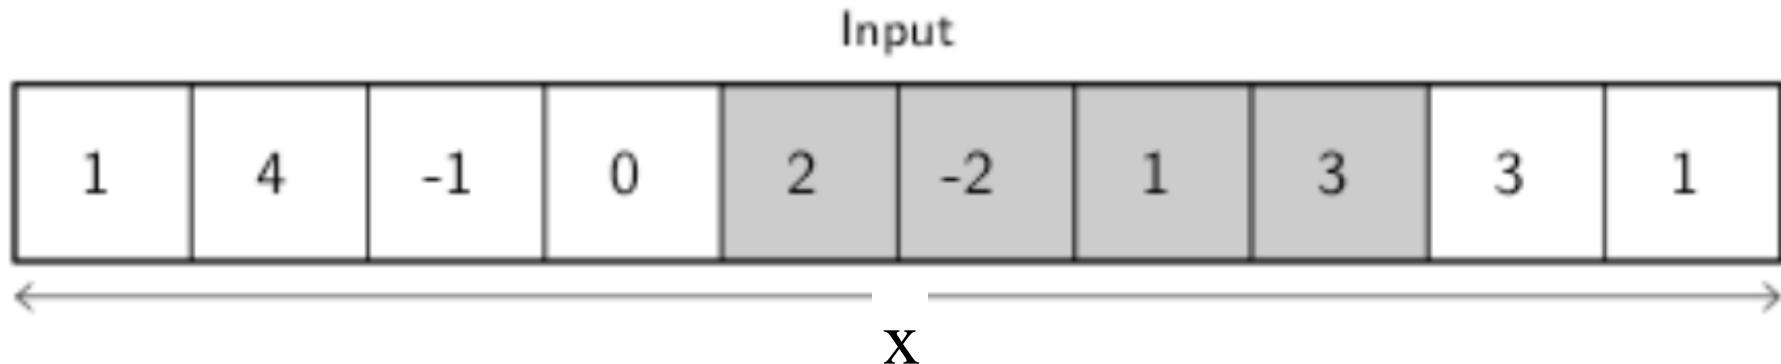
credit

1-D CONVOLUTION



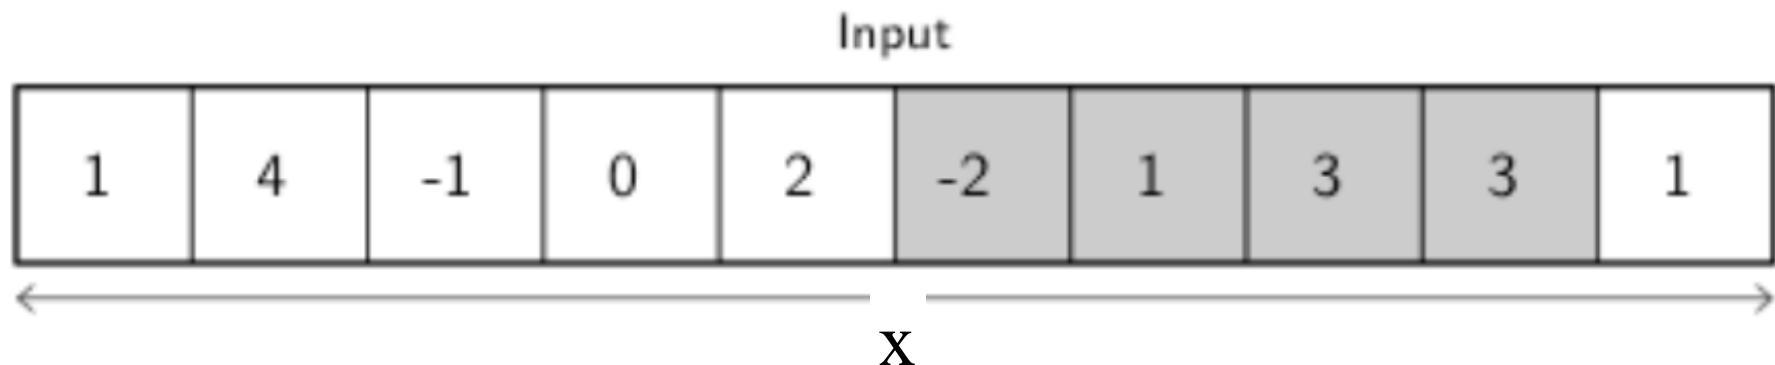
credit

1-D CONVOLUTION



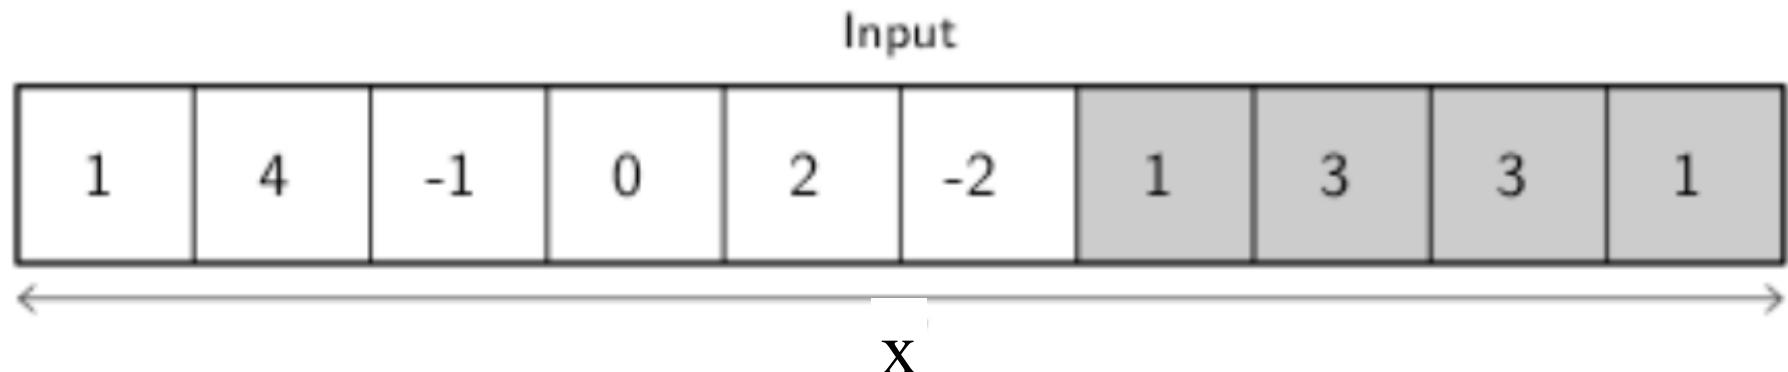
credit

1-D CONVOLUTION



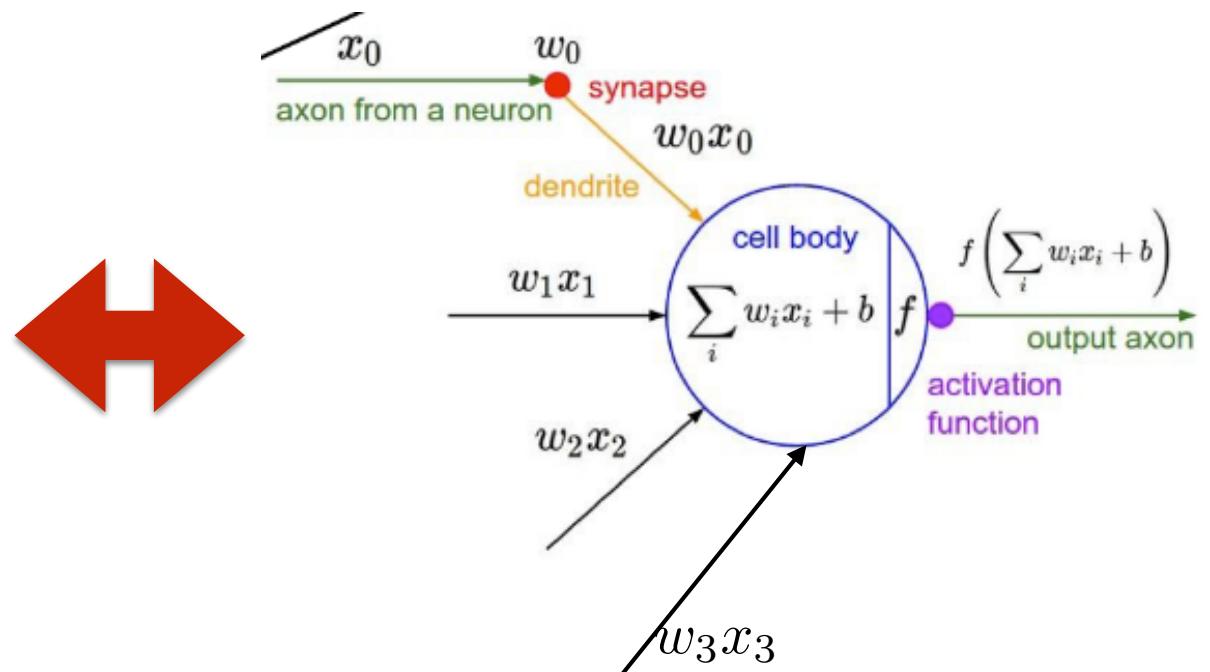
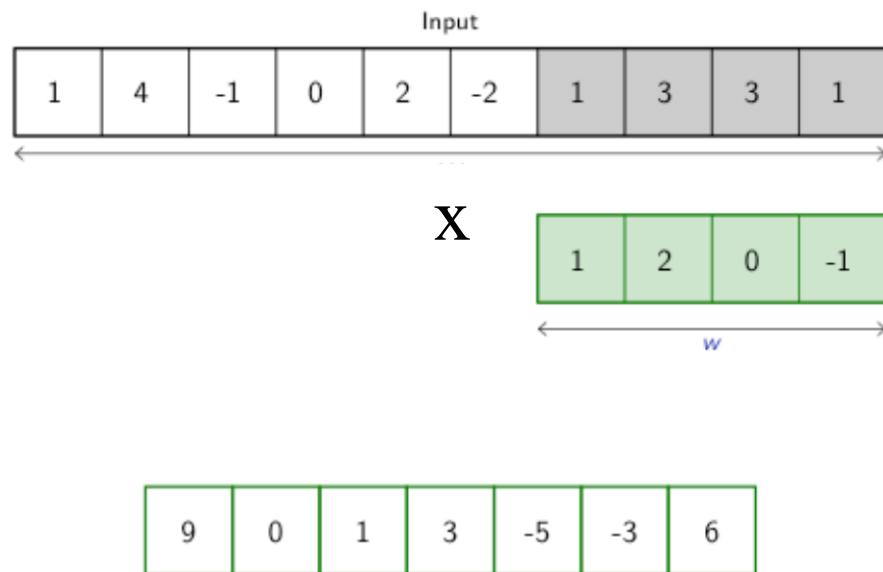
credit

1-D CONVOLUTION

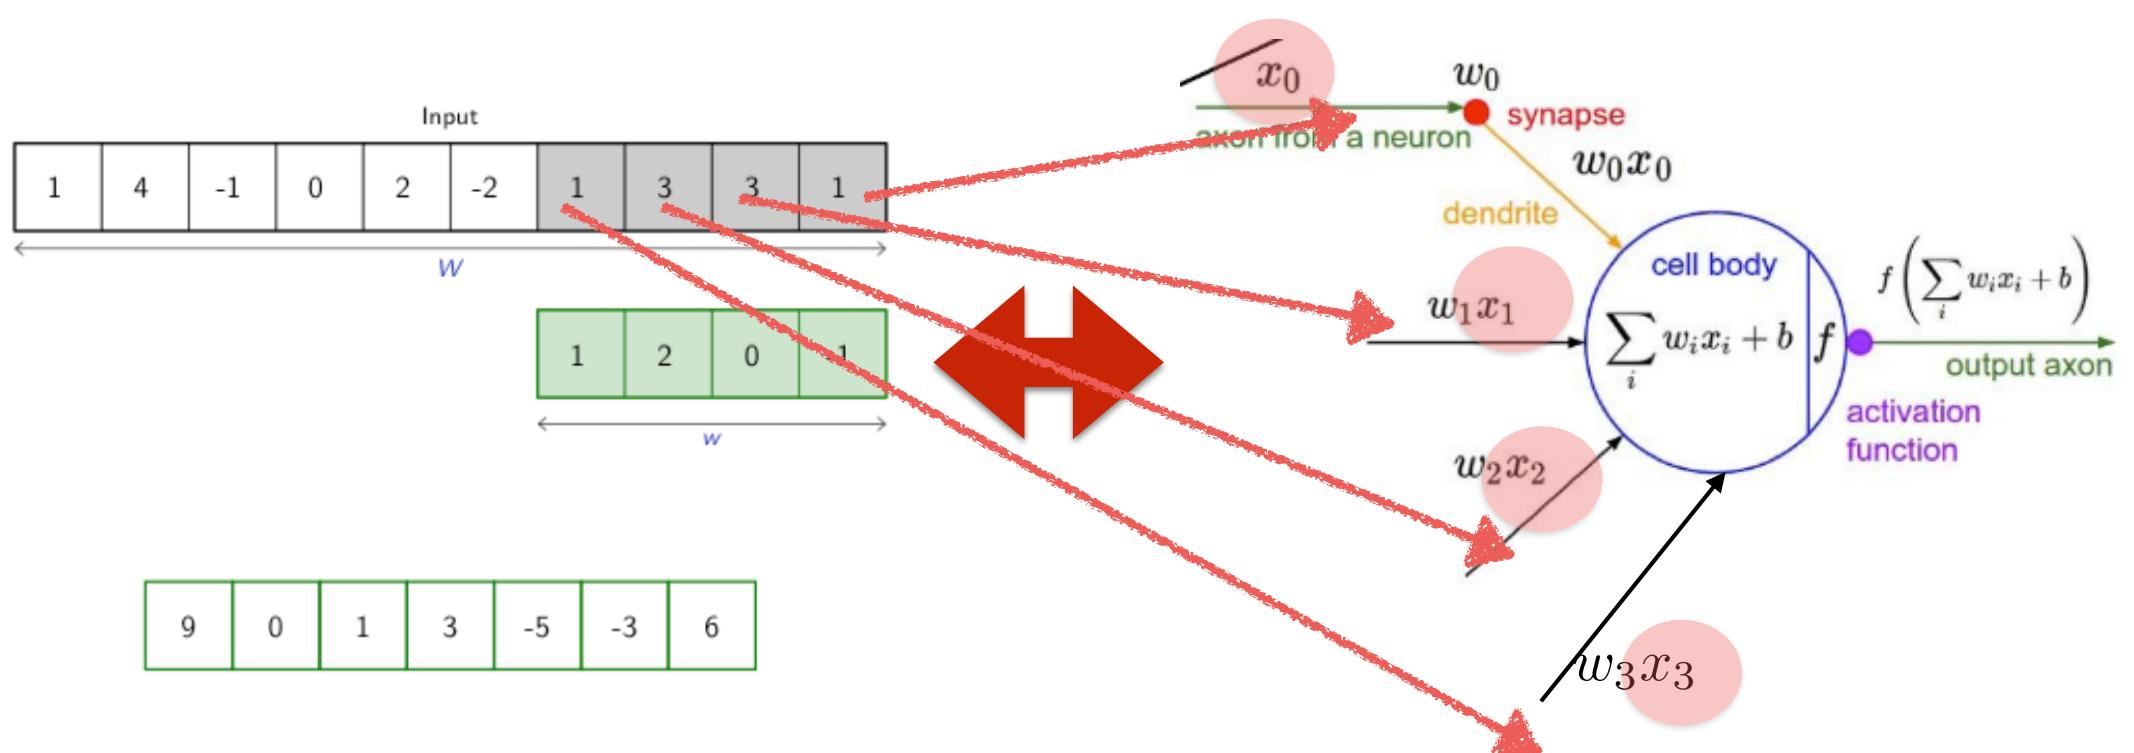


credit

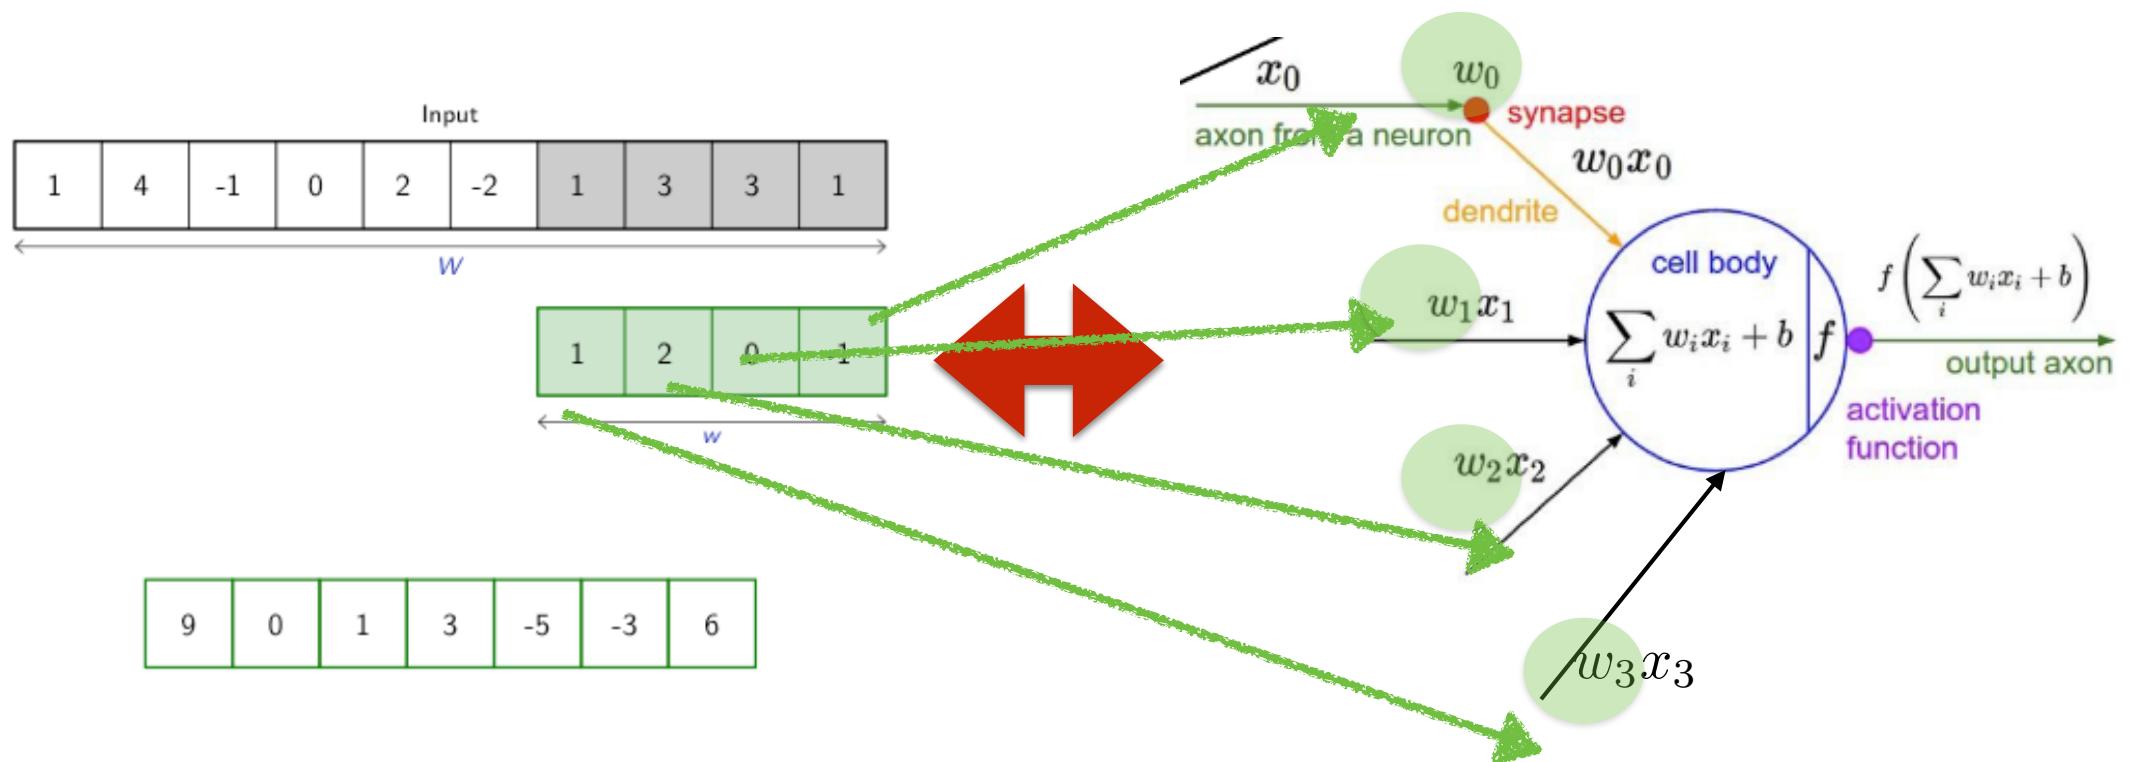
THE CONVOLUTION BUILDING BLOCK OPERATION IS EQUIVALENT TO A NEURON WITH AS MANY INPUTS AS KERNEL ELEMENTS AND WEIGHTS EQUAL TO THE KERNEL



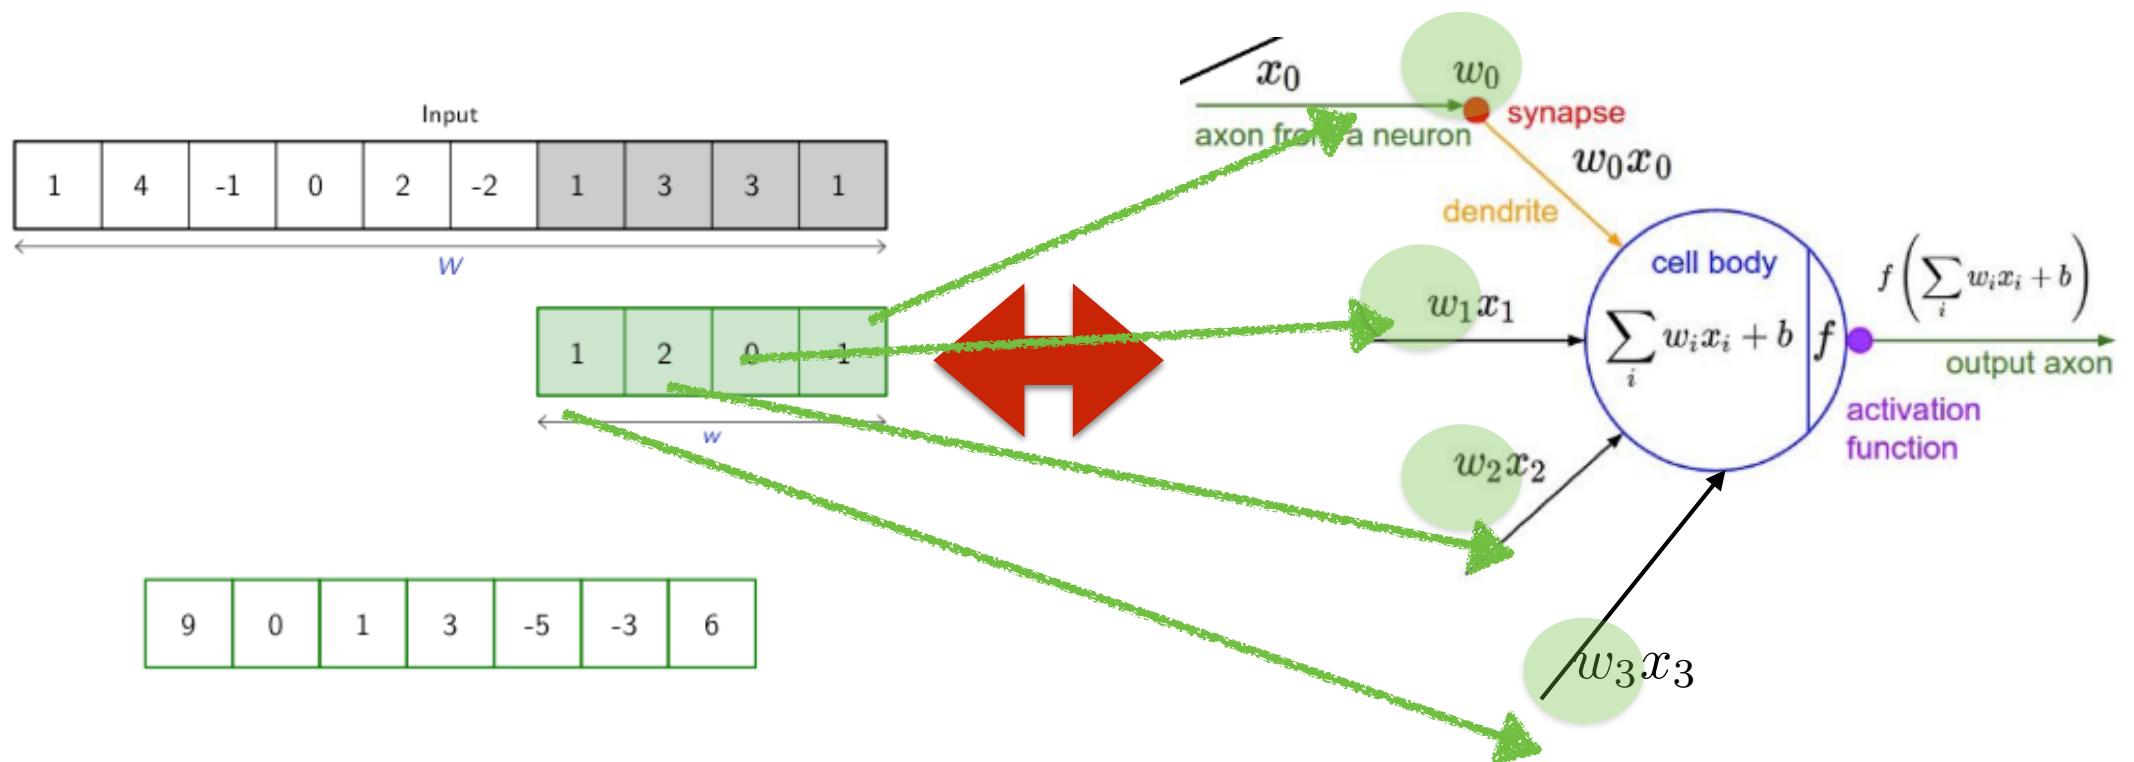
THE CONVOLUTION BUILDING BLOCK OPERATION IS EQUIVALENT TO A NEURON WITH AS MANY INPUTS AS KERNEL ELEMENTS AND WEIGHTS EQUAL TO THE KERNEL



THE CONVOLUTION BUILDING BLOCK OPERATION IS EQUIVALENT TO A NEURON WITH AS MANY INPUTS AS KERNEL ELEMENTS AND WEIGHTS EQUAL TO THE KERNEL



THE CONVOLUTION BUILDING BLOCK OPERATION IS EQUIVALENT TO A NEURON WITH AS MANY INPUTS AS KERNEL ELEMENTS AND WEIGHTS EQUAL TO THE KERNEL



WITH THE ADVANTAGE THAT THE SAME WEIGHTS ARE APPLIED TO ALL THE SIGNAL: TRANSLATION INVARIANCE

2-D CONVOLUTION

SAME IDEA, BUT THE KERNEL IS NOW 2D

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3

KERNEL

INPUT (IMAGE)

OUTPUT

Credit: animations from https://github.com/vdumoulin/conv_arithmetic

2-D CONVOLUTION

SAME IDEA, BUT THE KERNEL IS NOW 2D

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3

IN THE EXAMPLE: EACH 3x3 REGION GENERATES AN OUTPUT

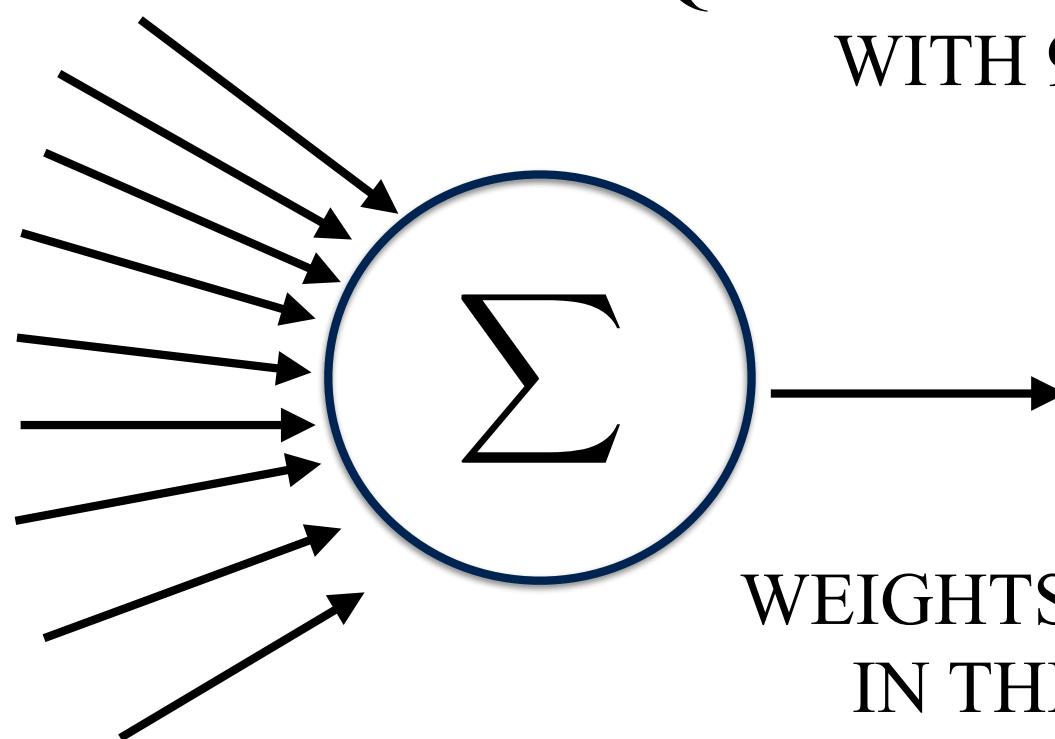
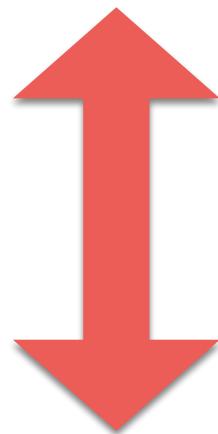
$$Size_{output} = Size_{input} - Size_{kernel} + 1$$

Credit: animations from https://github.com/vdumoulin/conv_arithmetic

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3



EQUIVALENT TO A NEURON
WITH 9 INPUTS

WEIGHTS ARE CODED
IN THE KERNEL

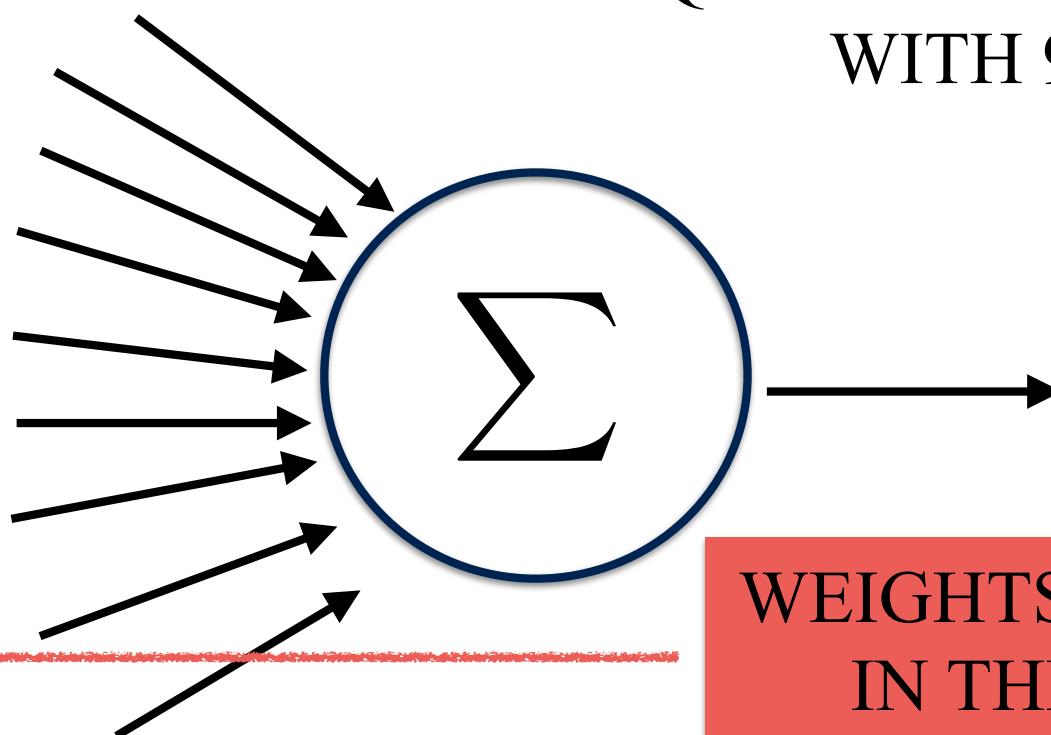
1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3



EQUIVALENT TO A NEURON
WITH 9 INPUTS



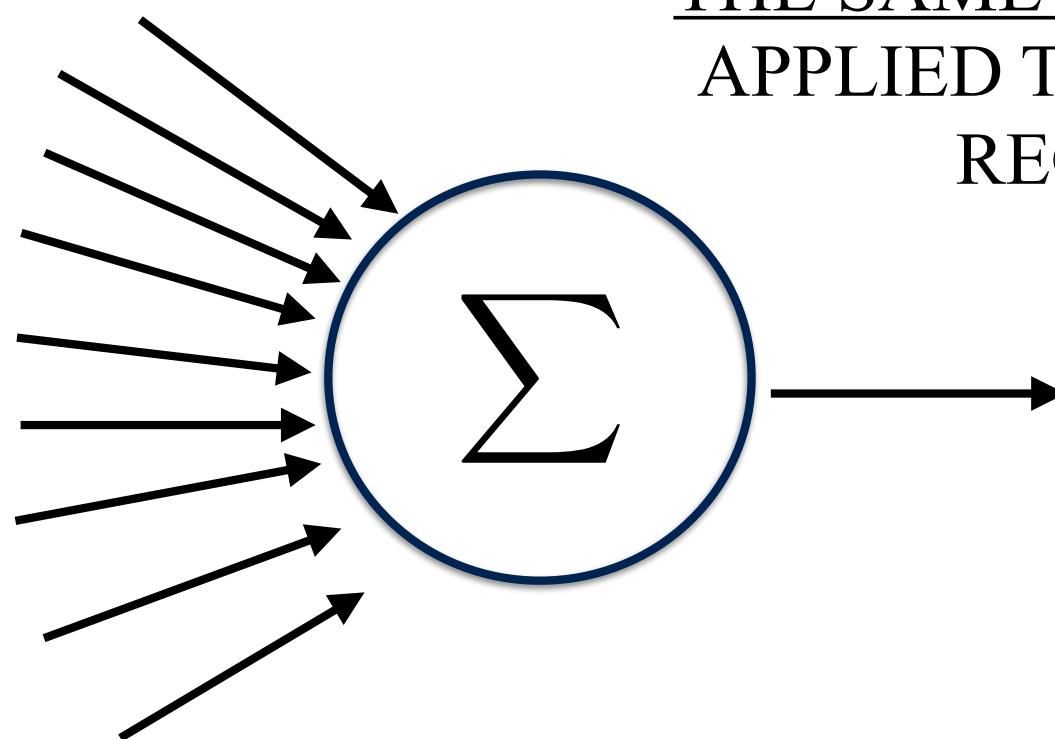
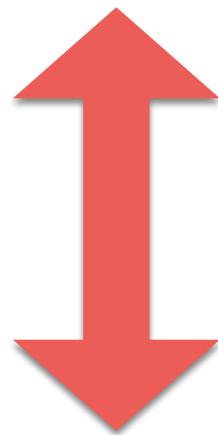
THIS IS WHAT
THE
NETWORK
LEARNS!

WEIGHTS ARE CODED
IN THE KERNEL

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3

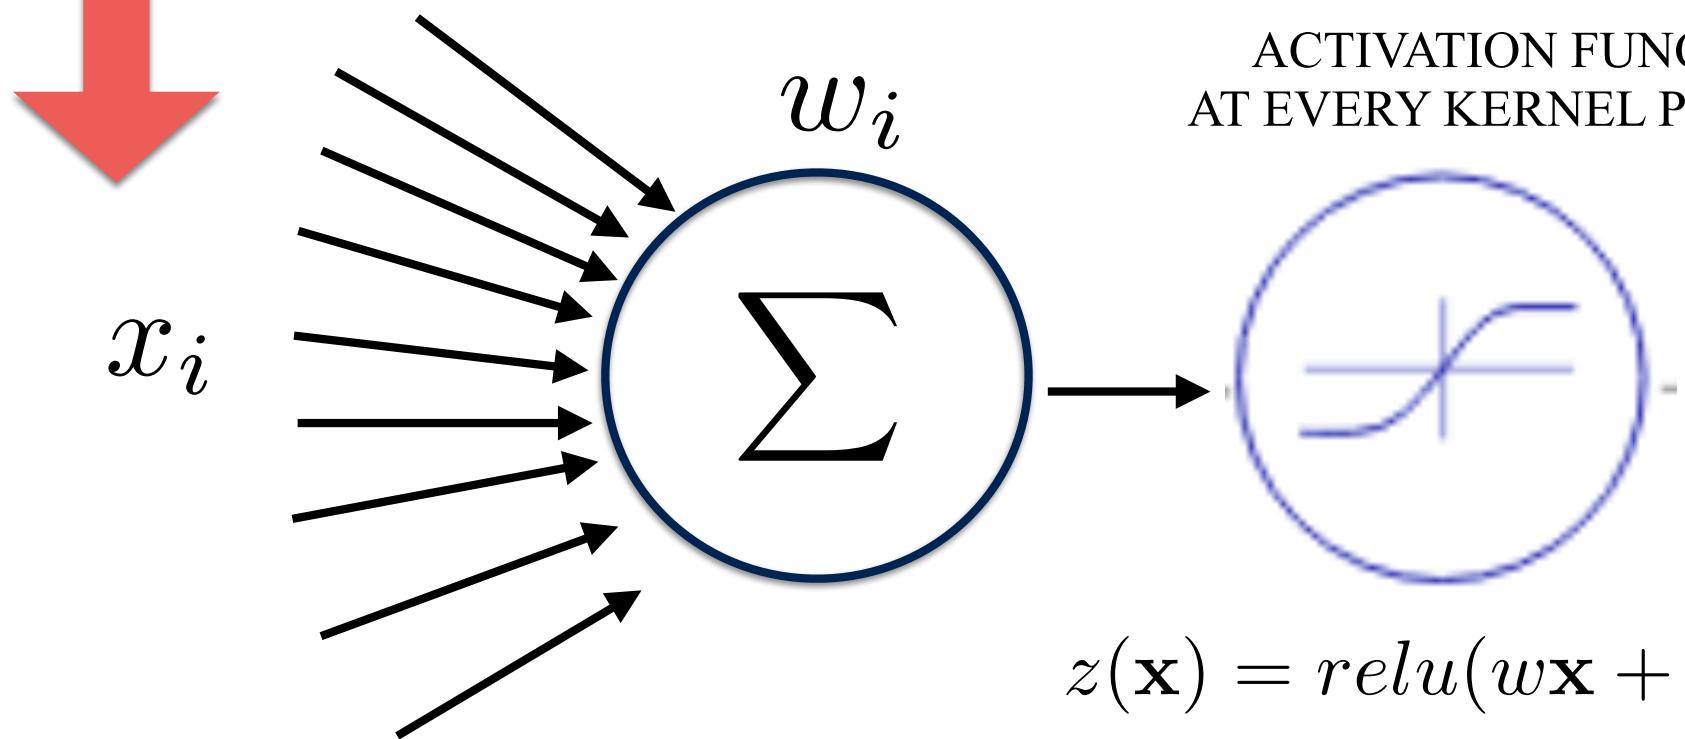


THE KEY IS AGAIN THAT
THE SAME WEIGHTS ARE
APPLIED TO ALL IMAGE
REGIONS

x_i		
-------	--	--

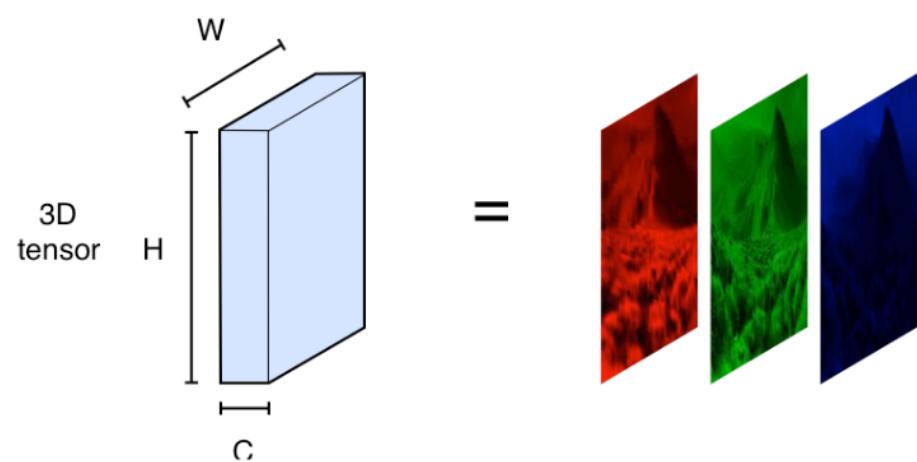
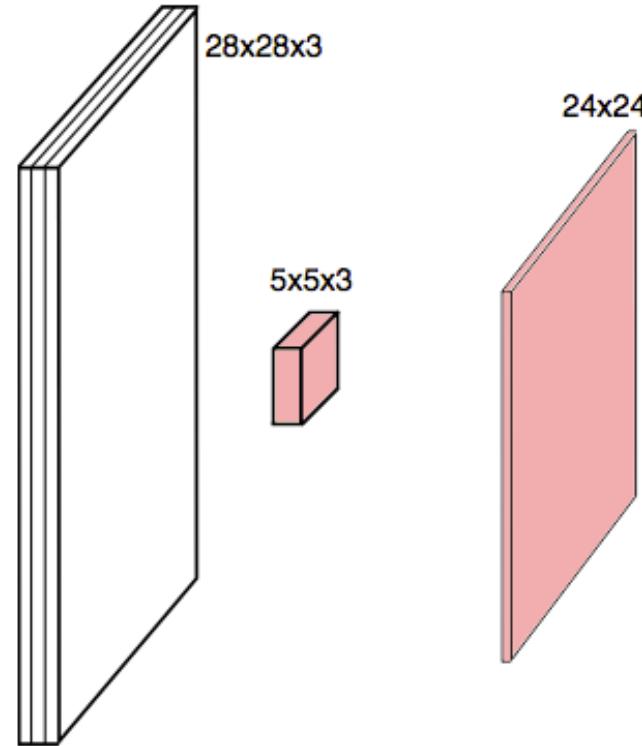
w_i		
-------	--	--

[weights]



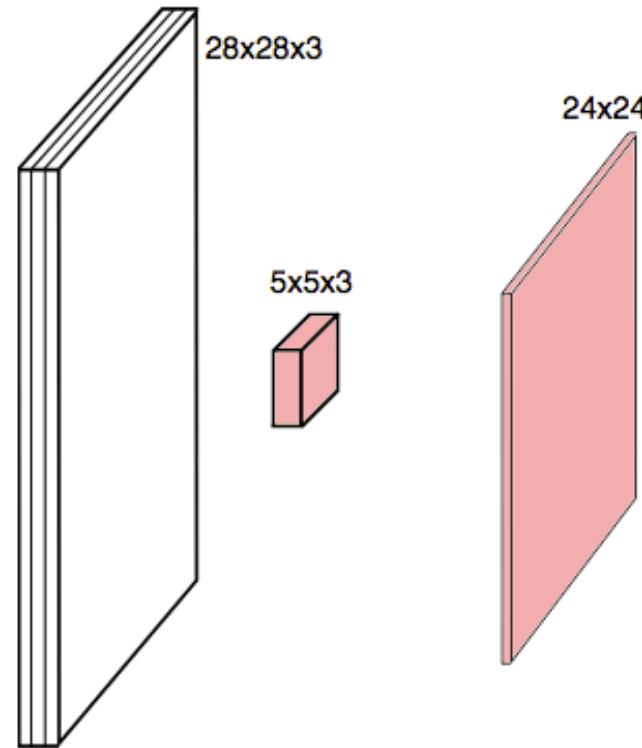
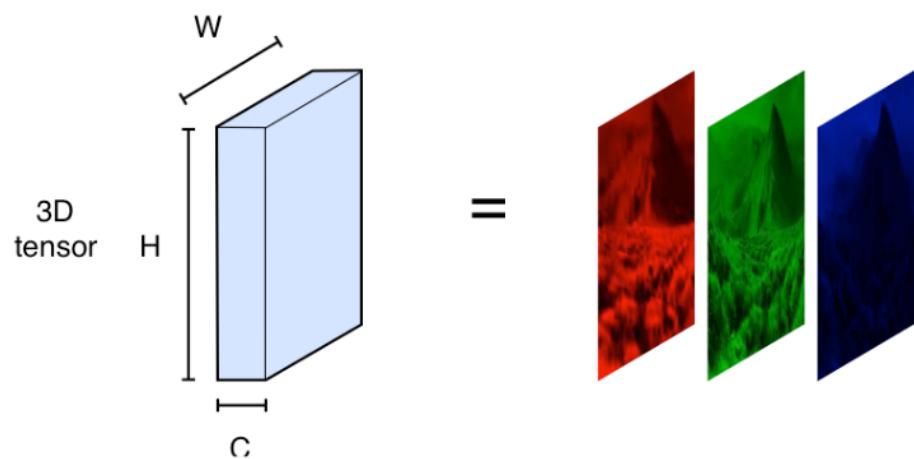
CONVOLUTIONS CAN ALSO BE COMPUTED ACROSS CHANNELS (OR COLORS)

A COLOR IMAGE IS A
TENSOR
OF SIZE height x width x
channels



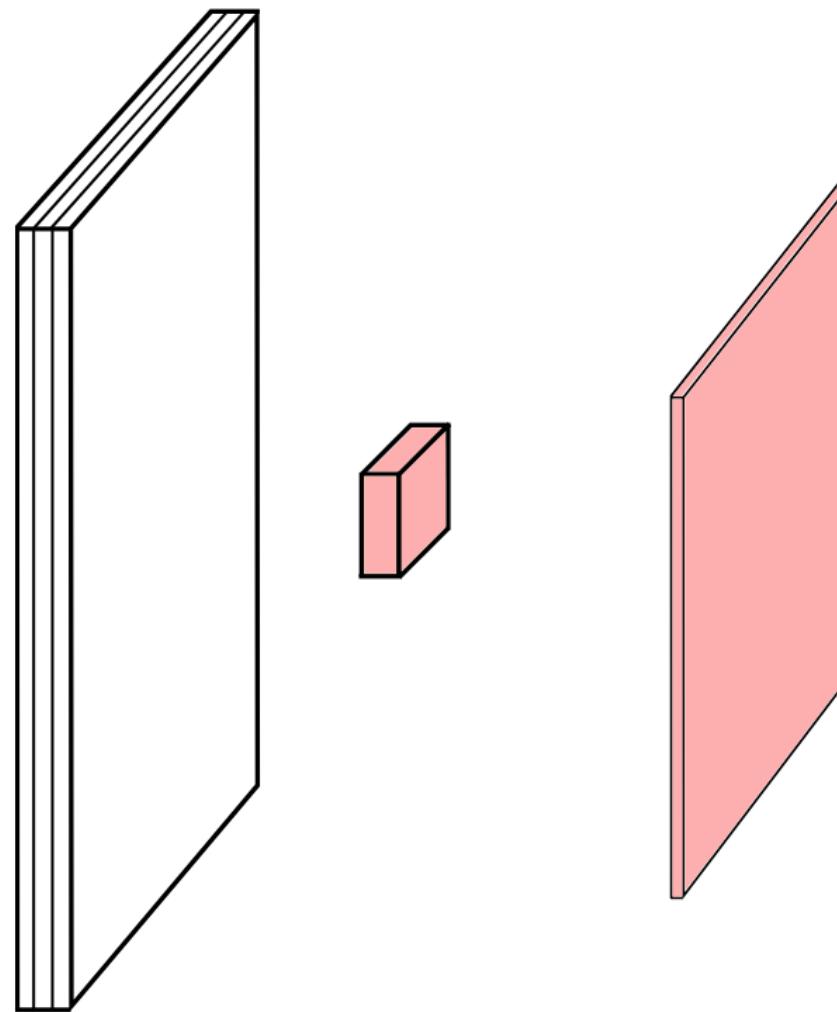
CONVOLUTIONS CAN ALSO BE COMPUTED ACROSS CHANNELS (OR COLORS)

A COLOR IMAGE IS A
TENSOR
OF SIZE height x width x
channels



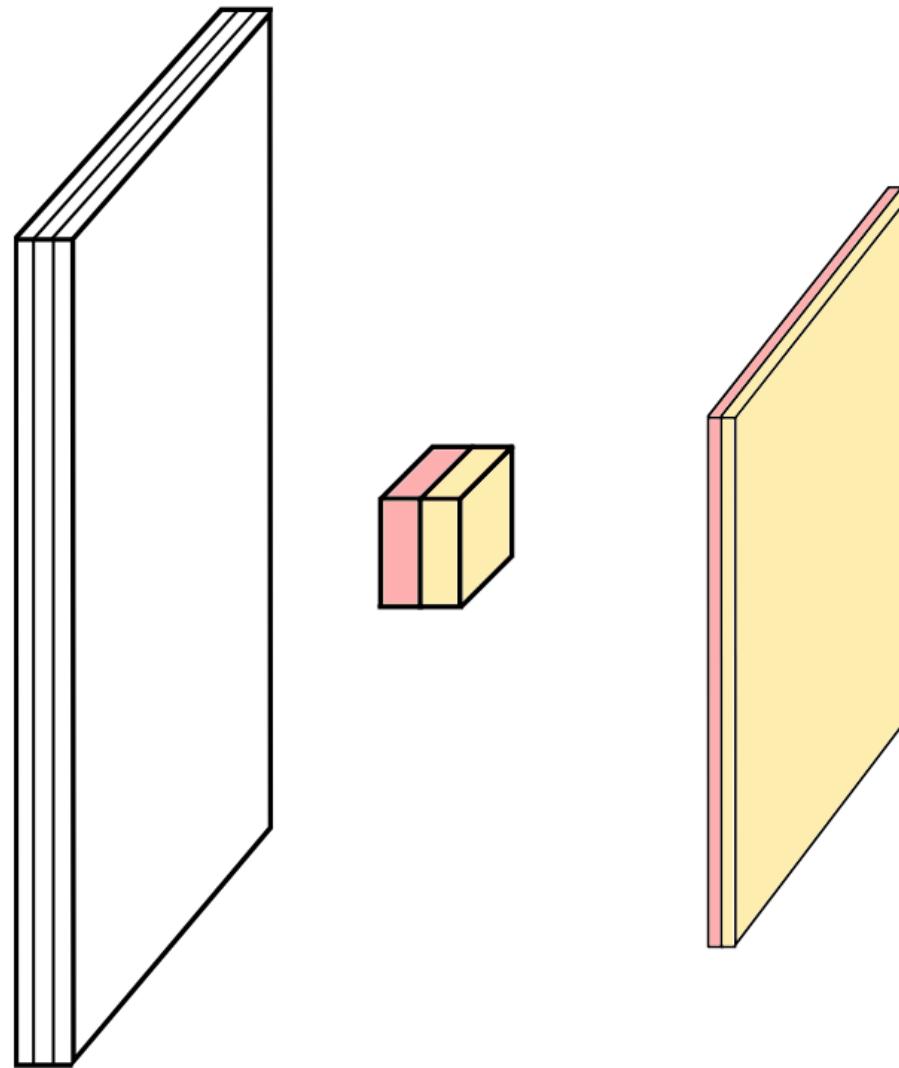
THEN THE KERNEL
HAS ALSO 3
CHANNELS

MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



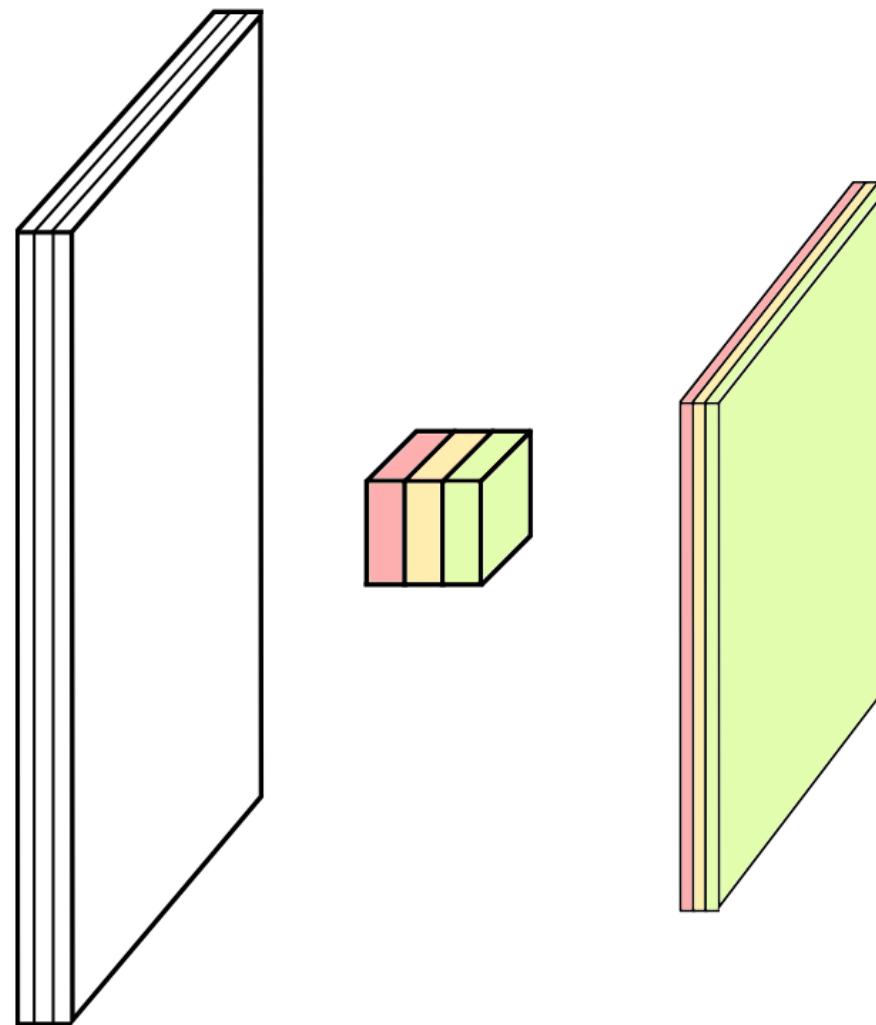
credit

MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



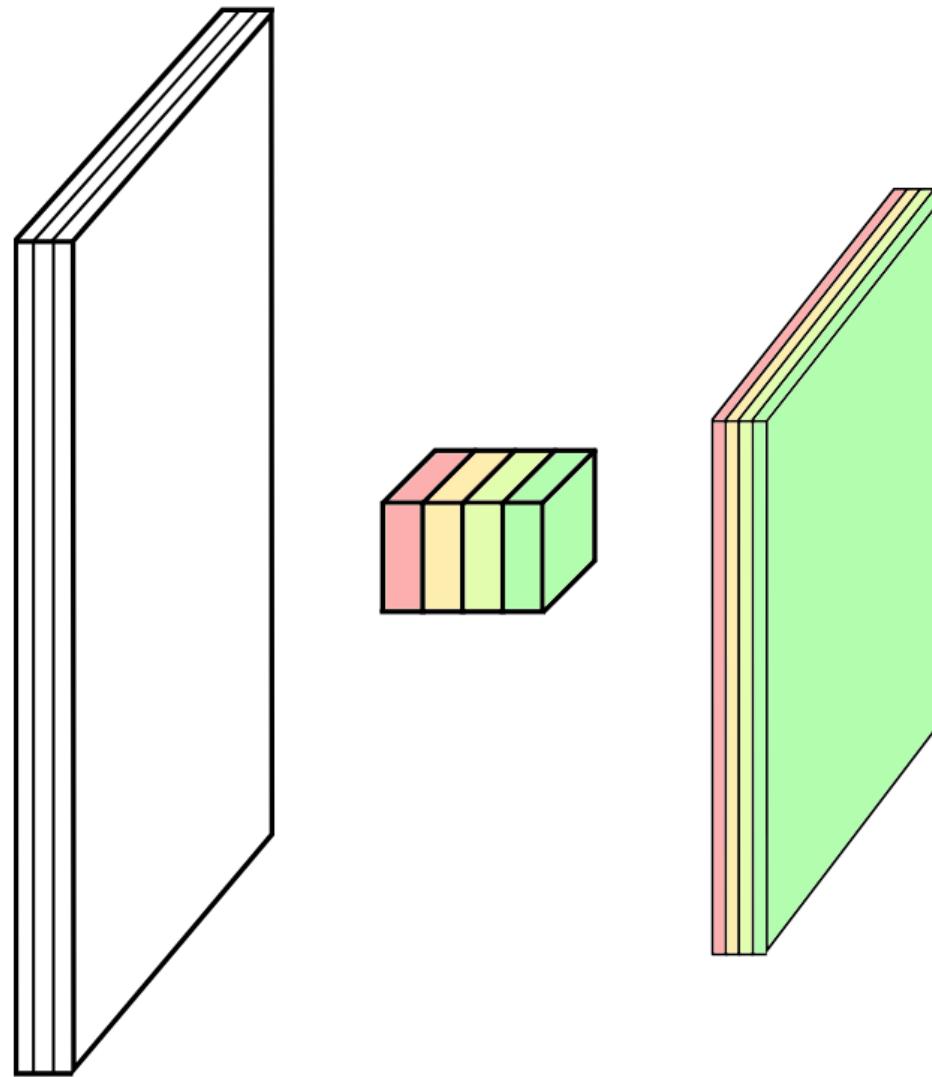
credit

MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



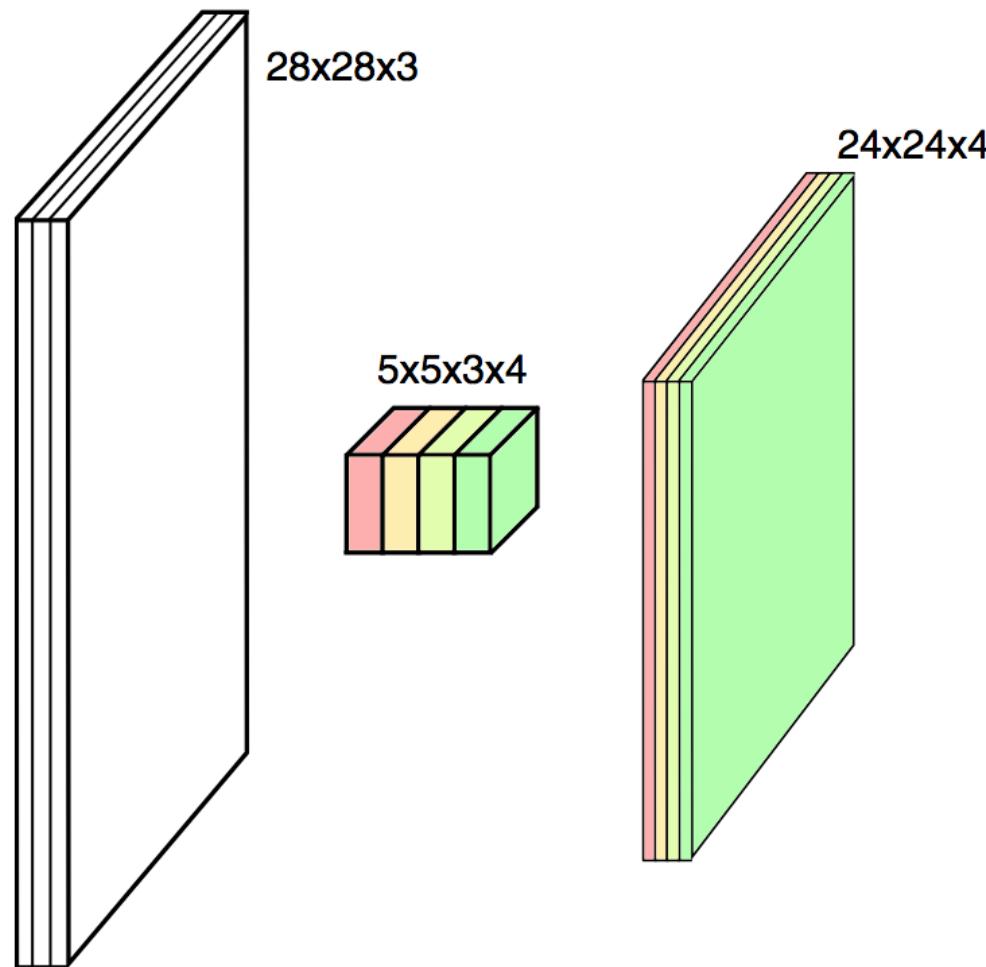
credit

MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



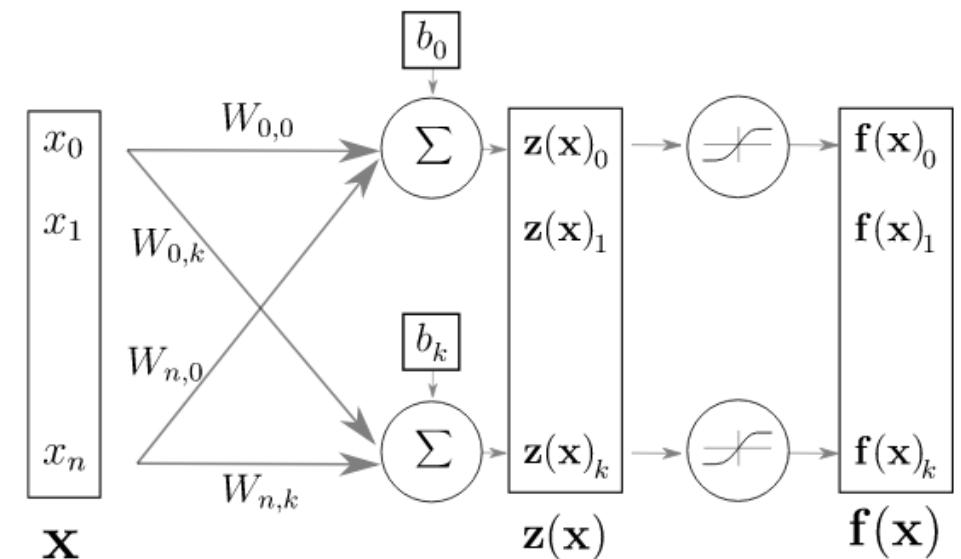
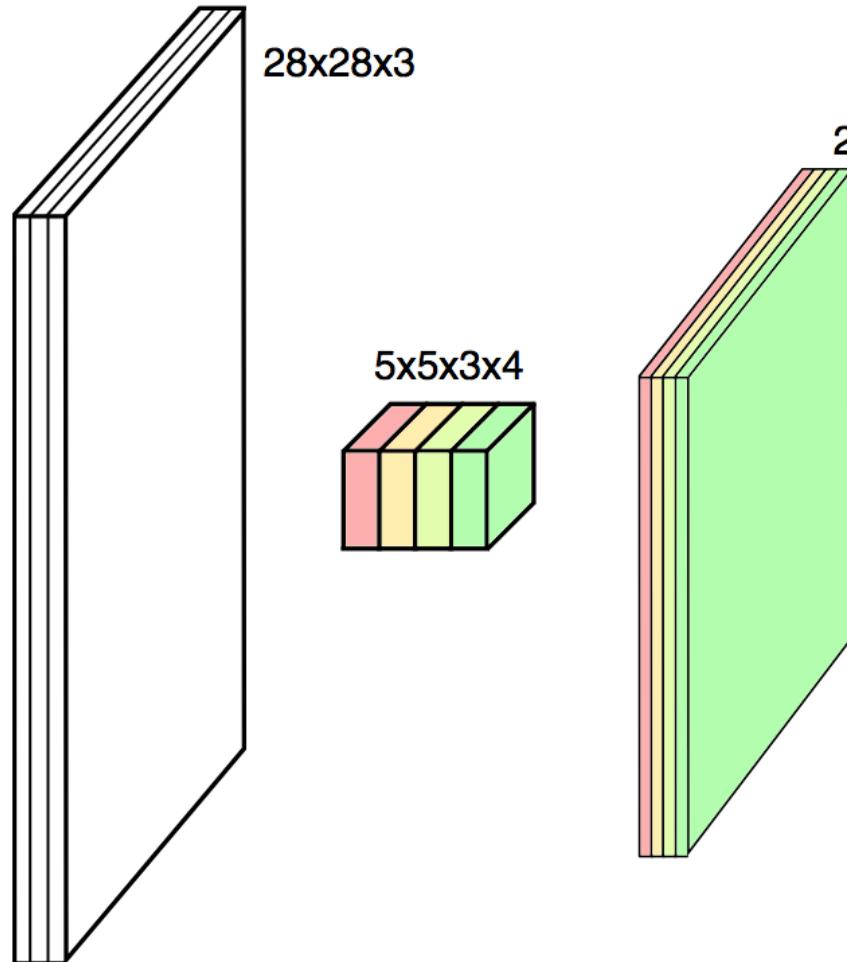
credit

MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



credit

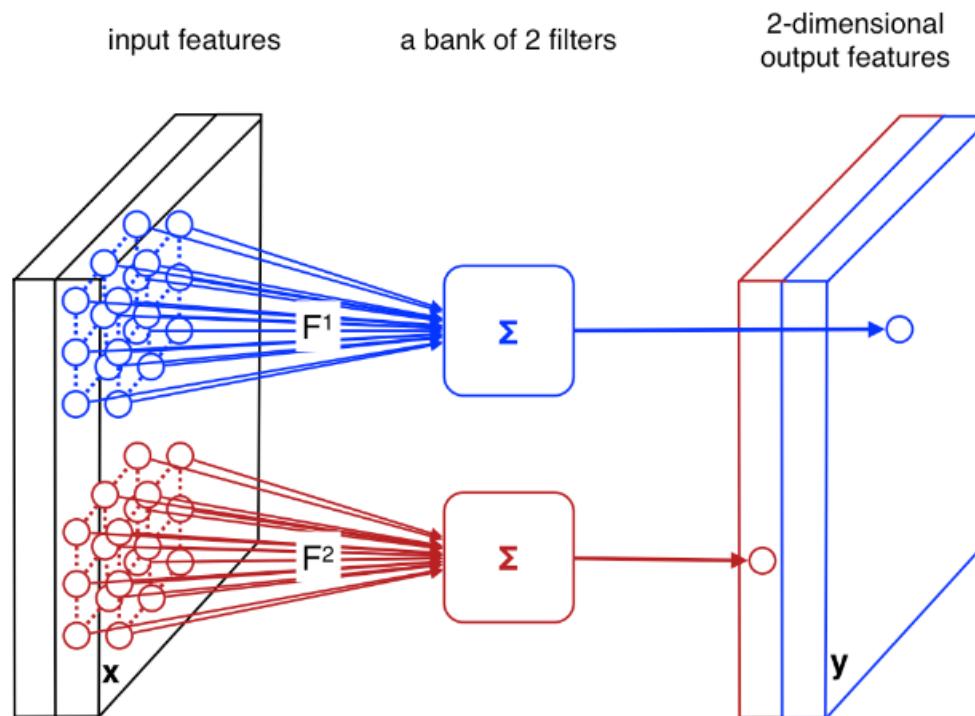
MULTIPLE CONVOLUTIONS WITH DIFFERENT KERNELS CAN BE PERFORMED



credit

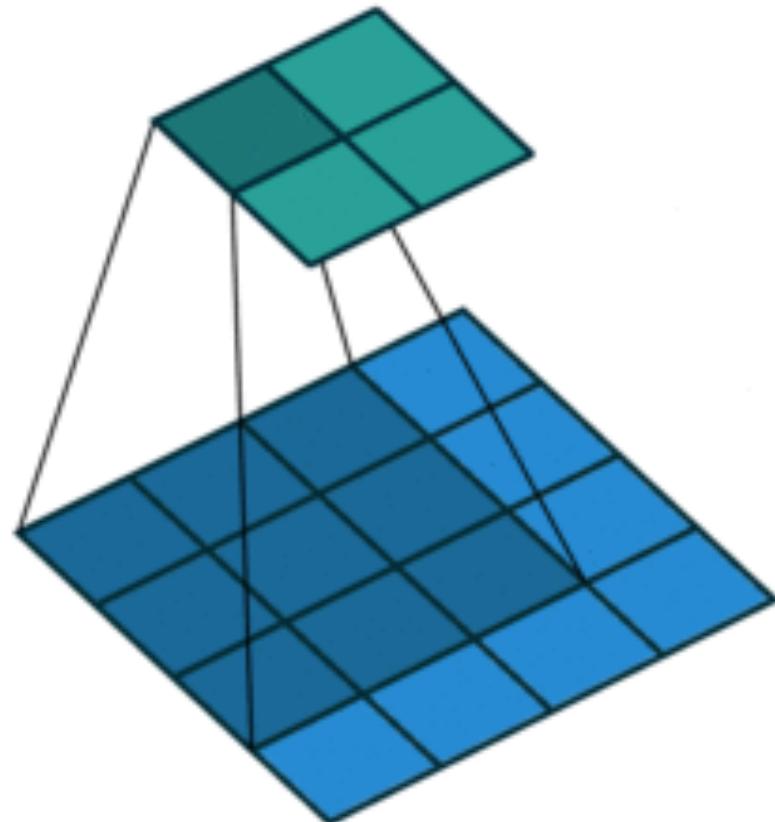
SINCE CONVOLUTIONS OUTPUT ONE SCALAR< THEY CAN BE SEEN AS AN INDIVIDUAL NEURON WITH A RECEPTIVE FIELD LIMITED TO THE KERNEL DIMENSIONS

THE SAME NEURON IS FIRED WITH DIFFERENT AREAS FROM THE INPUT

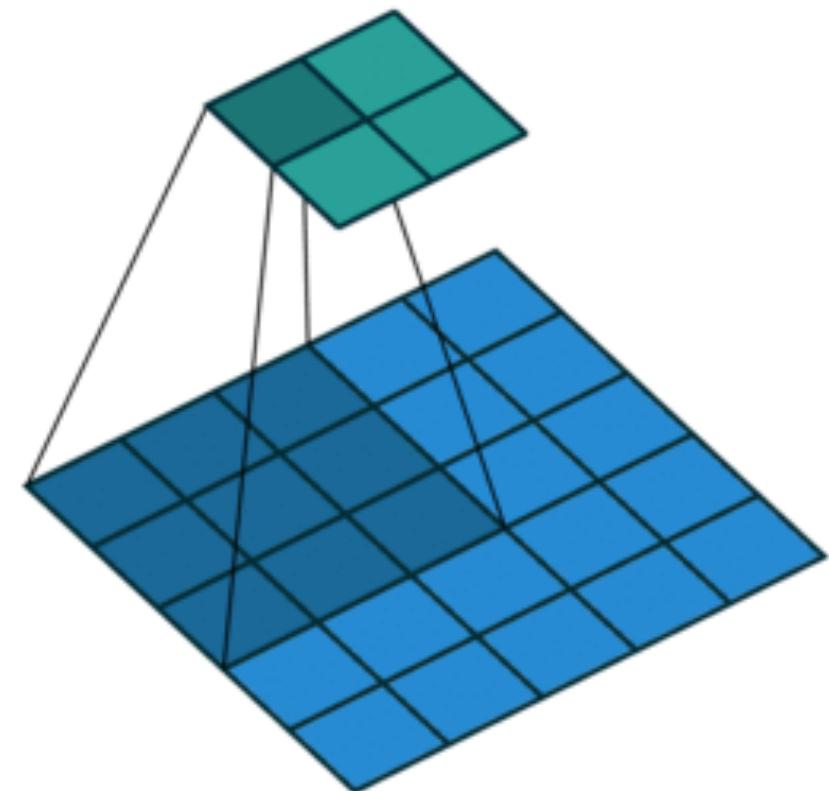


Credit

OPTIONS: STRIDES

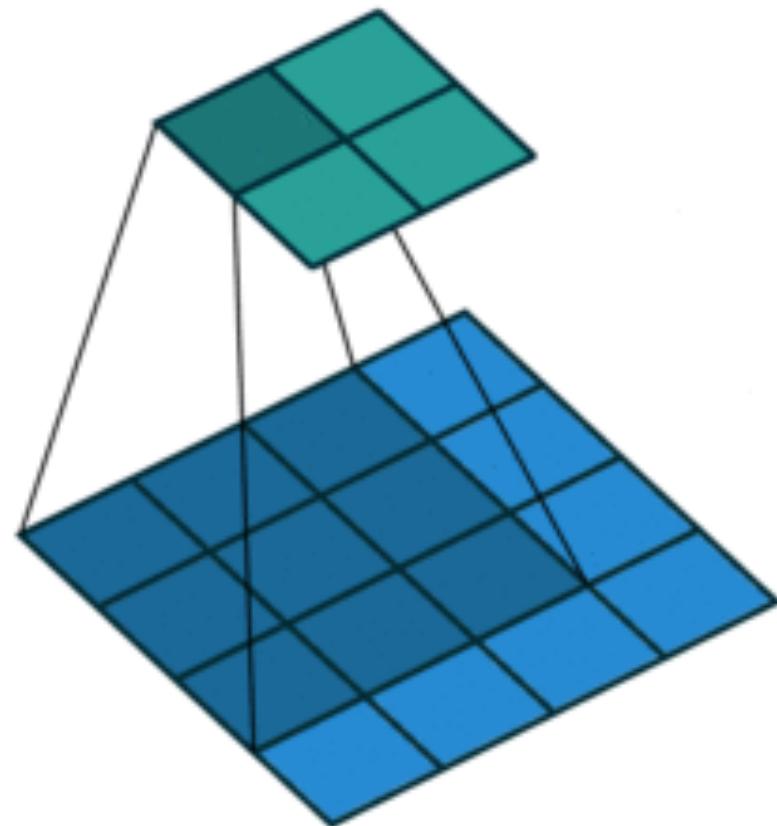


NO STRIDES

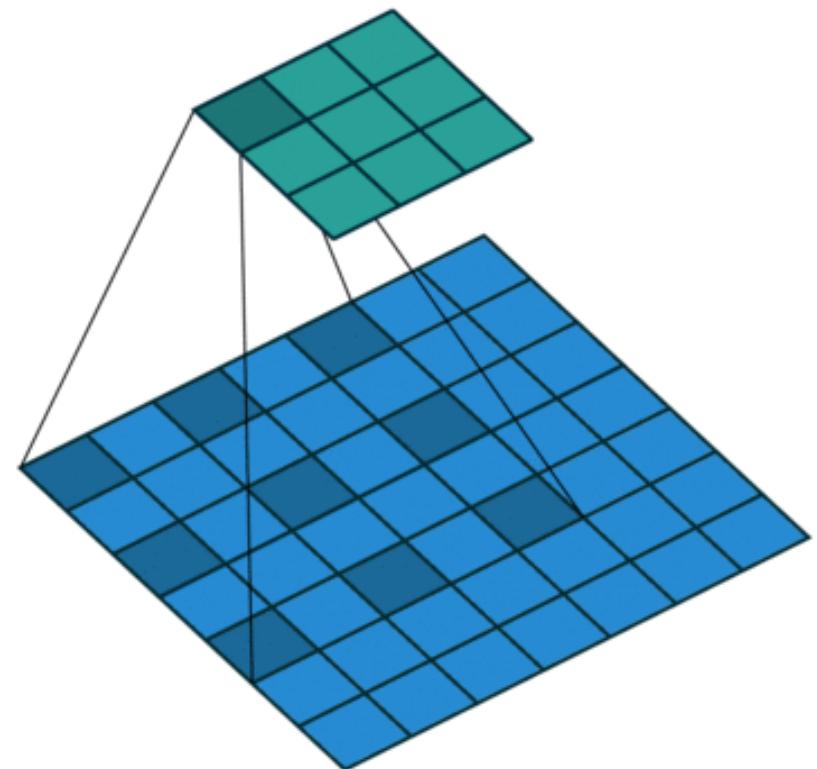


STRIDES

OPTIONS: DILATION

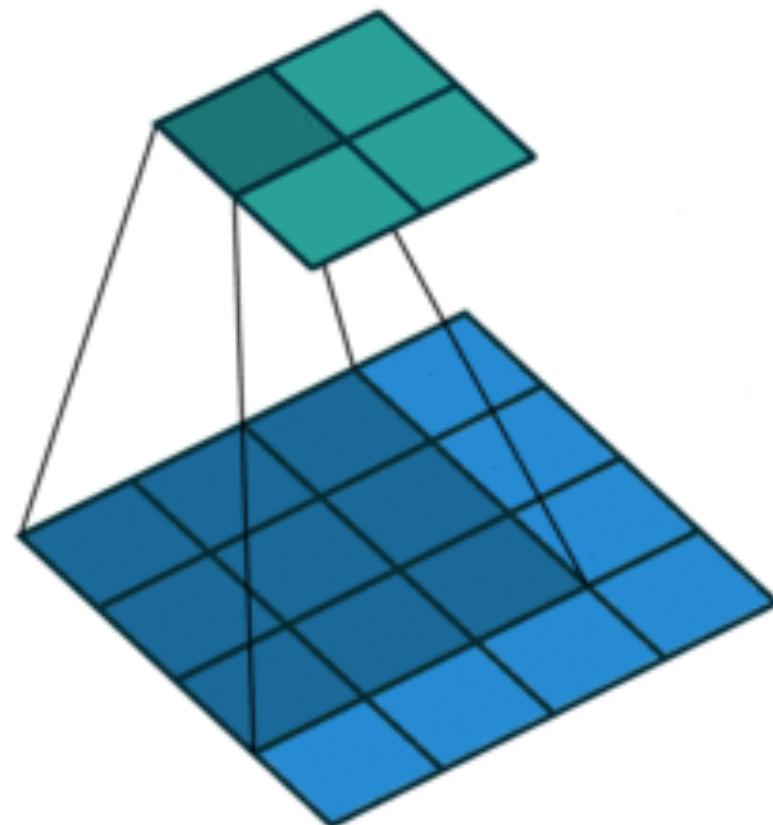


NO STRIDES

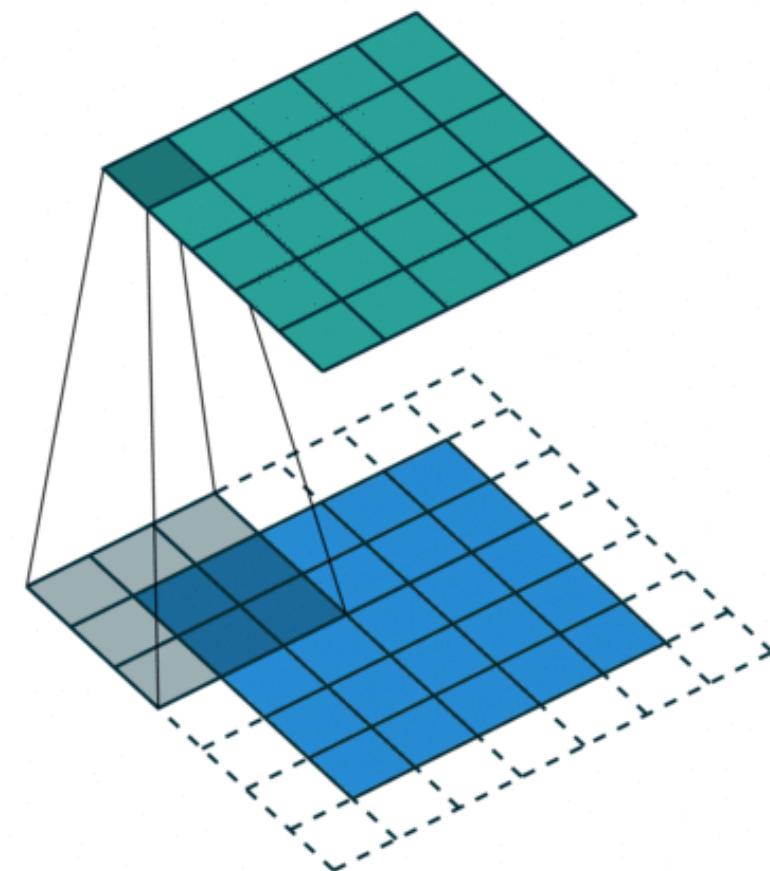


DILATION

OPTIONS: PADDING

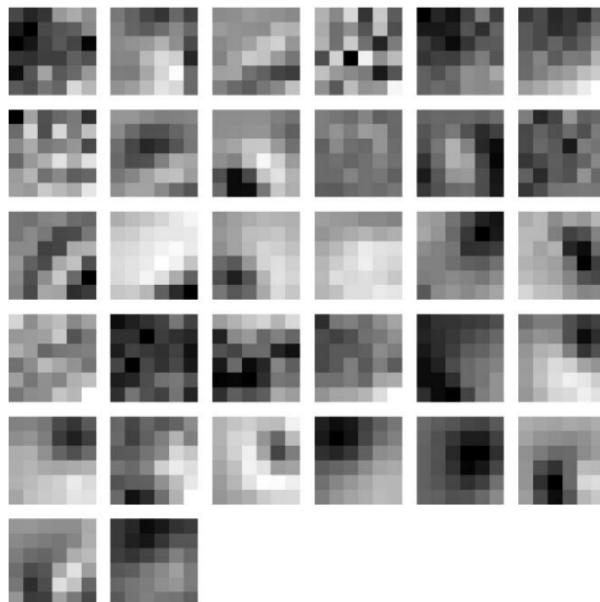


NO STRIDES

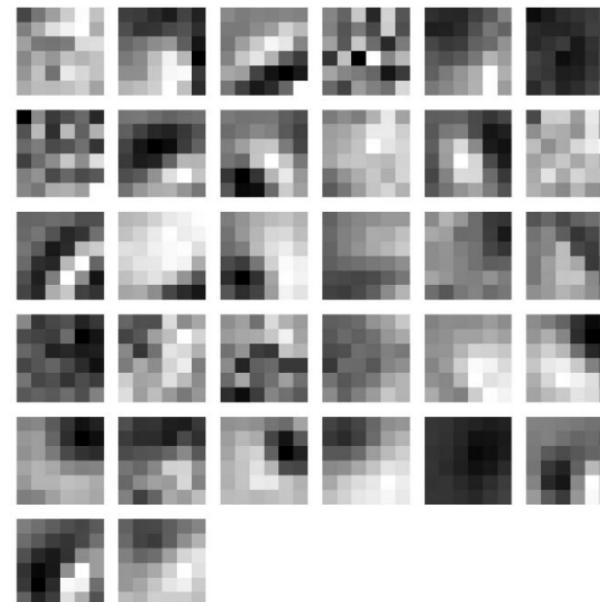


PADDING

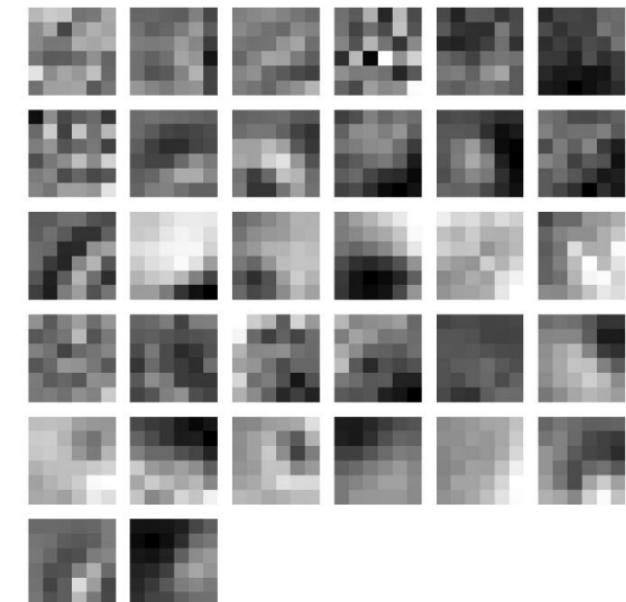
EXAMPLE OF 32 FILTERS LEARNED IN A CONVOLUTIONAL LAYER



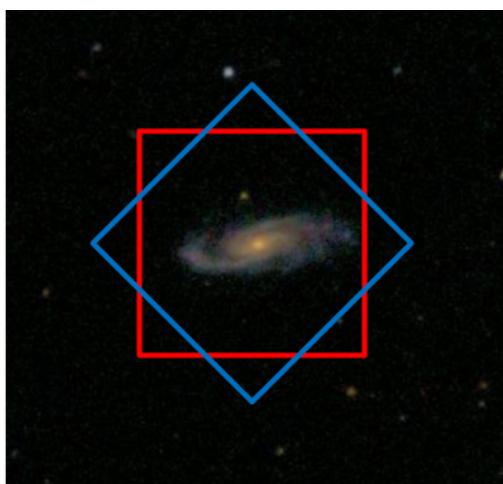
(a) red channel



(b) green channel



(c) blue channel



Dieleman+16

ESTIMATING SHAPES AND NUMBER OF PARAMETERS

KERNEL SHAPE:

$$(F, F, C^i, C^o)$$

PADDING:

$$P$$

STRIDES:

$$S$$

OUTPUT SIZE:

$$W_0 = (W^i - F + 2P)/S + 1$$

NUMBER OF PARAMETERS:

$$(F \times F \times C^i + 1) \times C^o$$



the number of parameters increases fast!

32 filters of 5*5 on a color image —> 2432 parameters to learn

POOLING

CONVOLUTIONS ARE OFTEN FOLLOWED BY AN OPERATION OF DOWNSAMPLING [POOLING]

VERY SIMPLE OPERATION - ONLY ONE OUT OF EVERY N PIXELS ARE KEPT

OFTEN MATCHED WITH AN INCREASE OF THE FEATURE CHANNELS

TYPES OF POOLING

SUM POOLING

$$y = \sum x_{uv}$$

SQUARE SUM POOLING

$$y = \sqrt{\sum x_{uv}^2}$$

MAX POOLING

$$y = \max(x_{uv})$$

TYPES OF POOLING

SUM POOLING

$$y = \sum x_{uv}$$

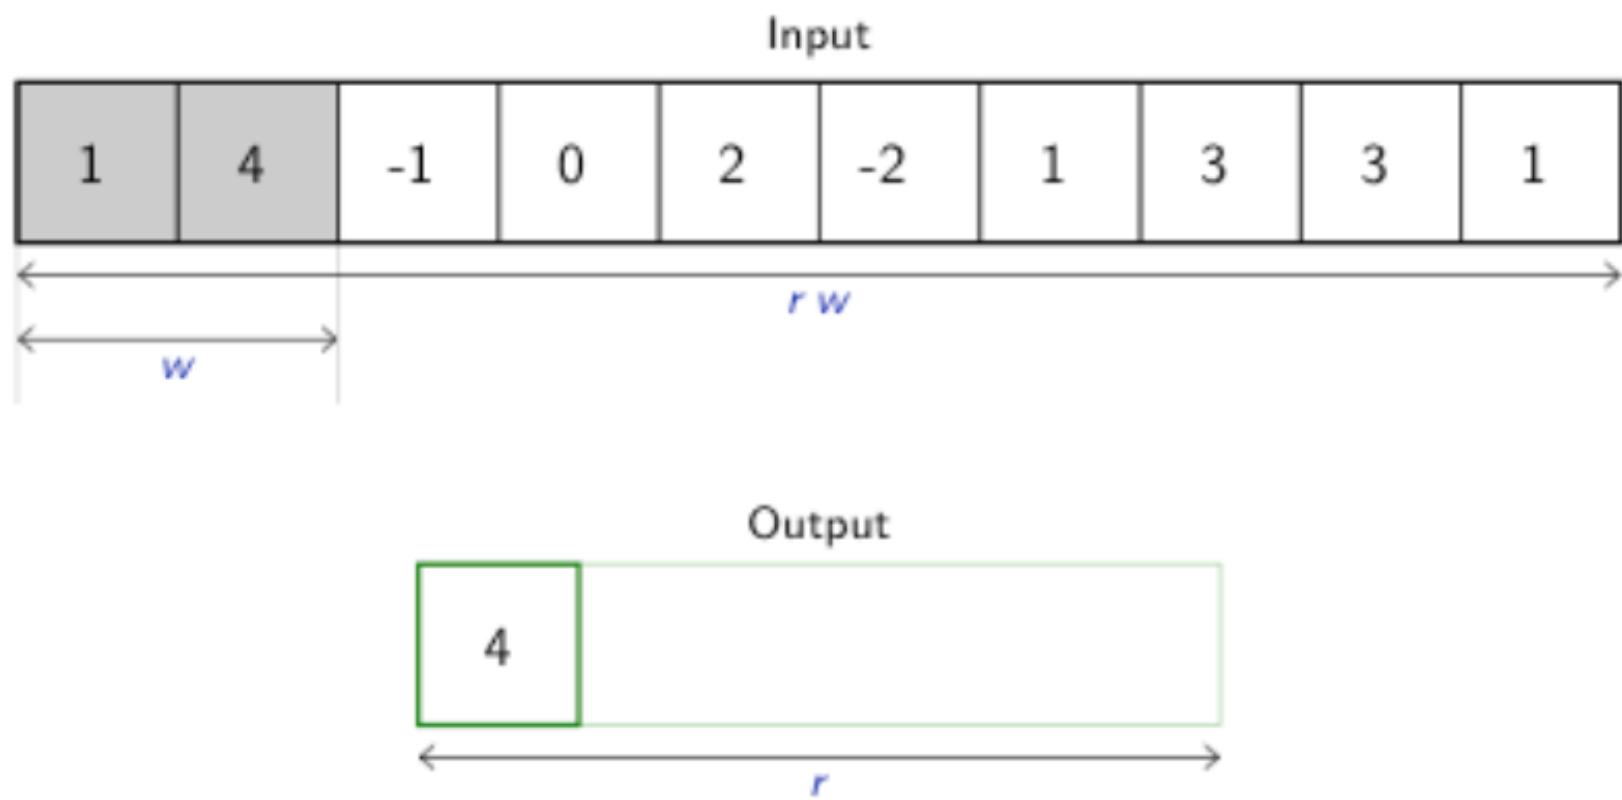
SQUARE SUM POOLING

$$y = \sqrt{\sum x_{uv}^2}$$

MAX POOLING

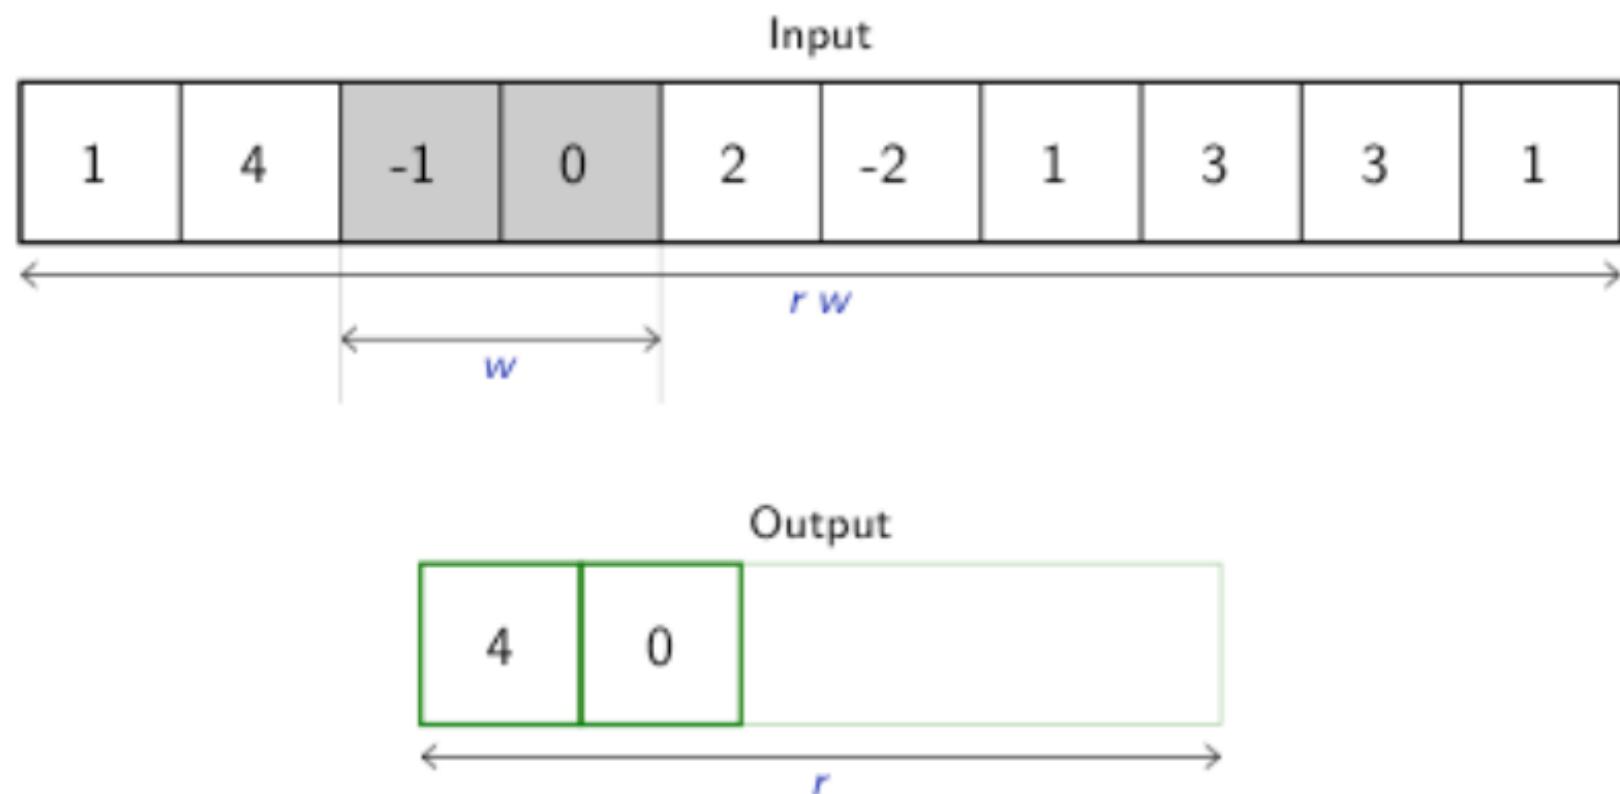
$$y = \max(x_{uv})$$

MAX POOLING 1D



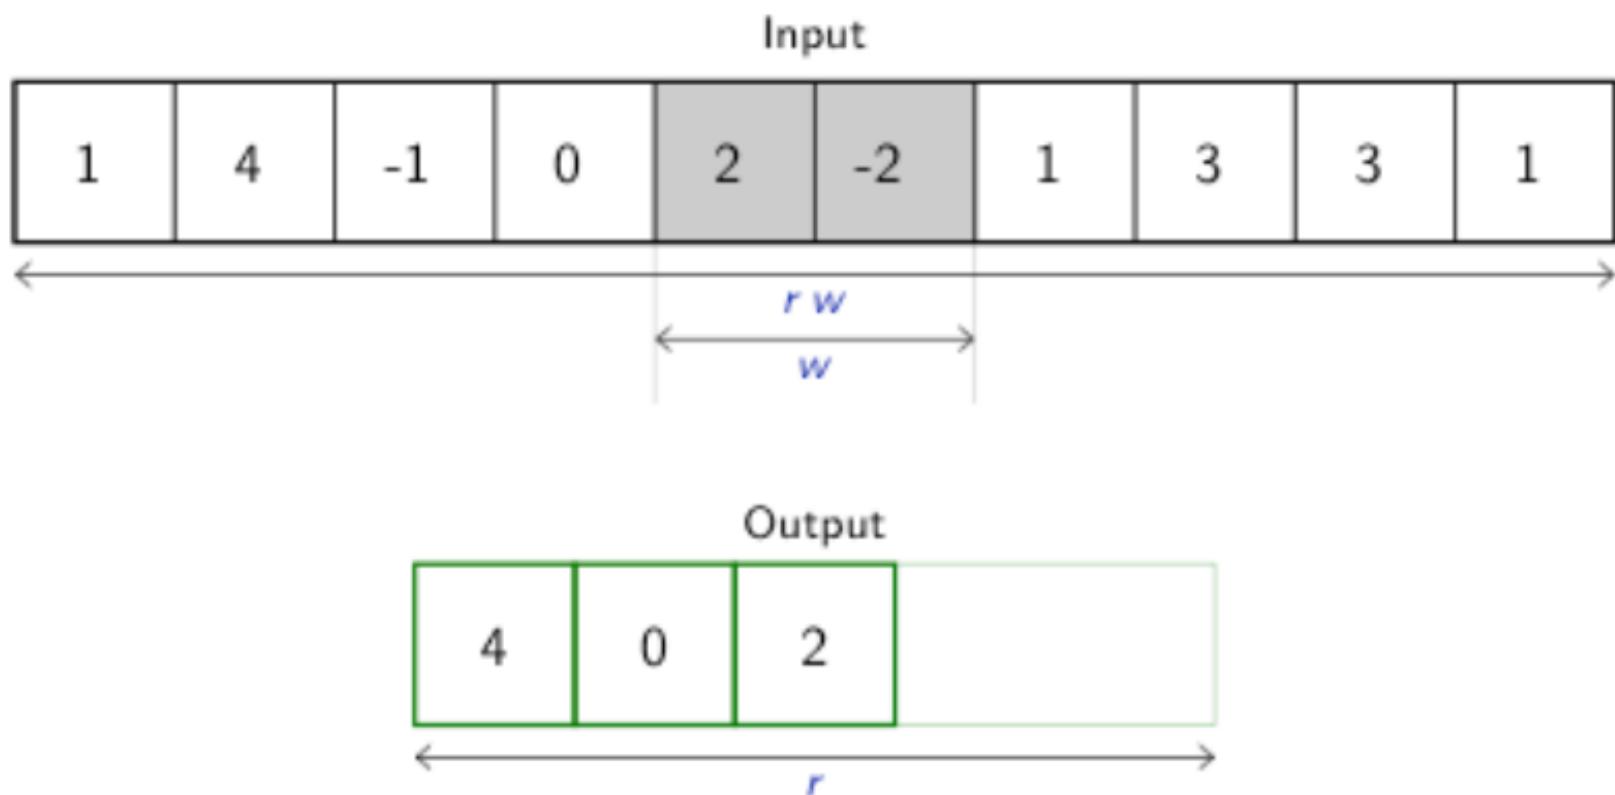
Credit: F. Fleuret

MAX POOLING 1D



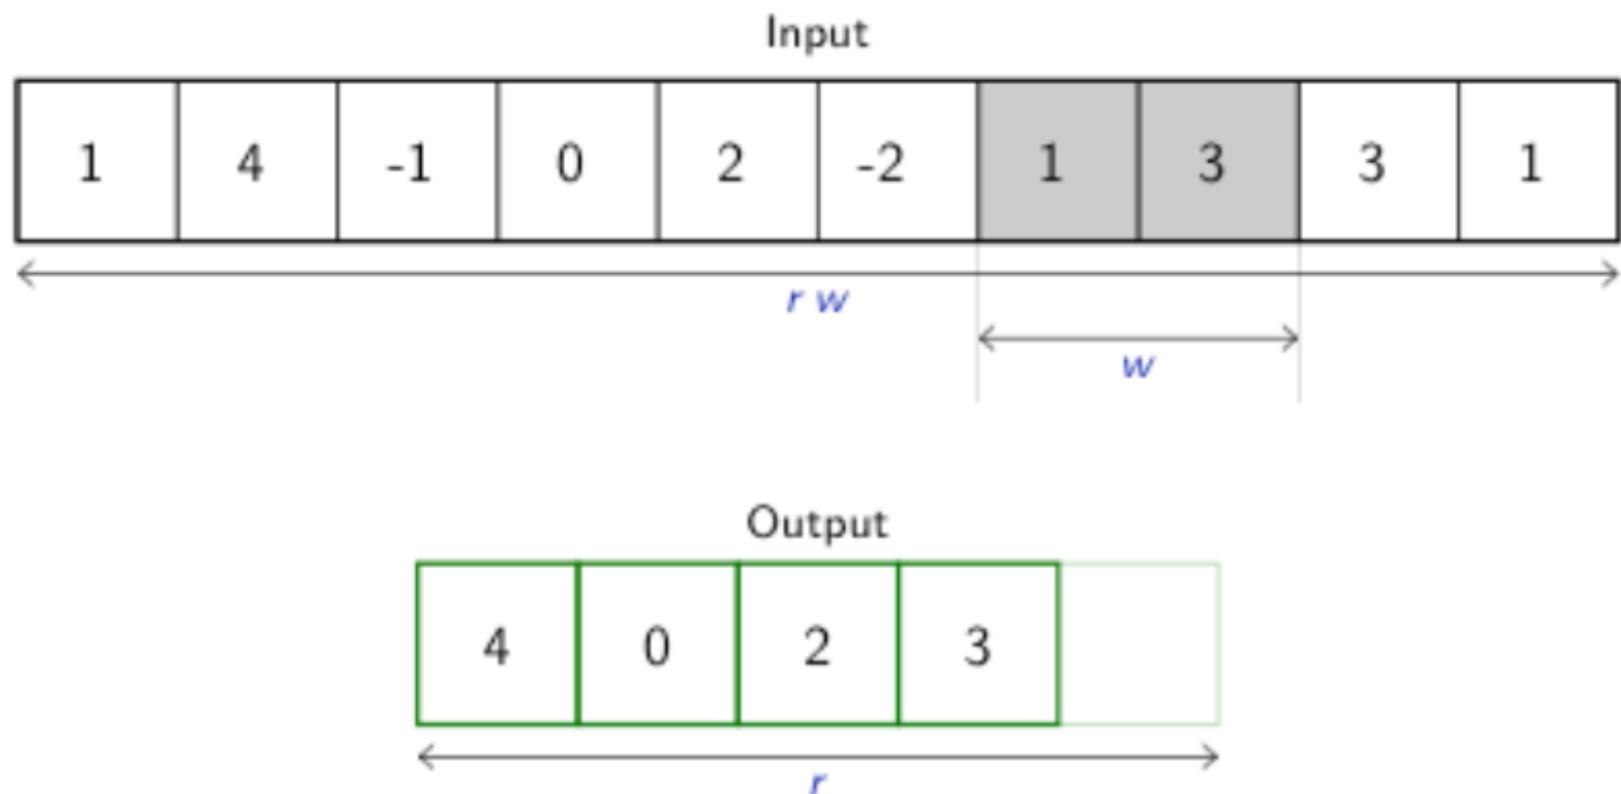
Credit: F. Fleuret

MAX POOLING 1D



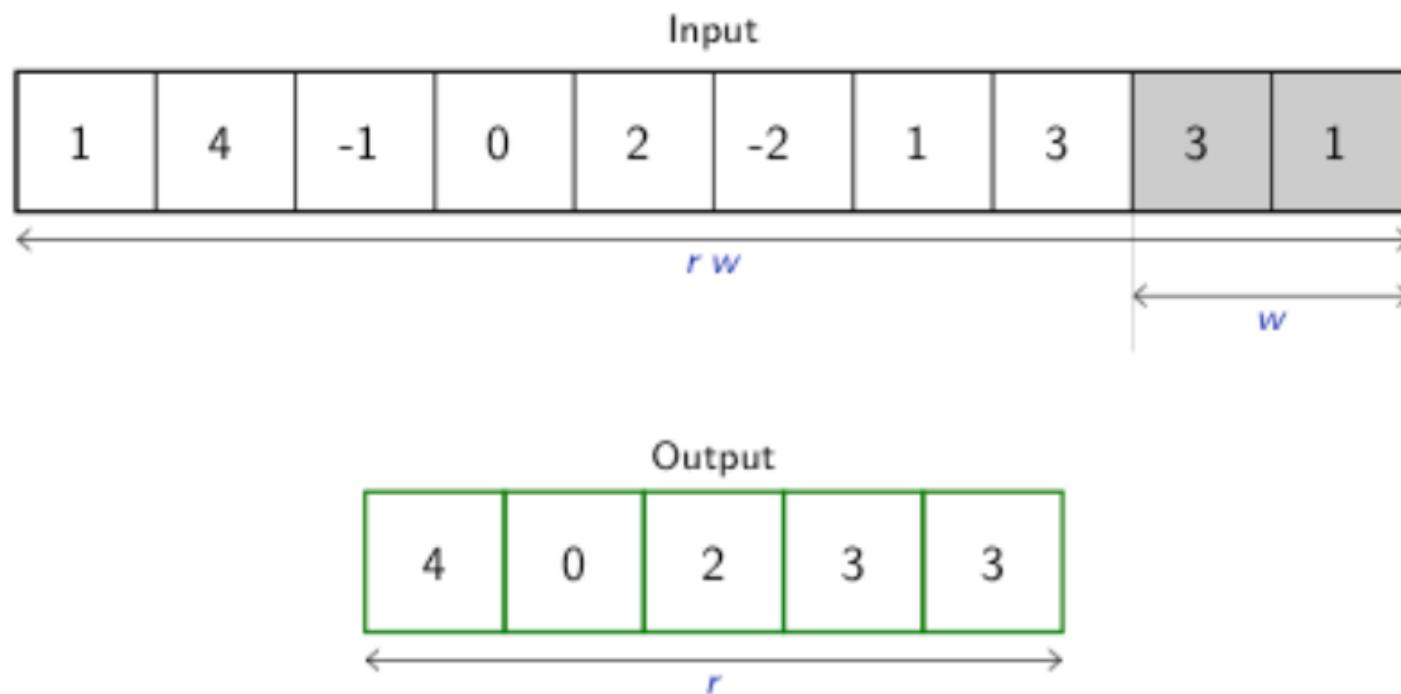
Credit: F. Fleuret

MAX POOLING 1D



Credit: F. Fleuret

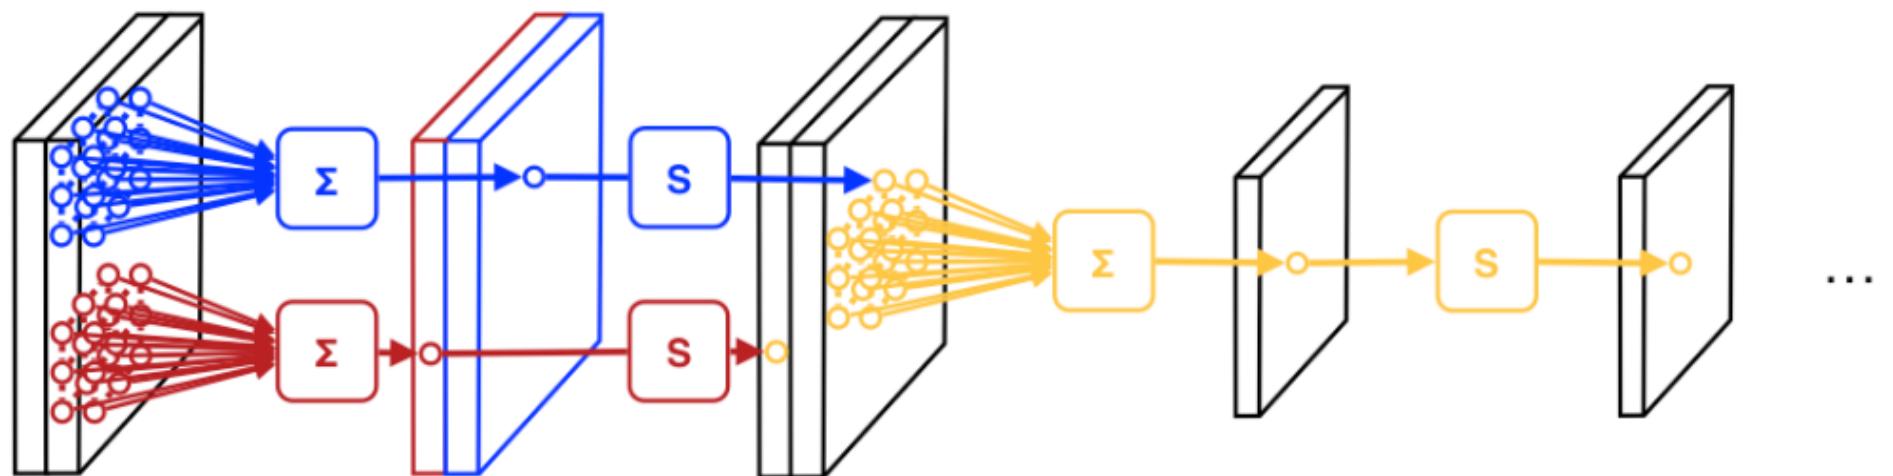
MAX POOLING 1D



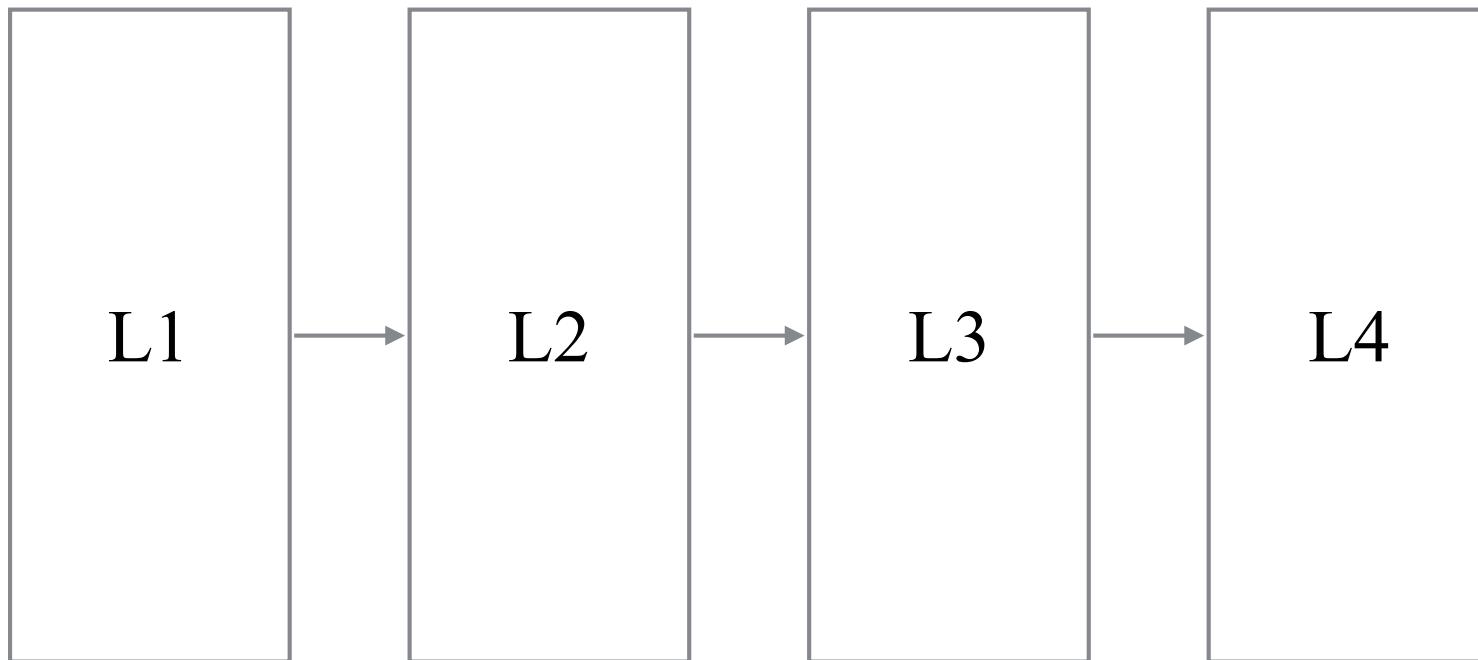
Credit: F. Fleuret

DOWNSAMPLING

DOWNSAMPLING IS APPLIED TO REDUCE THE OVERALL SIZE OF TENSORS AND CAPTURE LARGE SCALE STRUCTURE

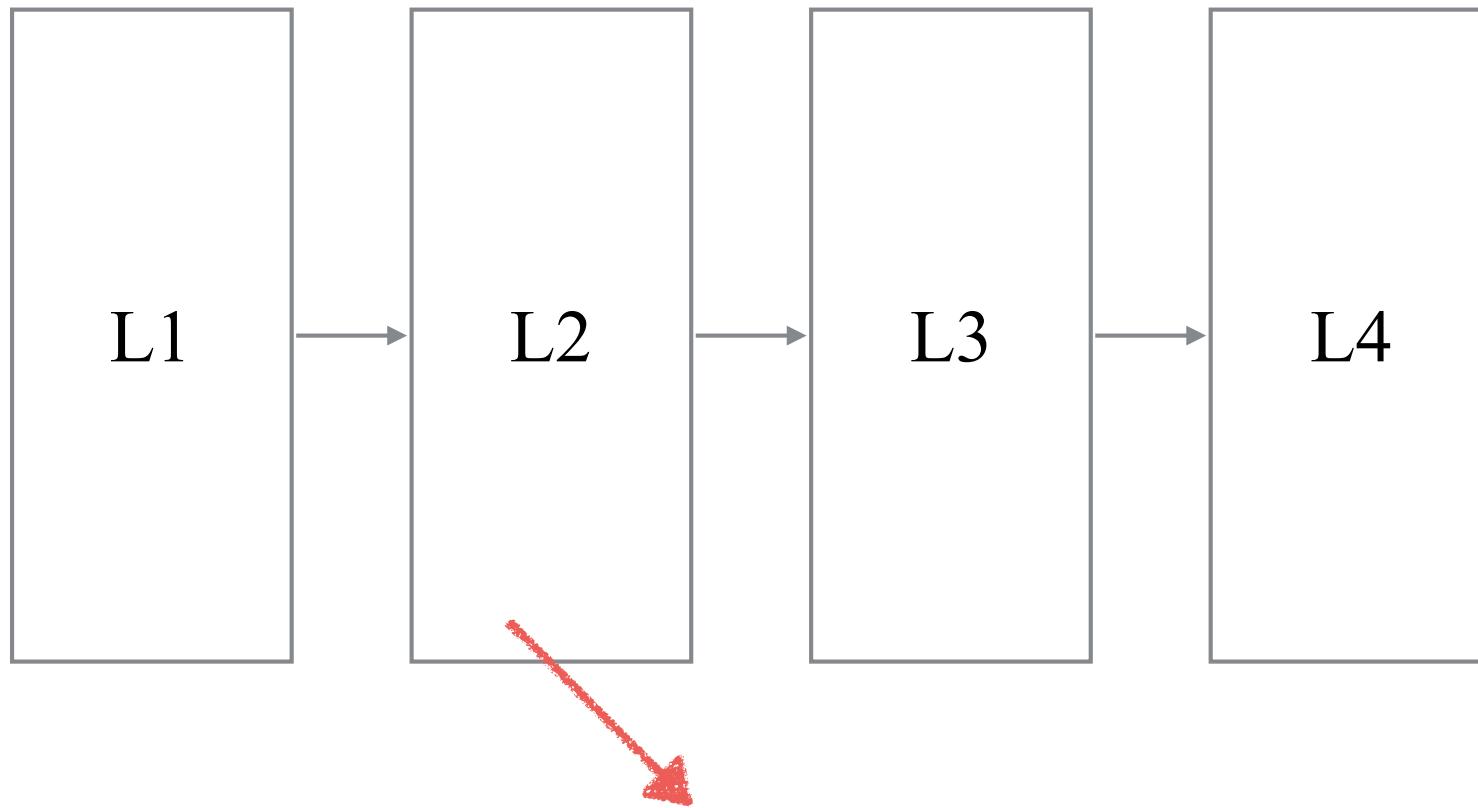


CONVNET OR CNN



A CONCATENATION OF MULTIPLE
CONVOLUTIONAL BLOCKS

CONVNET OR CNN



EACH BLOCK TYPICALLY MADE OF:

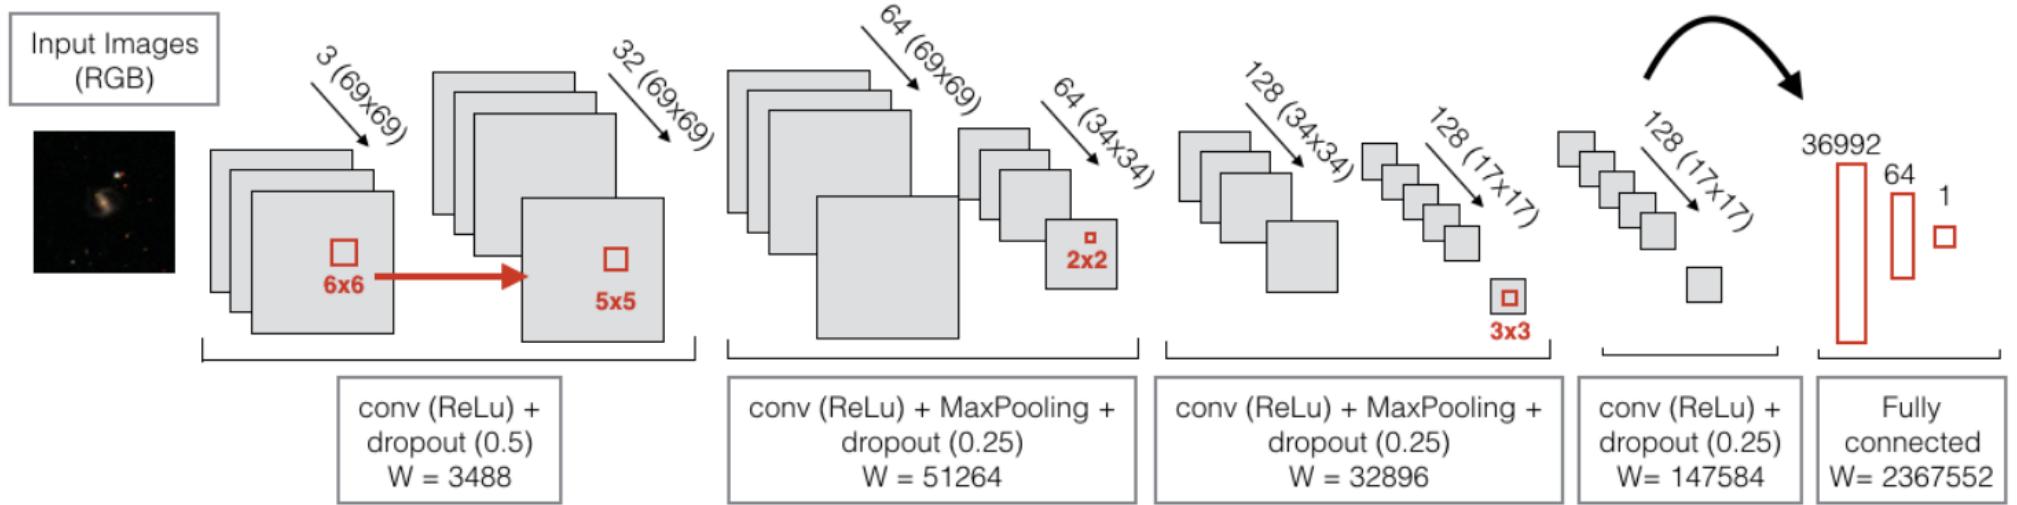
CONV

ACTIVATION

POOLING

(+dropout
for training)

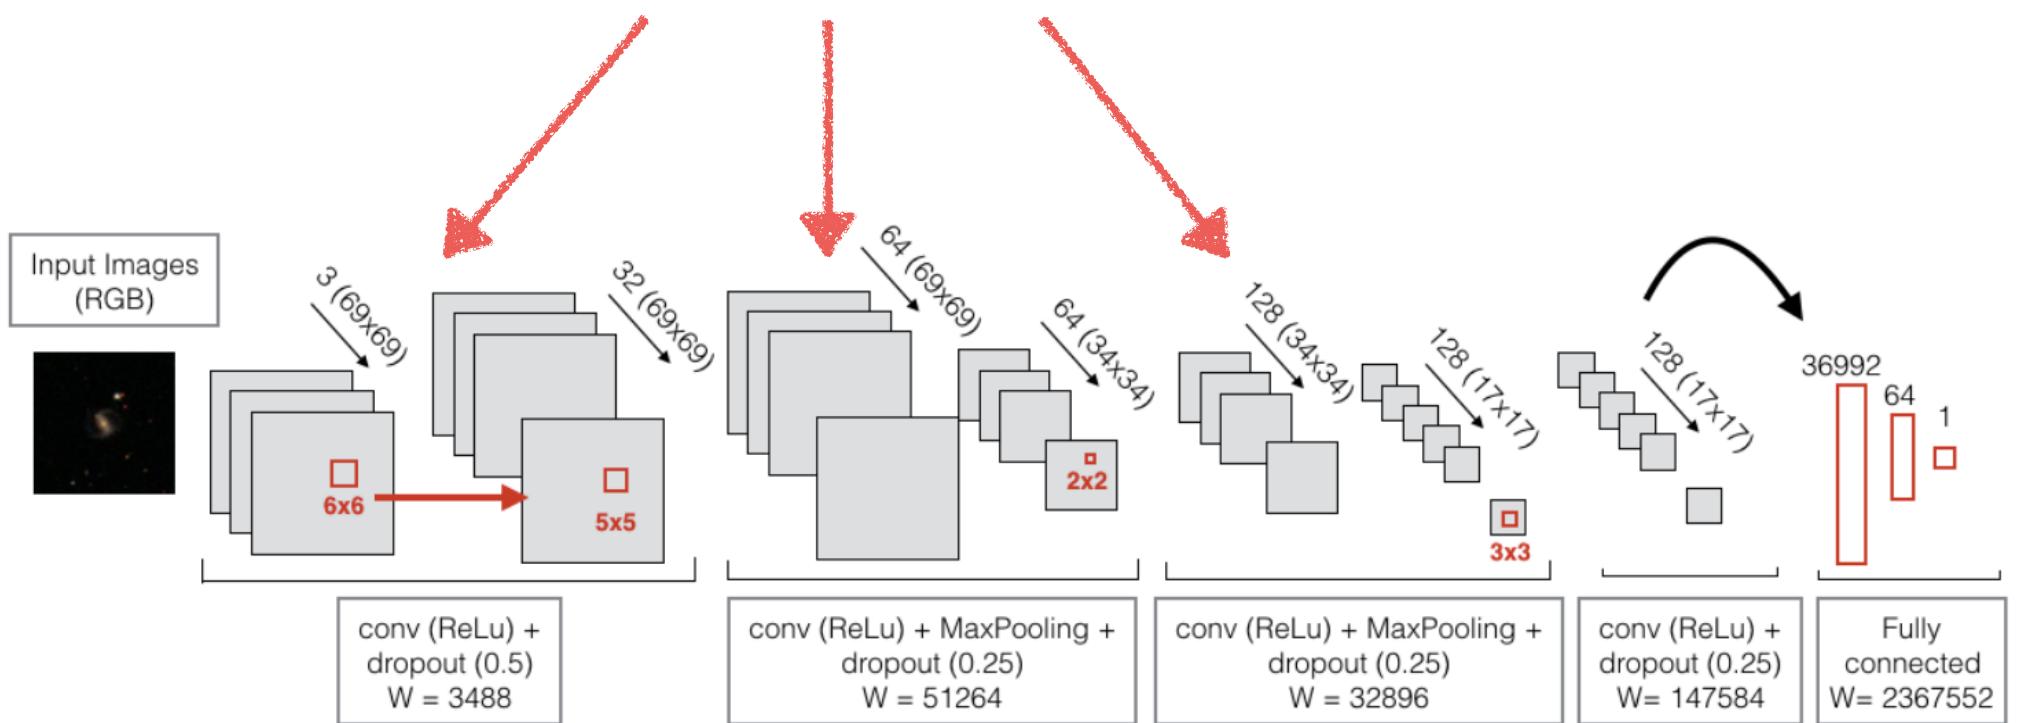
EXAMPLE OF VERY SIMPLE CNN



Dominguez-Sanchez+18

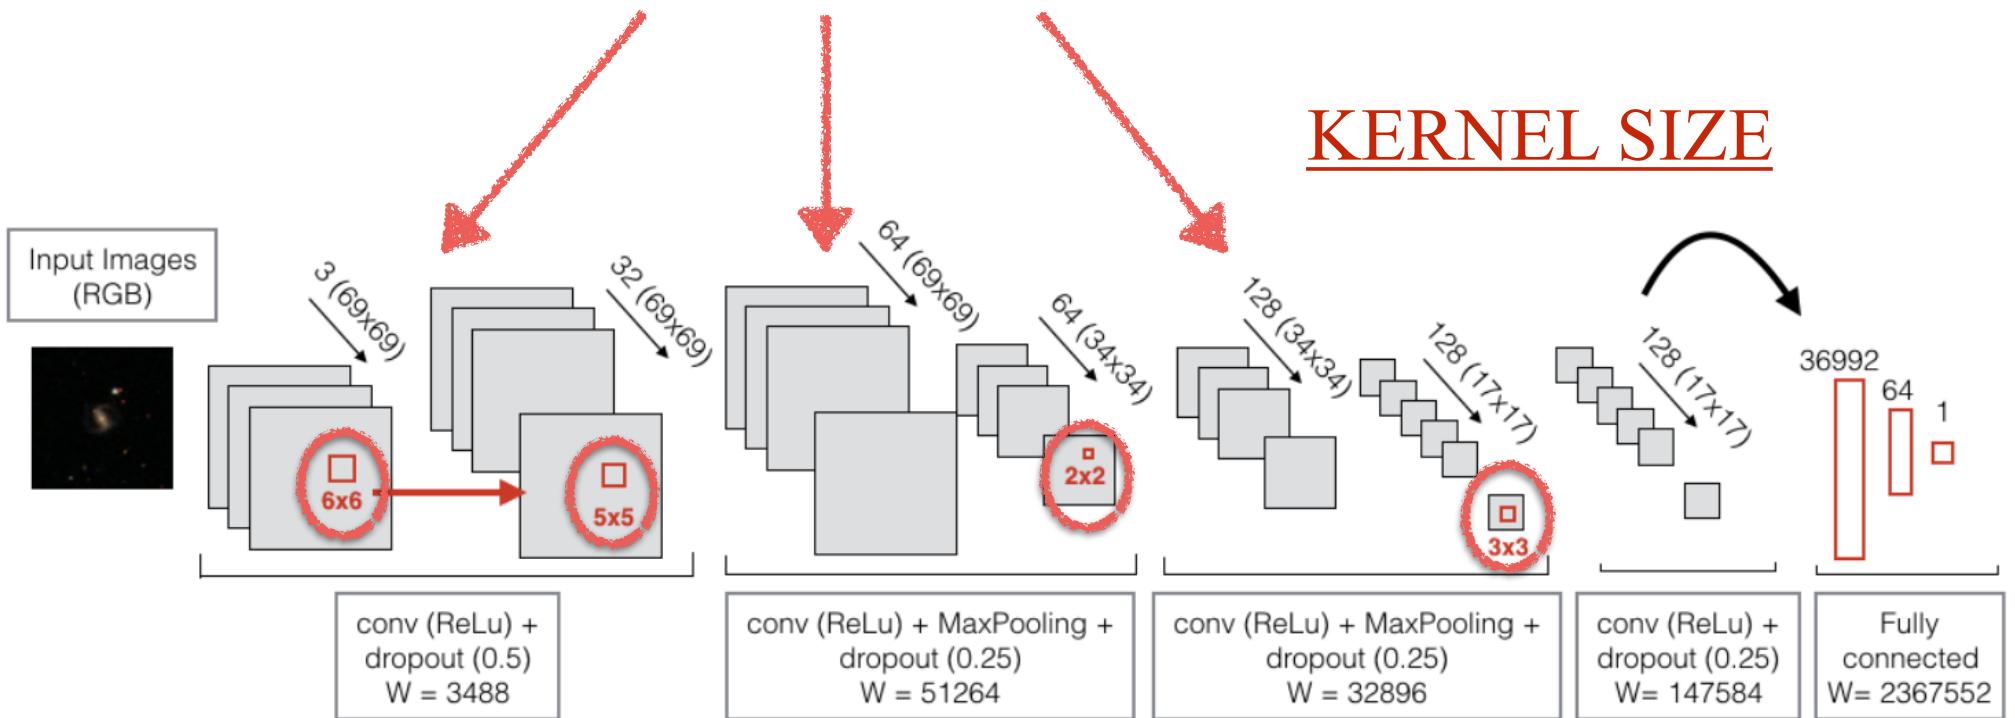
EXAMPLE OF VERY SIMPLE CNN

3 convolutional layers



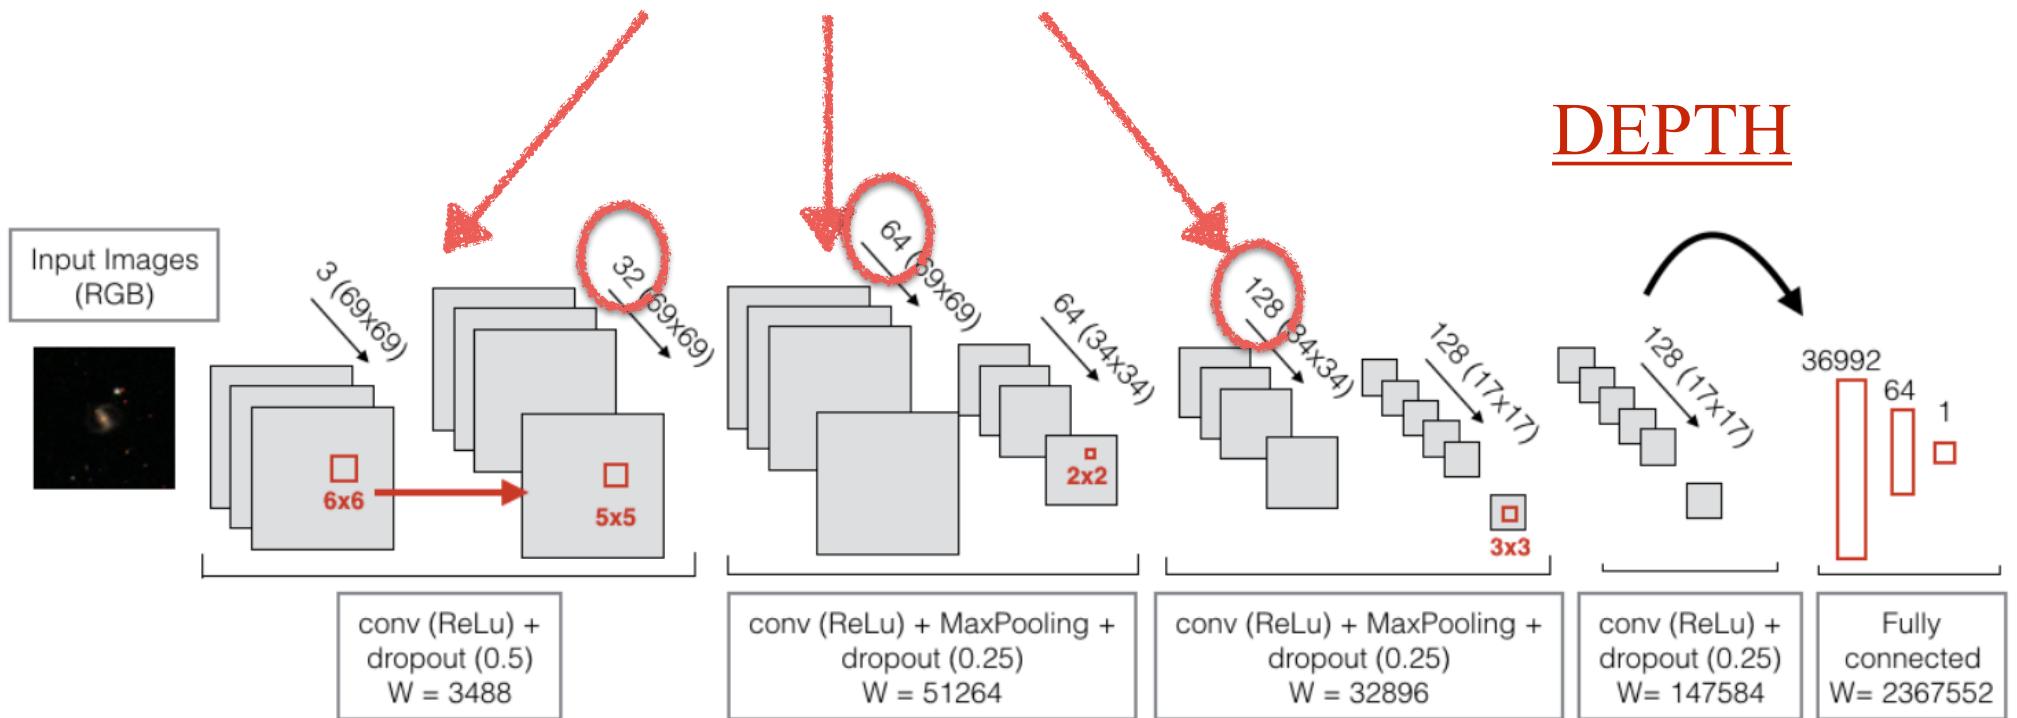
EXAMPLE OF VERY SIMPLE CNN

3 convolutional layers



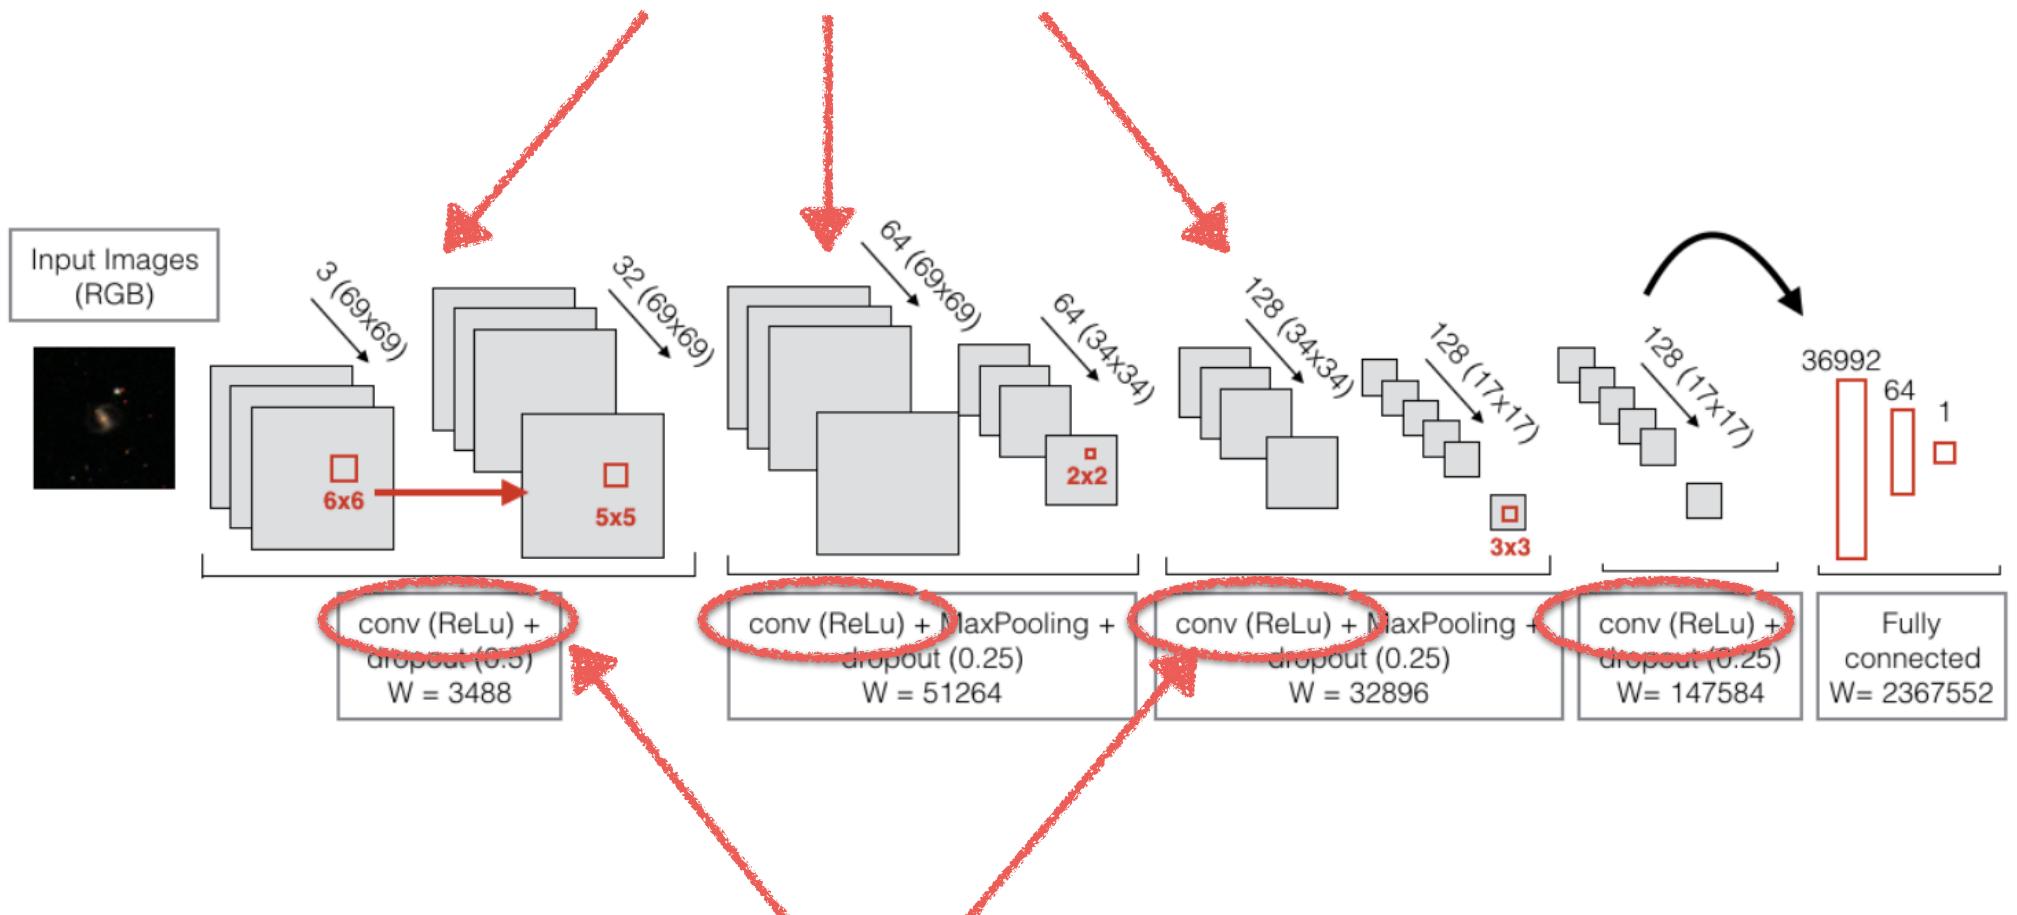
EXAMPLE OF VERY SIMPLE CNN

3 convolutional layers



EXAMPLE OF VERY SIMPLE CNN

3 convolutional layers

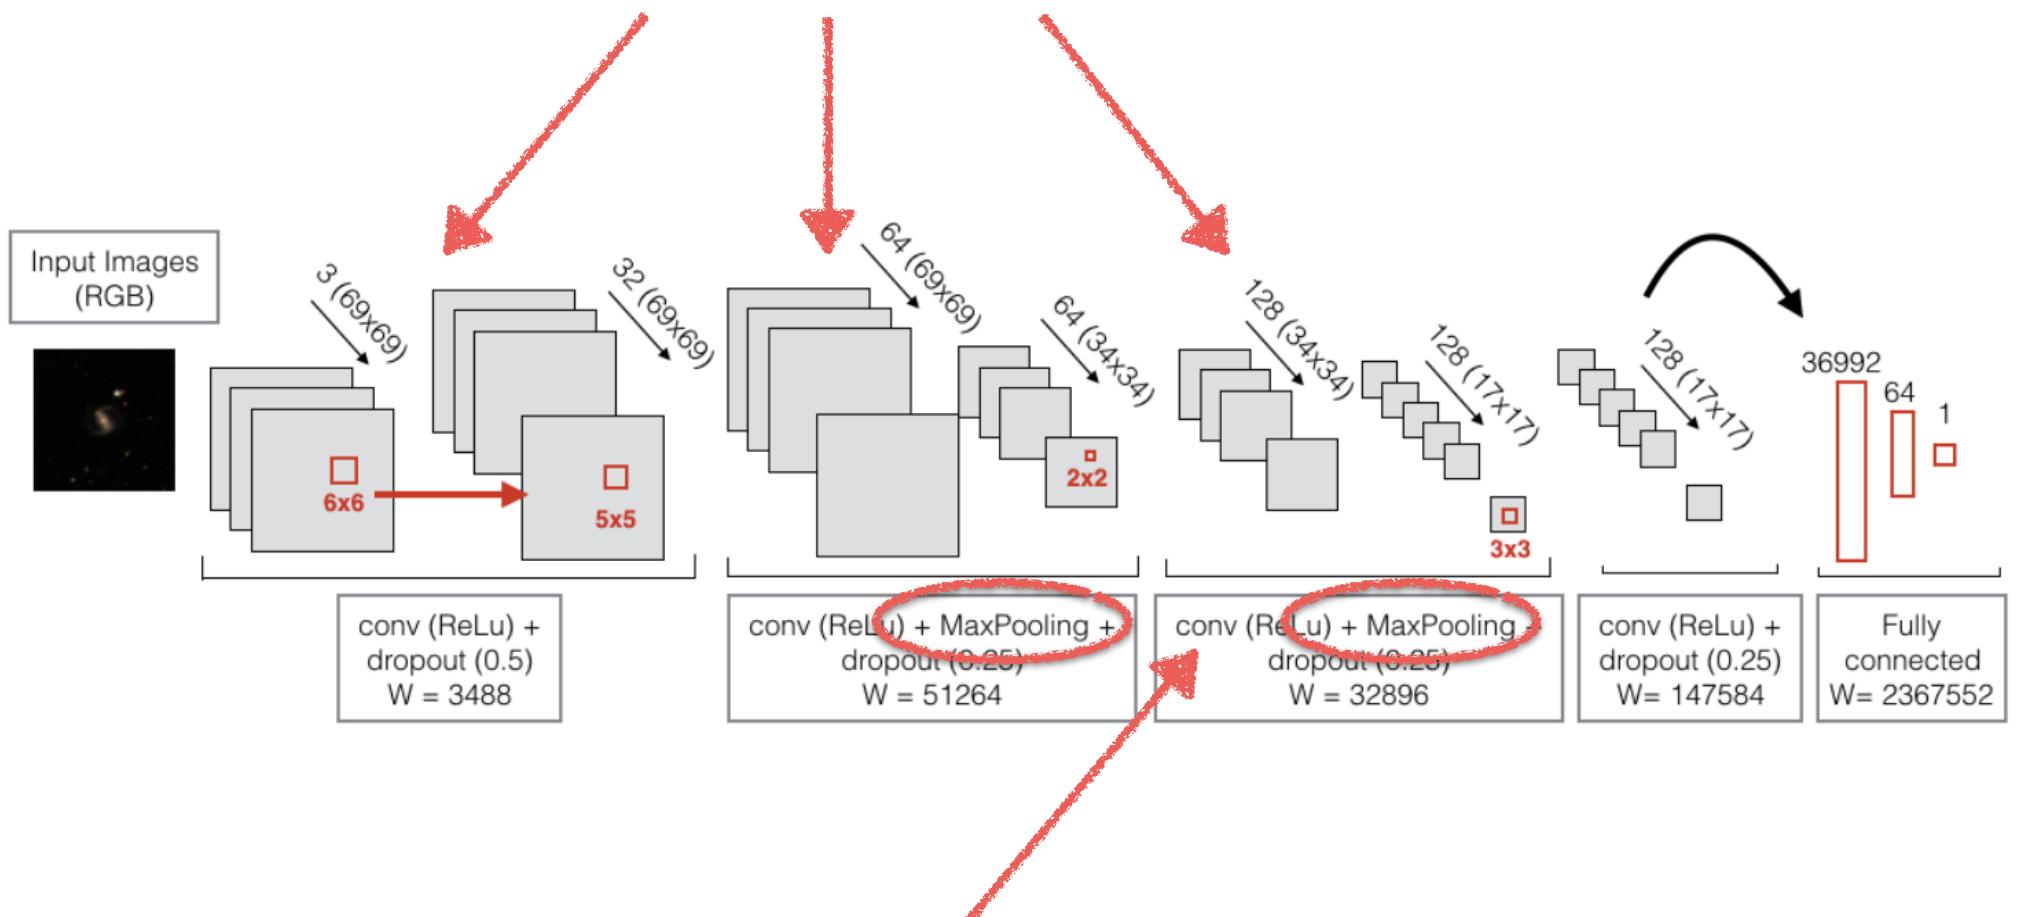


ReLU activation

Dominguez-Sanchez+18

EXAMPLE OF VERY SIMPLE CNN

3 convolutional layers

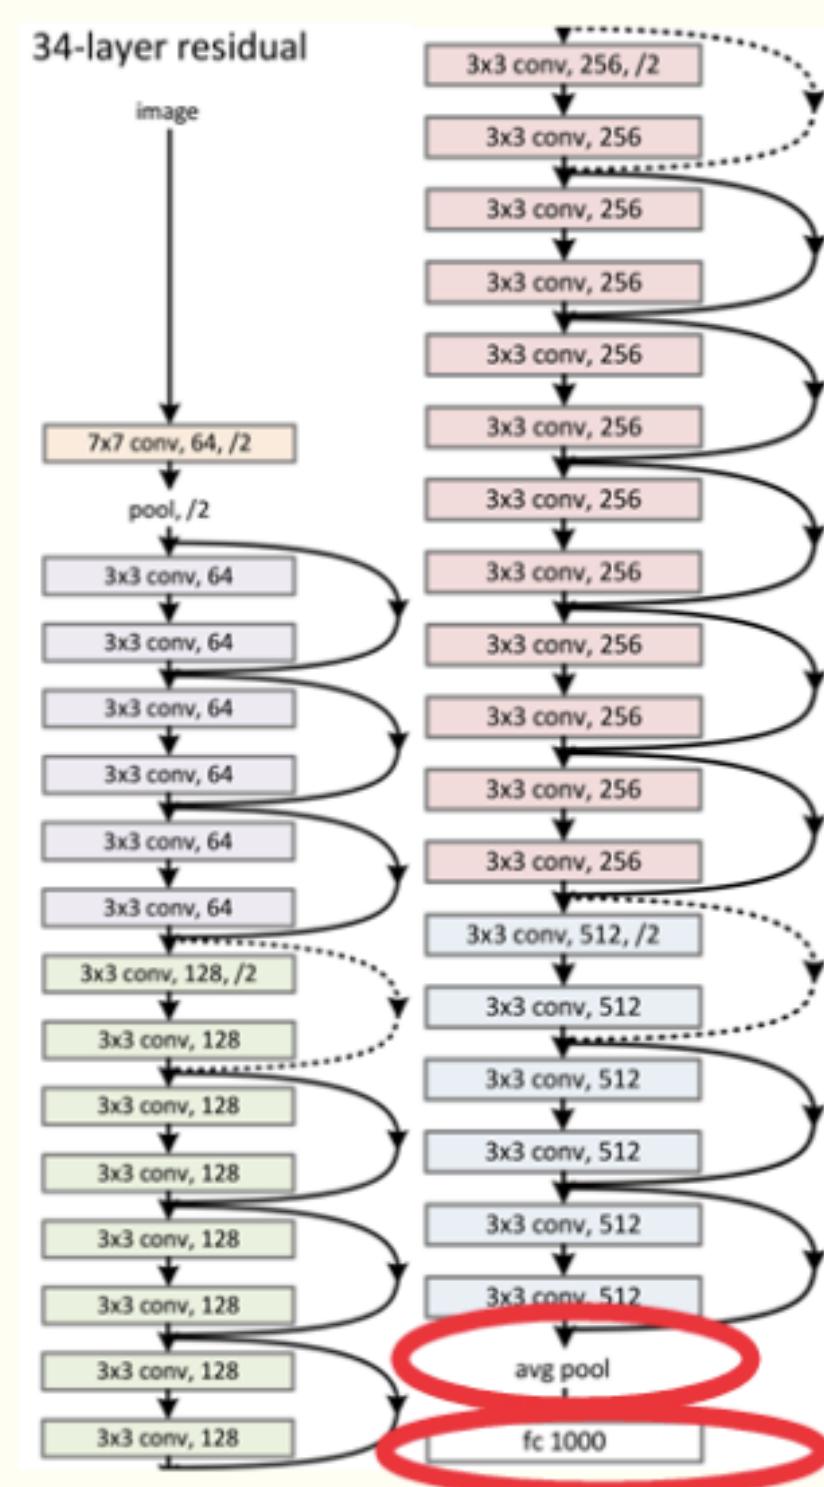


Pooling

Dominguez-Sanchez+18

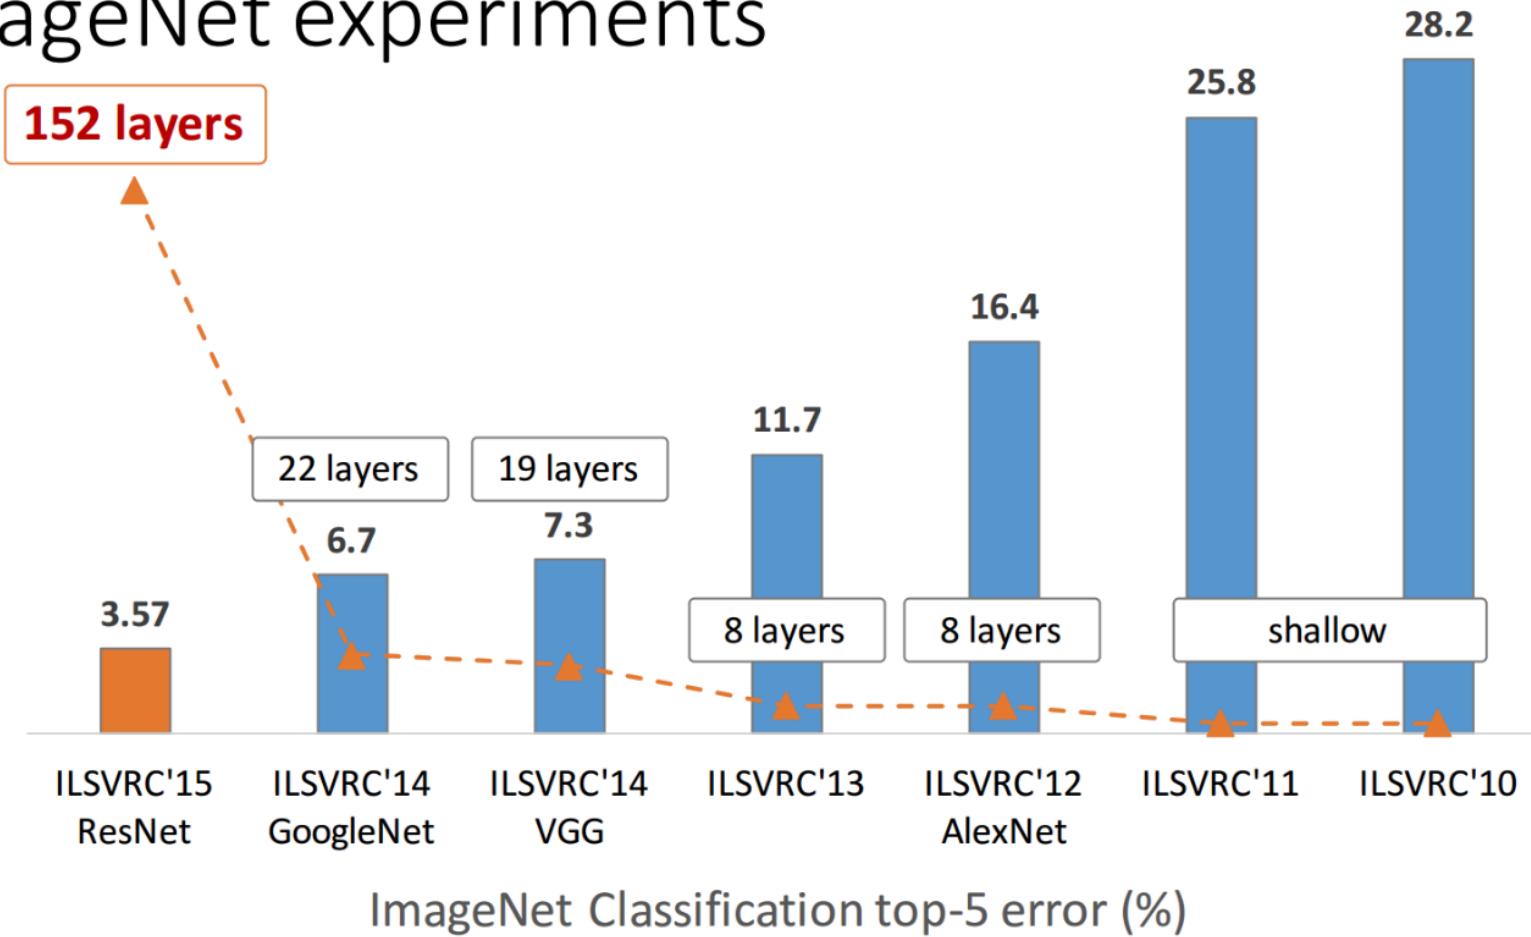
IN THE REAL LIFE..

RESNET



DEEPER TENDS TO BE BETTER...

ImageNet experiments



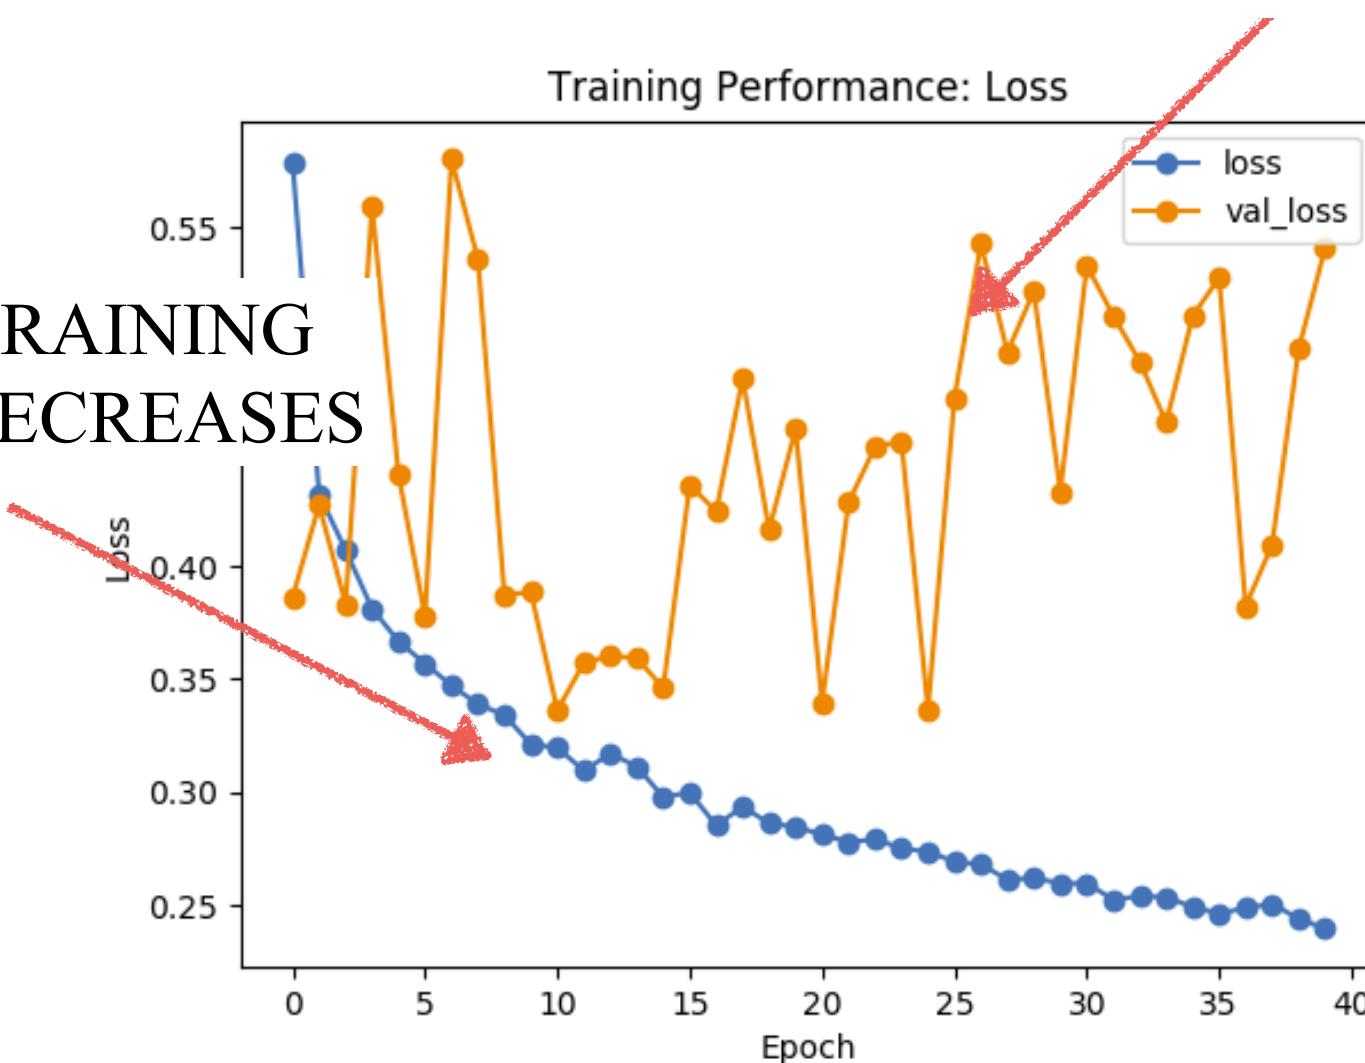
THE PROBLEMS OF GOING “TOO DEEP”

- DEEP NETWORKS ARE MORE DIFFICULT TO OPTIMIZE
- NEED MORE DATA - MORE SUBJECT TO OVERFITTING
- AND ALSO NEED MORE TIME ...

OVER-FITTING

THE TEST STAYS CONSTANT
OR INCREASES

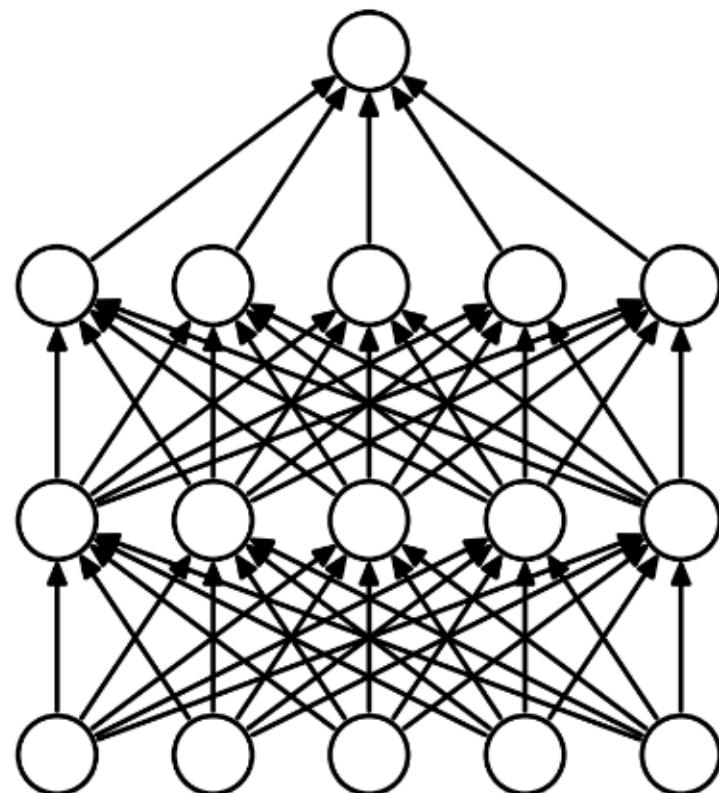
THE TRAINING
LOSS DECREASES



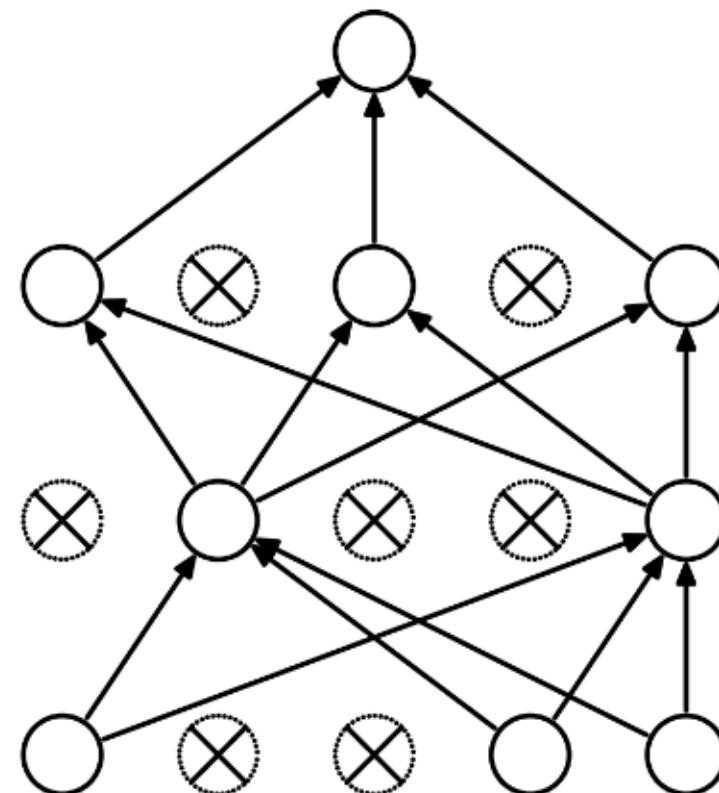
DROPOUT

[Hinton+12]

- THE IDEA IS TO REMOVE NEURONS RANDOMLY DURING THE TRAINING
- ALL NEURONS ARE PUT BACK DURING THE TEST PHASE



(a) Standard Neural Net



(b) After applying dropout.

DROPOUT

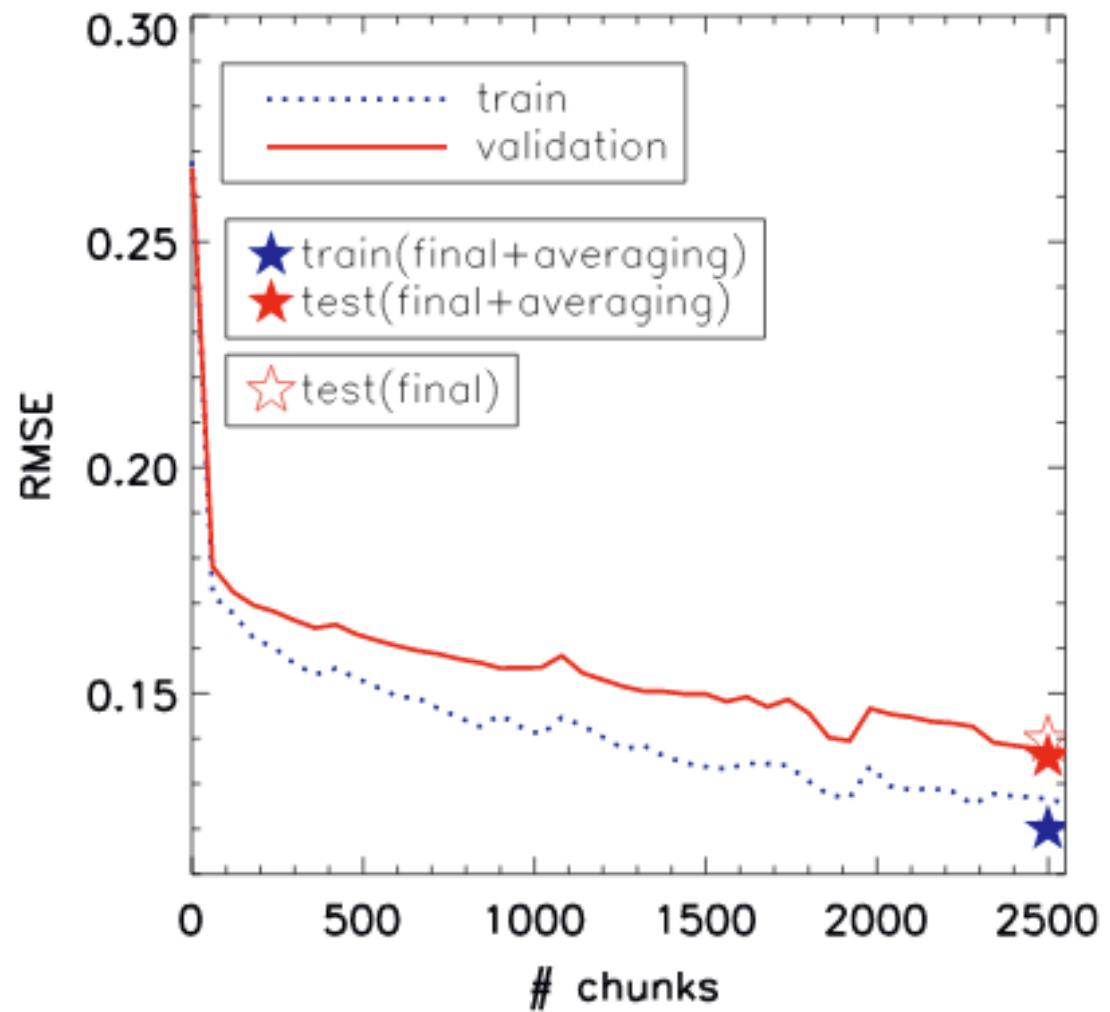
WHY DOES IT WORK?

1. SINCE NEURONS ARE REMOVED RANDOMLY, IT AVOIDS CO-ADAPTATION AMONG THEMSELVES

2. DIFFERENT SETS OF NEURONS WHICH ARE SWITCHED OFF, REPRESENT A DIFFERENT ARCHITECTURE AND ALL THESE DIFFERENT ARCHITECTURES ARE TRAINED IN PARALLEL. FOR N NEURONS ATTACHED TO DROPOUT, THE NUMBER OF SUBSET ARCHITECTURES FORMED IS 2^N . SO IT AMOUNTS TO PREDICTION BEING AVERAGED OVER THESE ENSEMBLES OF MODELS.

DROPOUT

WITH A LITTLE BIT
OF DROPOUT



BATCH NORMALIZATION

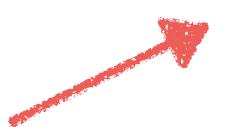
[SZEGEDY+15]

ANOTHER SOLUTION TO KEEP REASONABLE VALUES OF
THE ACTIVATIONS IN DEEP NETWORKS

BATCH NORMALIZATION PREVENTS LOW OR LARGE
VALUES BY RE-NORMALIZING THE VALUES BEFORE
ACTIVATION FOR EVERY BATCH

$$\hat{y}_i = \gamma \frac{y_i - E(y_i)}{\sigma(y_i)} + \beta$$

INPUT 

NORMALIZED INPUT 

SCATTER 

BATCH NORMALIZATION

[SZEGEDY+15]

BATCH NORMALIZATION SPEEDS UP AND STABILIZES TRAINING

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

BATCH NORMALIZATION

[SZEGEDY+15]

IN KERAS, IT IS IMPLEMENTED AS AN ADDITIONAL LAYER

BatchNormalization

[\[source\]](#)

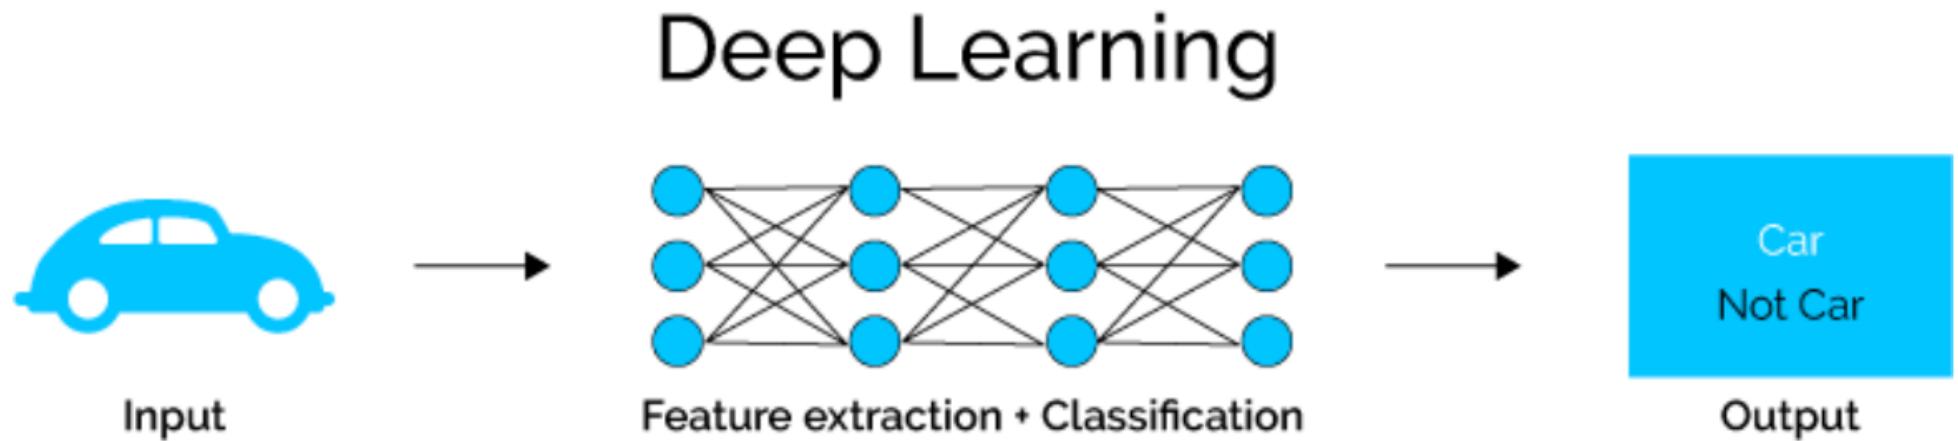
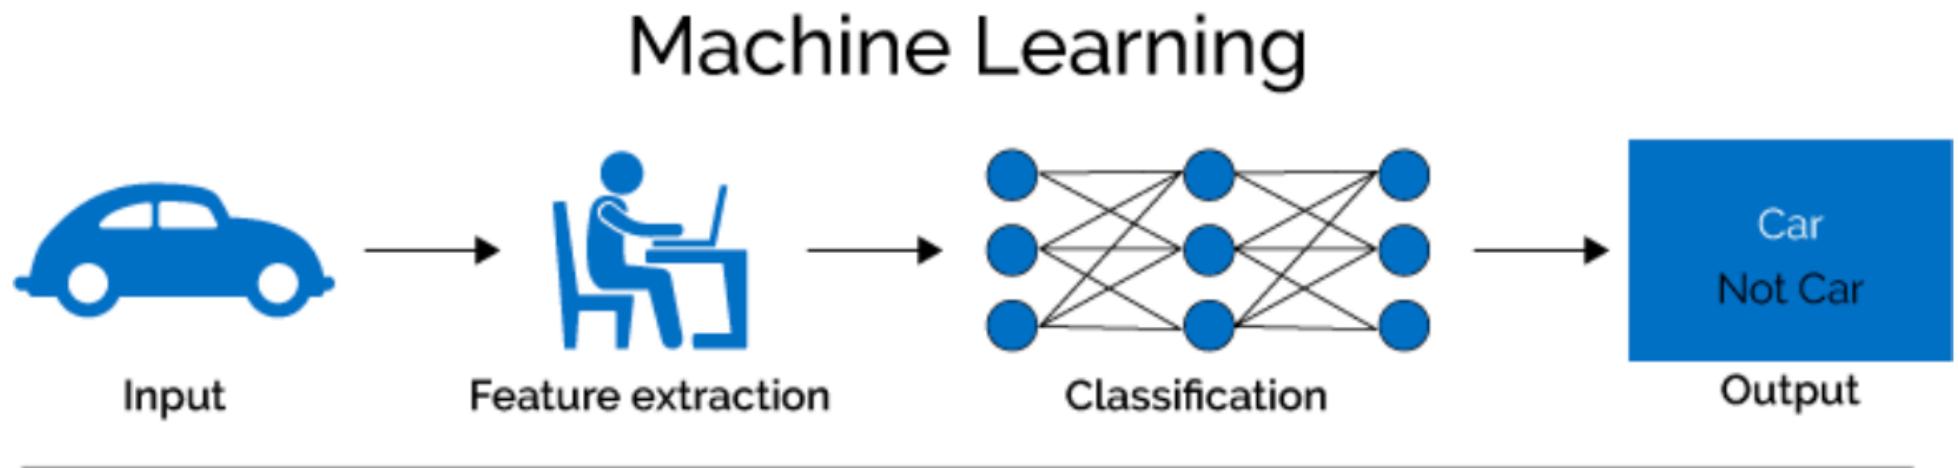
```
keras.layers.BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001, center=True, scale=True, beta_initializer
```

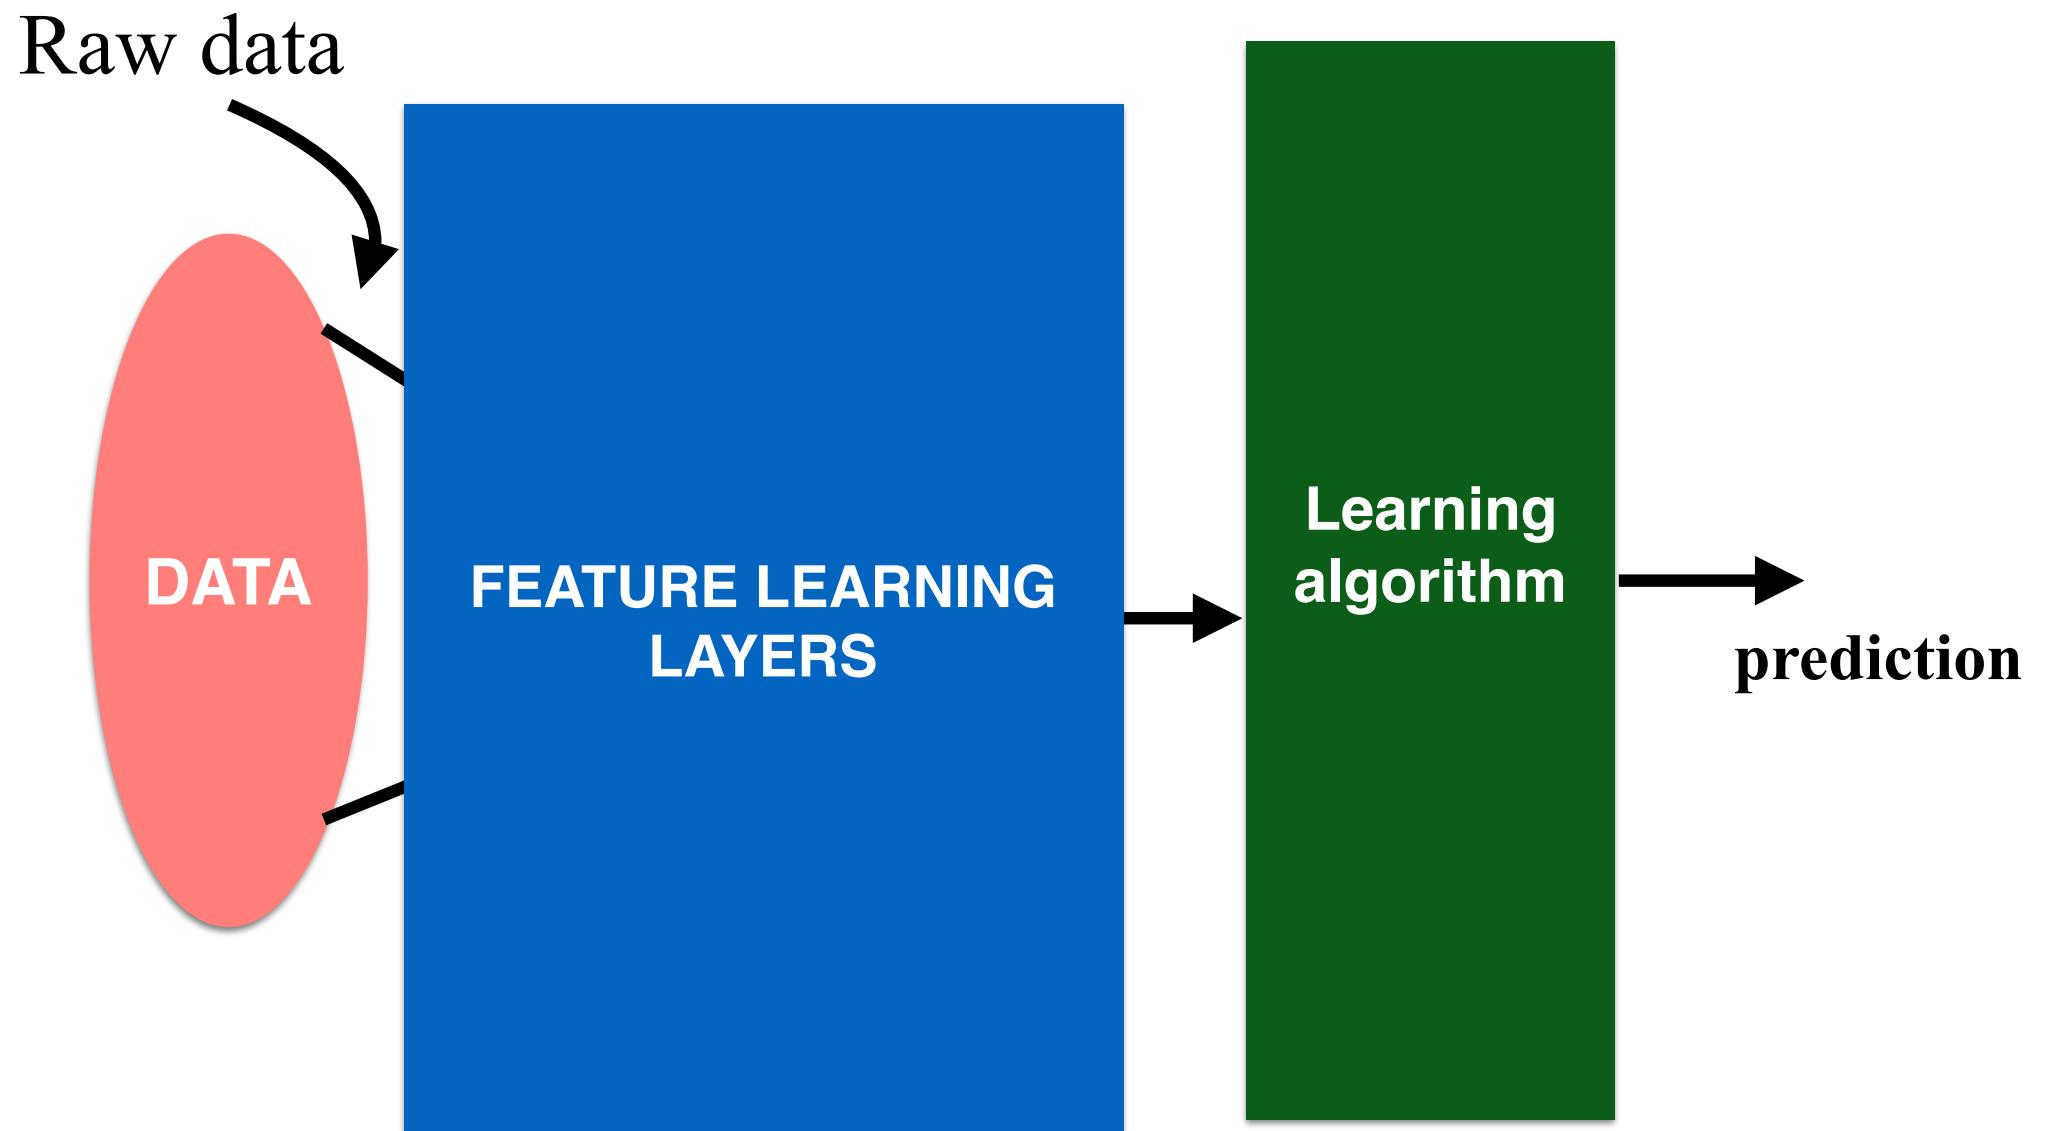
Batch normalization layer (Ioffe and Szegedy, 2014).

Normalize the activations of the previous layer at each batch, i.e. applies a transformation that maintains the mean activation close to 0 and the activation standard deviation close to 1.

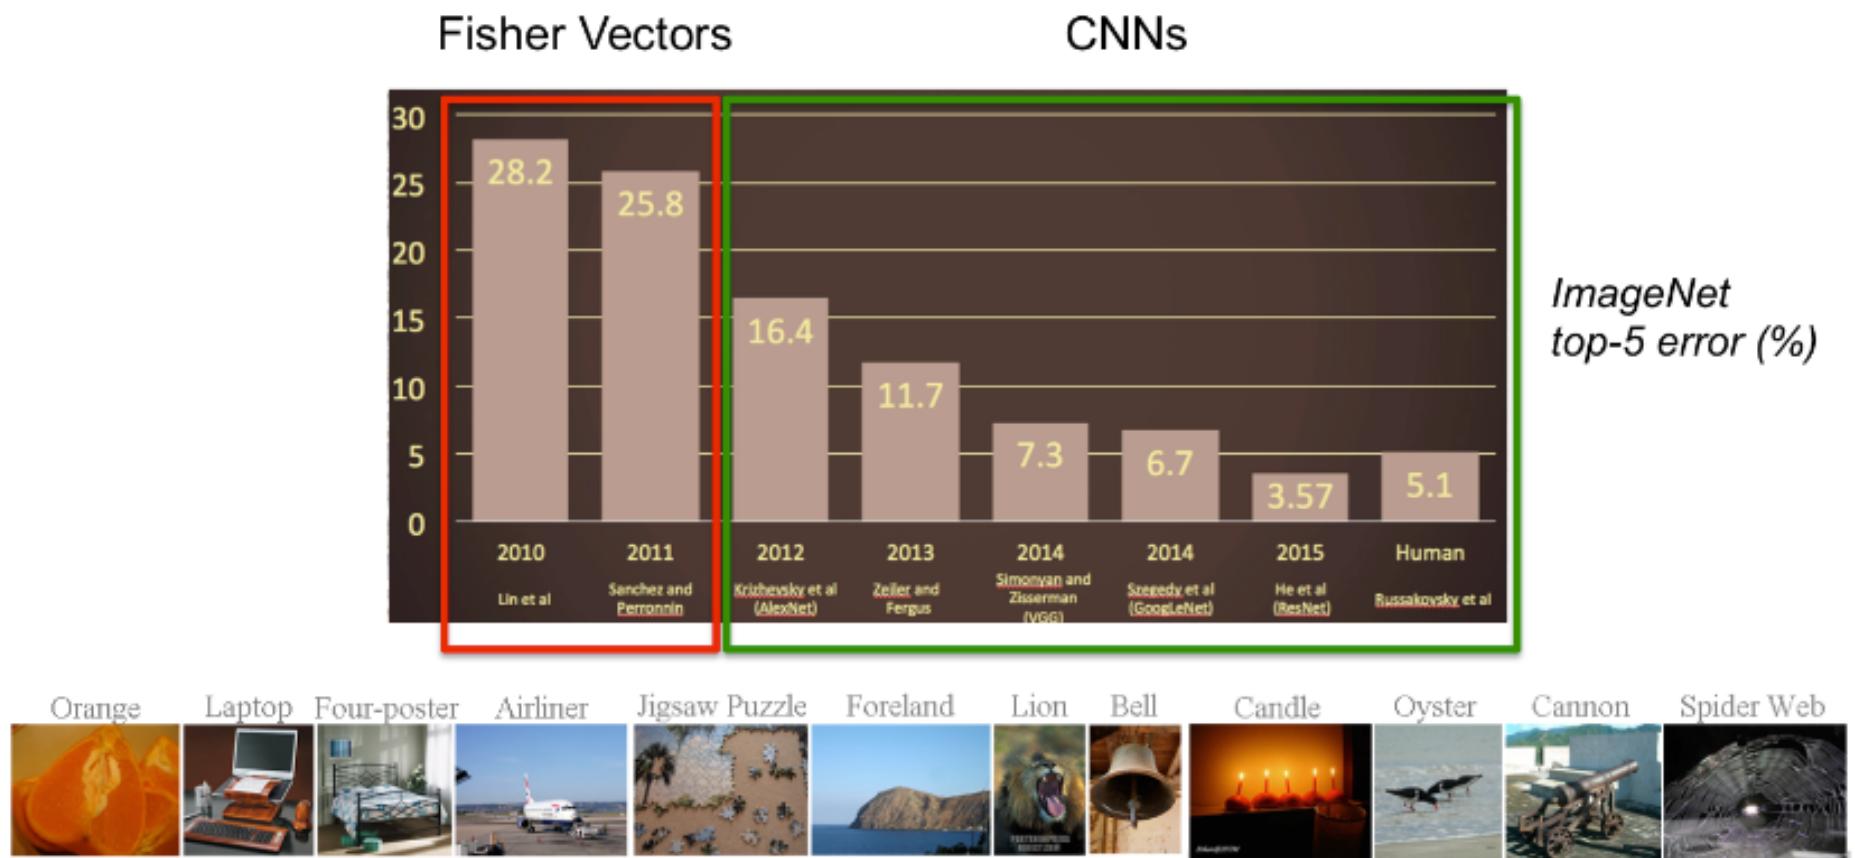
Arguments

THIS IS A CHANGE OF PARADIGM!





THIS IS A CHANGE OF PARADIGM!



ALSO FOR GALAXY MORPHOLOGY

SVMs

[HUERTAS-COMPANY+14]

AUTOMATIC

Late-Type

13

Early-Type

87

75

25

Early-Type



Late-Type



CNNs

[HUERTAS-COMPANY+15b]

AUTOMATIC

SPHEROID IRR DISK PS Unc

0.2 0.0 0.3 0.4 0.8

PS

IRR

DISK

SPHEROID

Unc

0.5 0.0 0.1 0.2 0.3 0.4 0.8

IRR

DISK

PS

SPHEROID

Unc

0.2 0.0 0.1 0.2 0.3 0.4 0.8

PS

SPHEROID

Unc

0.3 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.4 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.5 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.6 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.7 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.8 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

0.9 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

1.0 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

1.1 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

1.2 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

Unc

1.3 0.0 0.1 0.2 0.3 0.4 0.8

SPHEROID

SPHEROID

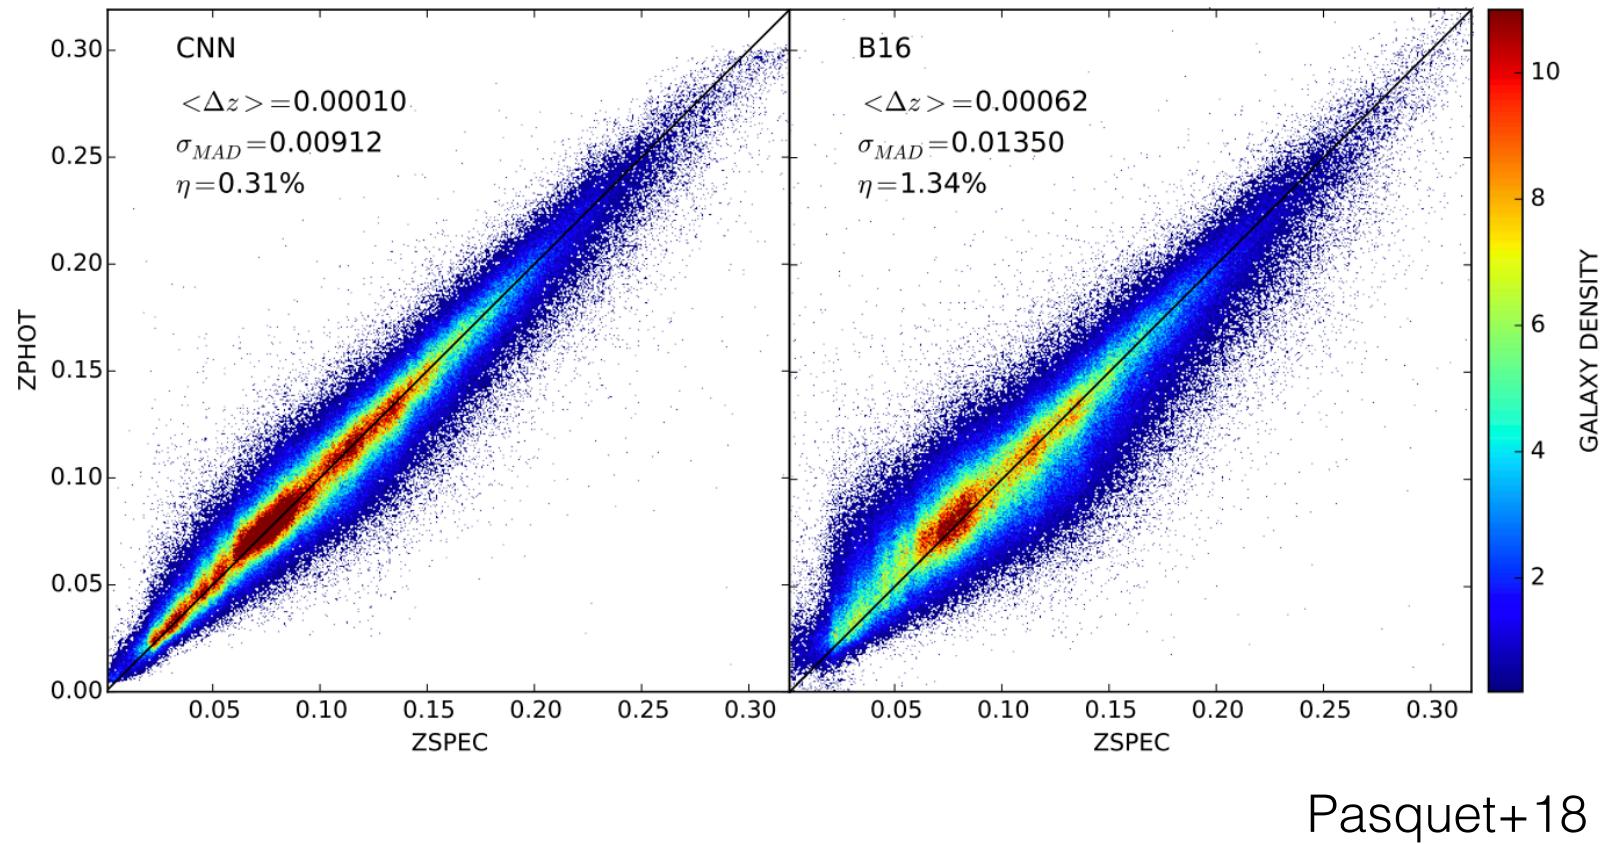
Unc

VISUAL



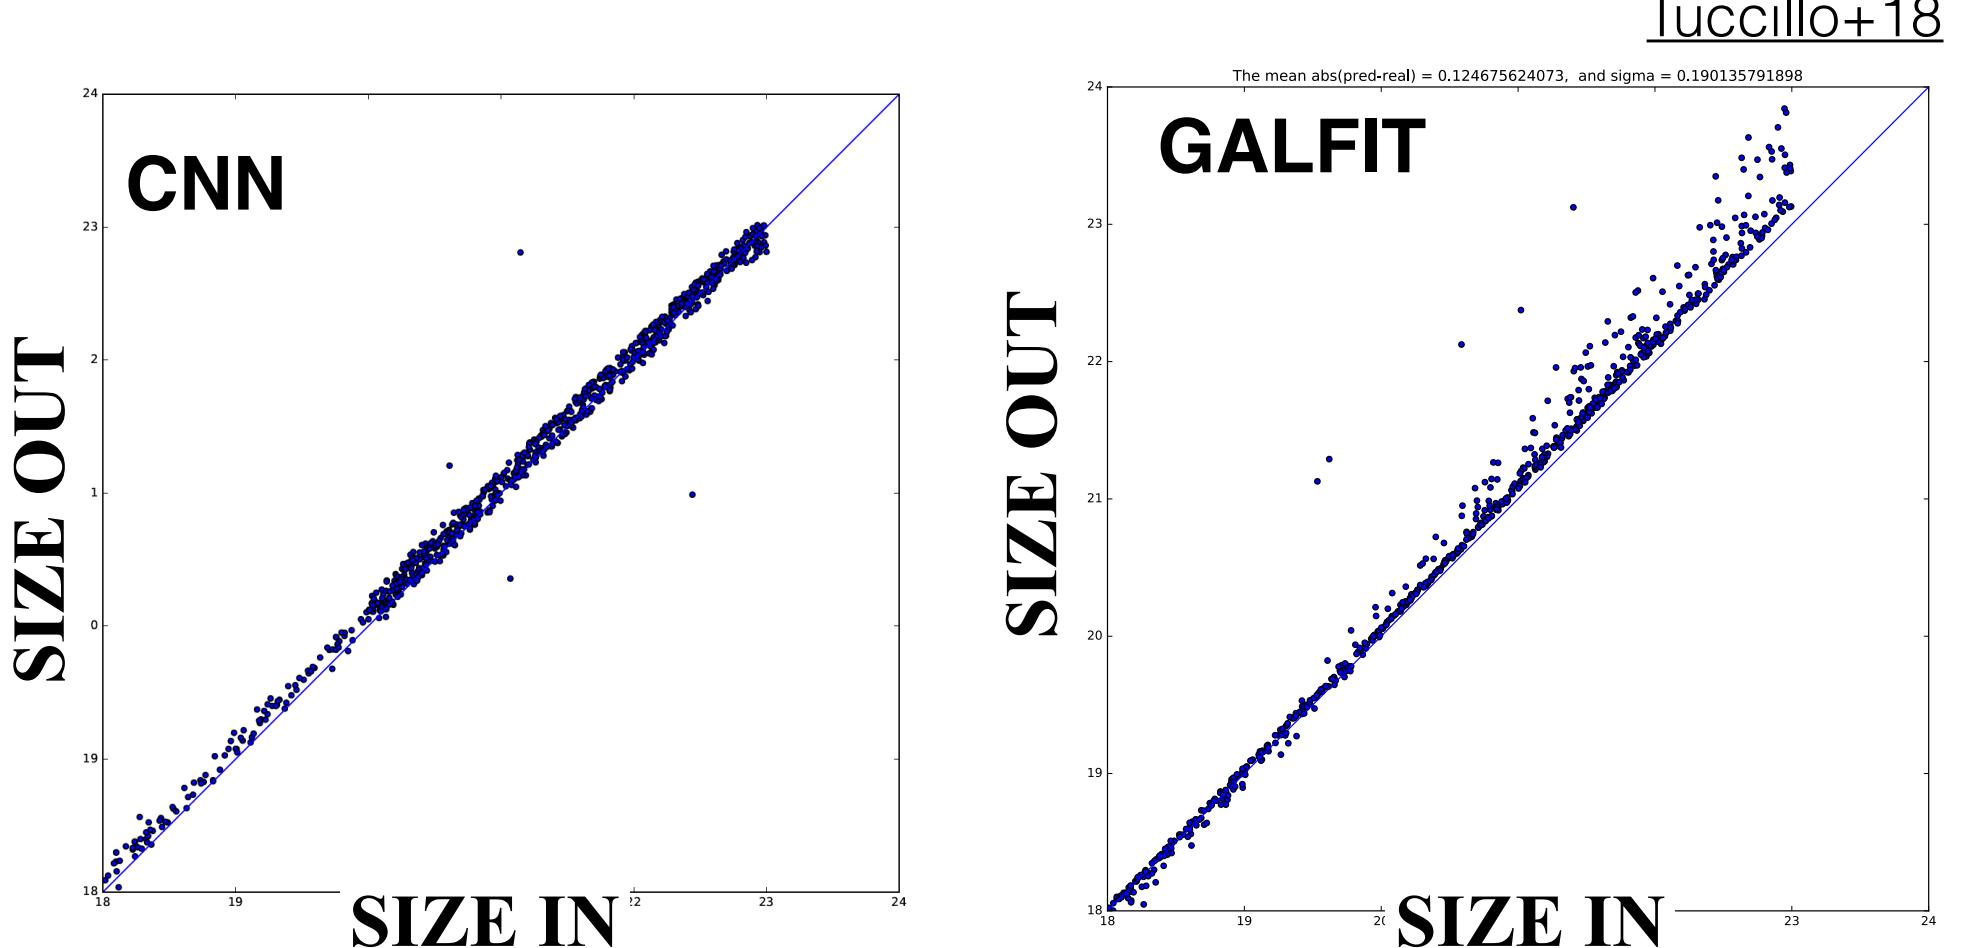
VISUAL

PHOTOMETRIC REDSHIFTS



AUTOMATICALLY COMBINING MORPHOLOGY AND COLOR
FOR PHOTOZ ESTIMATION

ESTIMATING SIZES OF GALAXIES FROM IMAGES (MODEL FITTING)



VERY SIMILAR RESULTS ON THE SAME SIMULATIONS, BUT
CNNs are several orders of magnitude faster [3.5 hrs vs. <1 sec for
~1000 objects]

THE PRICE TO PAY?

1. LARGE NUMBER OF PARAMETERS IMPLIES LARGE DATASETS TO TRAIN
2. LOOSE EVEN MORE DEGREE OF CONTROL OF WHAT THE ALGORITHM IS DOING SINCE THE FEATURE EXTRACTION PROCESS BECOMES UNSUPERVISED

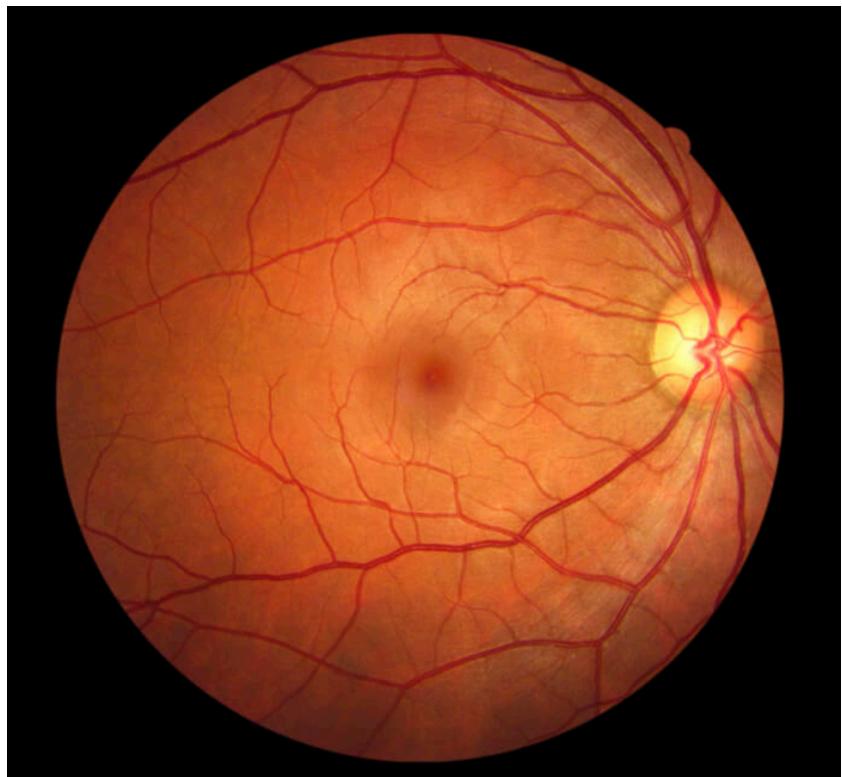
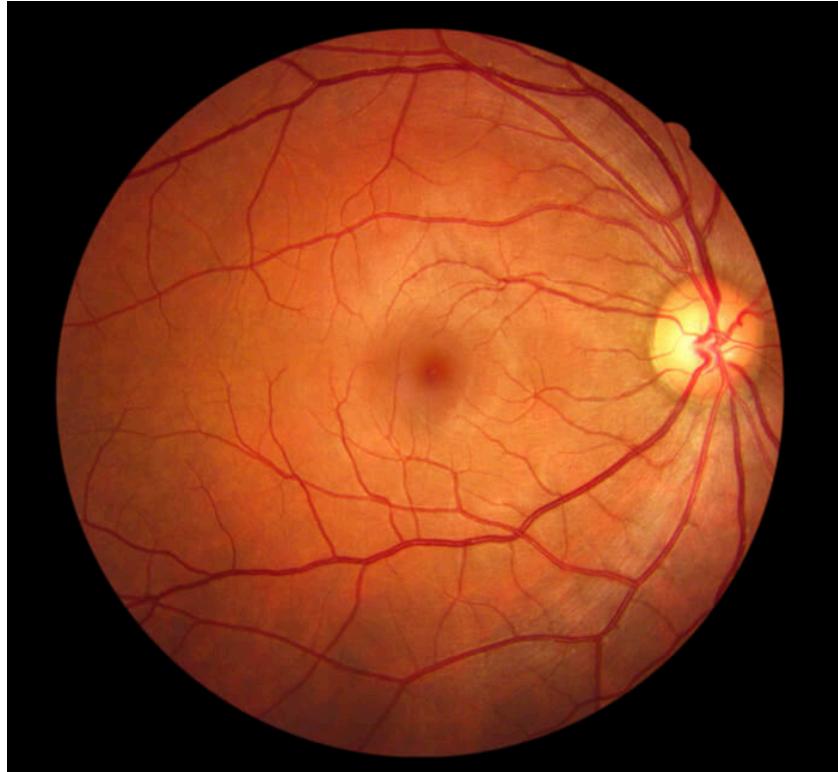


IMAGE OF THE BACK OF THE EYE





**DEEP LEARNING CAN
IDENTIFY
THE PATIENT'S
GENDER WITH 95%
ACCURACY**

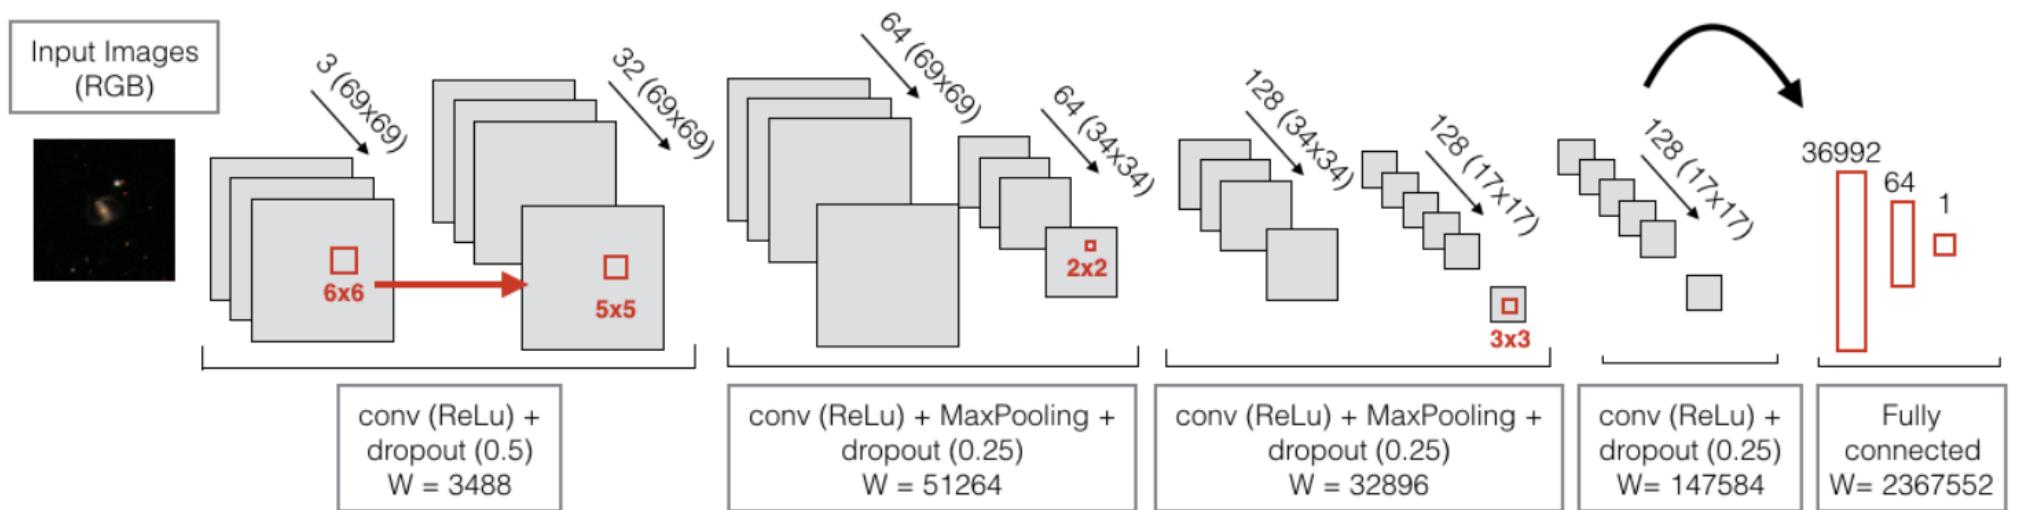
IMAGE OF THE BACK OF THE EYE



PART IV: BEYOND CLASSIFICATION:

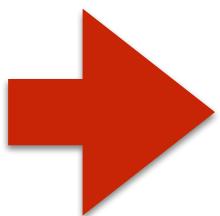
IMAGE2IMAGE NETWORKS

UP TO NOW CNNs MAP IMAGES (SIGNALS) INTO FLOATS



Dominguez-Sanchez+18

Classification has its limits

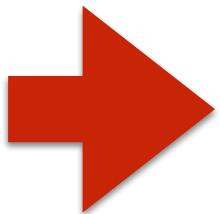


HOW DO I CLASSIFY THIS IMAGE?

Classification has its limits



classification



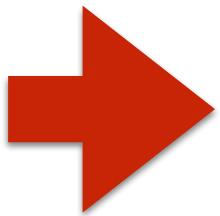
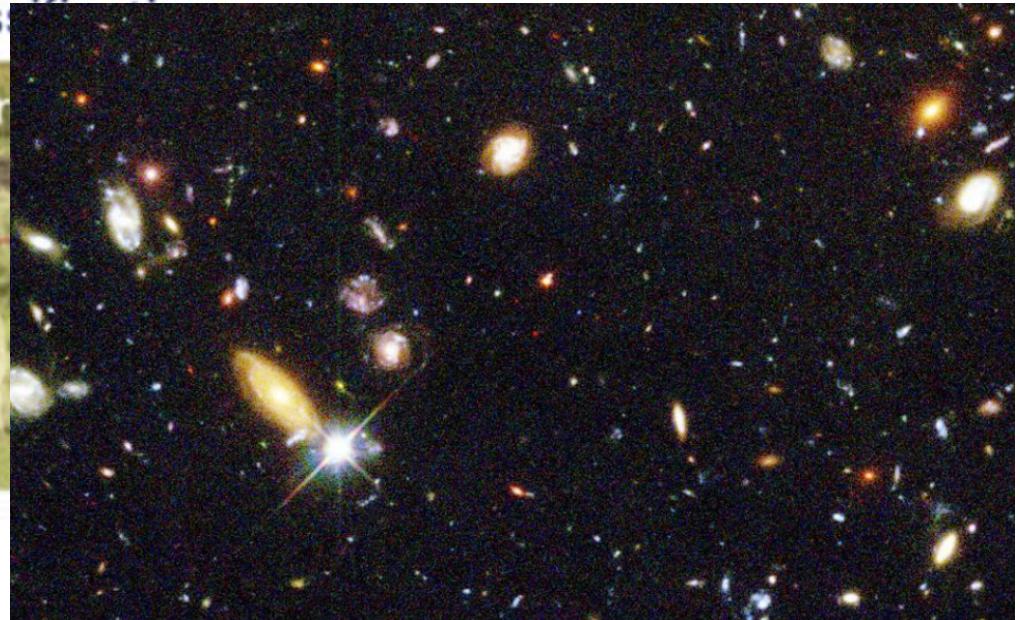
HOW DO I CLASSIFY THIS IMAGE?

Classification has its limits



clas

per



HOW DO I CLASSIFY THIS IMAGE?

Going beyond classification: increasing complexity

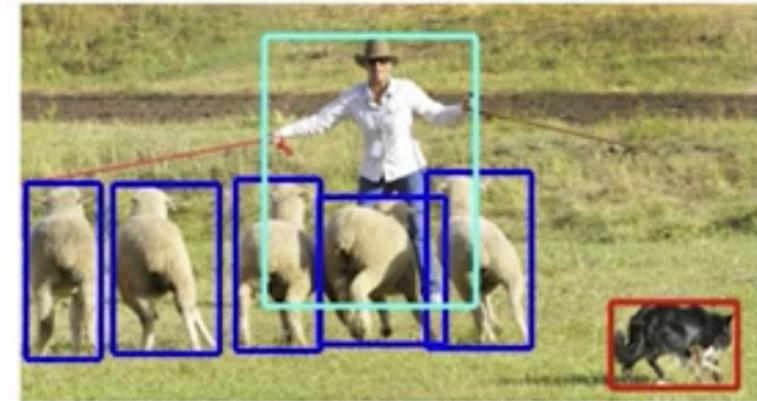
classification



semantic segmentation



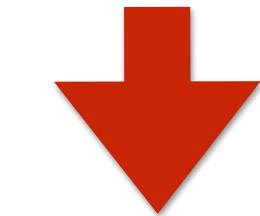
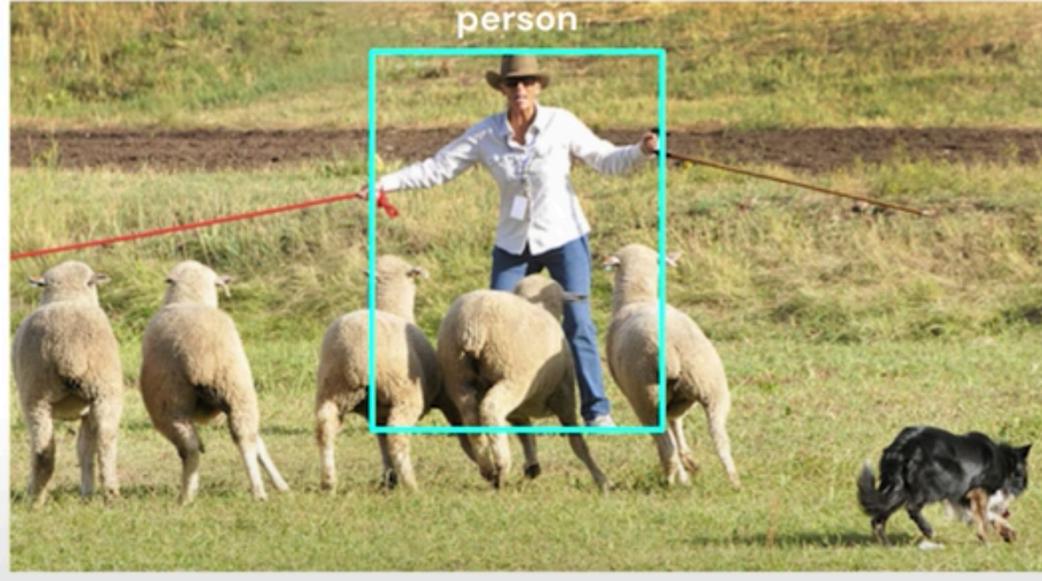
object detection



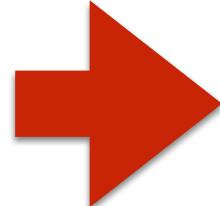
instance segmentation



LET'S GO A STEP FURTHER INTO SEMANTIC SEGMENTATION



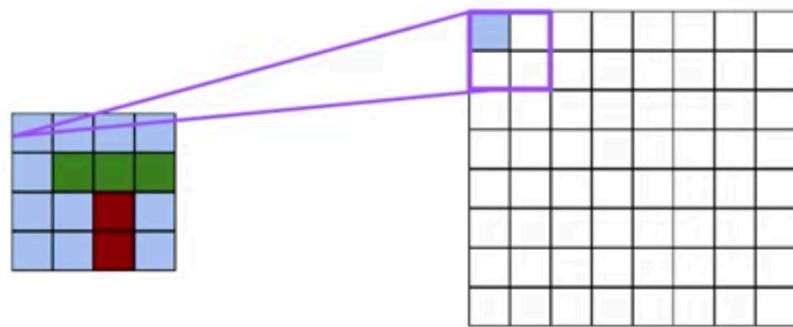
**BOUNDING BOXES
ARE NOT ALWAYS
GOOD
REPRESENTATIONS**



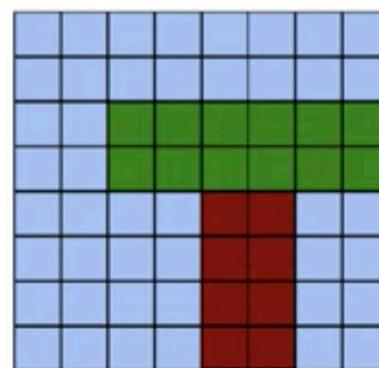
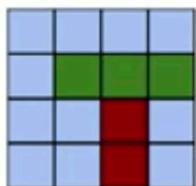
semantic segmentation



UNPOOLING OPERATION (INVERSE OF POOLING)



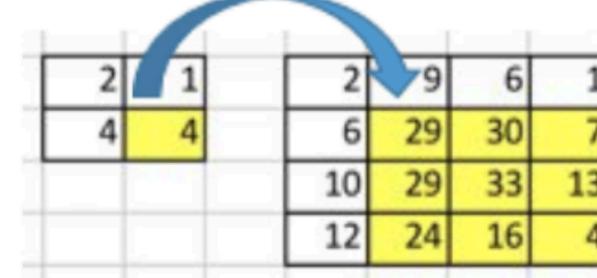
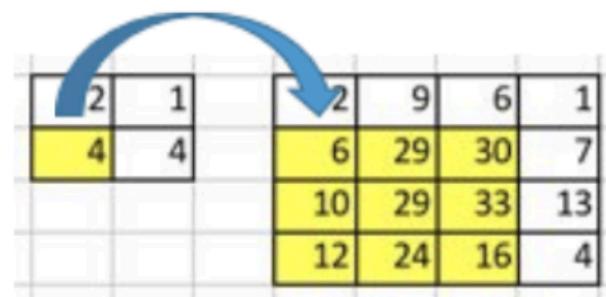
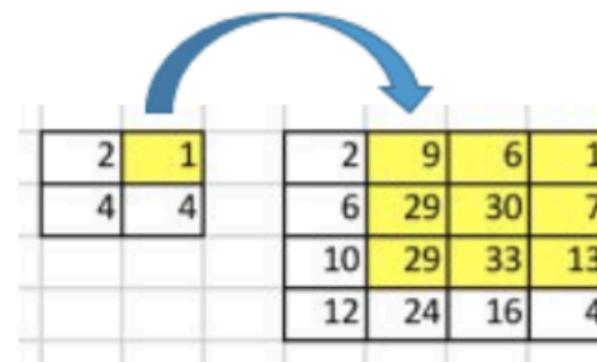
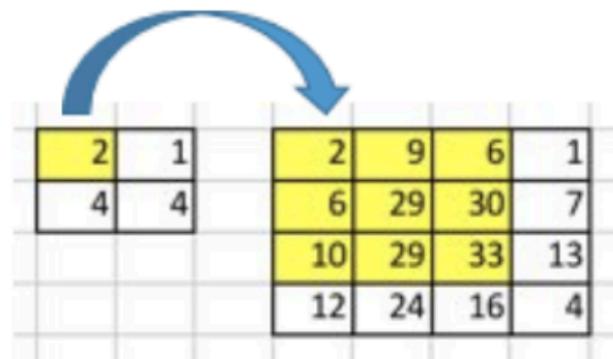
COPY PIXELS IN A
GIVEN WINDOW



GENERATES
LARGER IMAGES
FROM SMALLER
ONES

TRANSPOSED CONVOLUTION

ALLOWS TO INCREASE THE SIZE



Going Backward of Convolution

EXAMPLE TAKEN FROM HERE

CONVOLUTION MATRIX

	0	1	2
0	1	4	1
1	1	4	3
2	3	3	1

Kernel (3, 3)

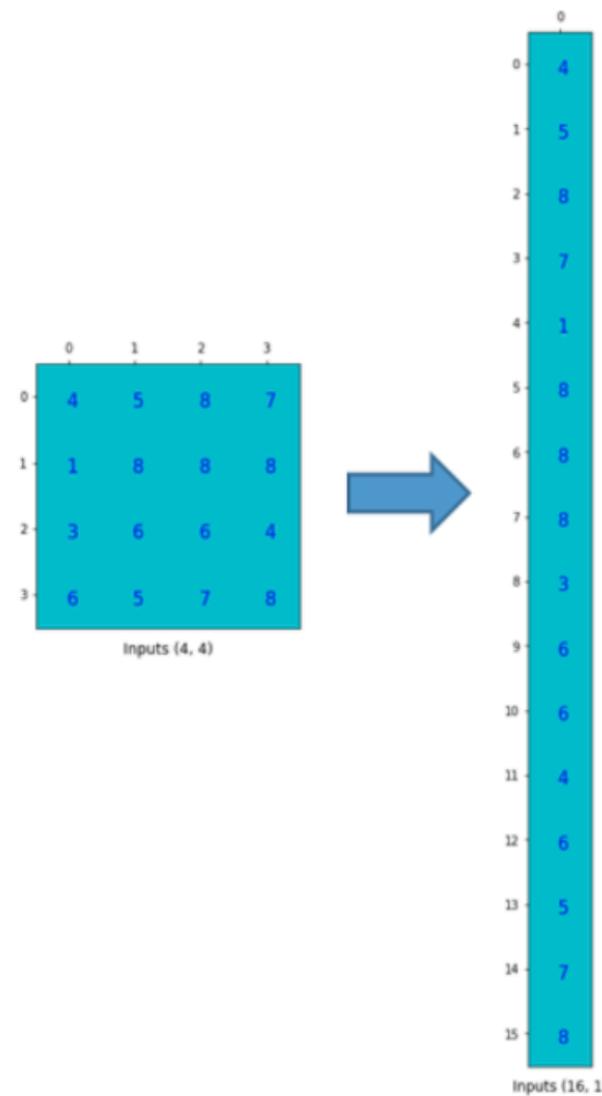
THE KERNEL CAN BE ARRANGED IN FORM OF A MATRIX:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0	0
1	0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0
2	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1	0
3	0	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1

Convolution Matrix (4, 16)

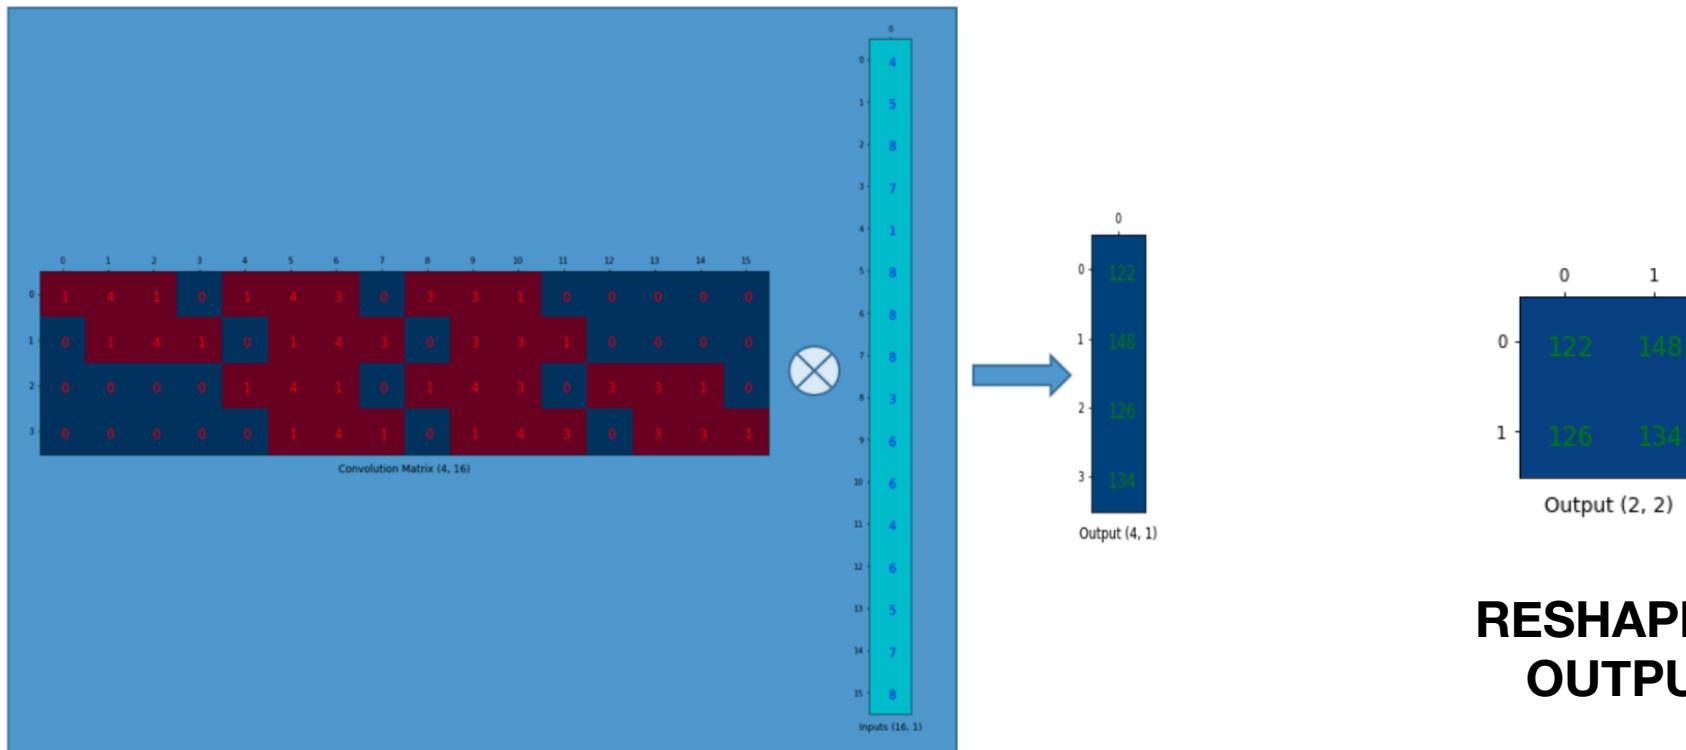
EXAMPLE TAKEN FROM HERE

THE INPUT IS FLATTENED INTO A COLUMN VECTOR



EXAMPLE TAKEN FROM HERE

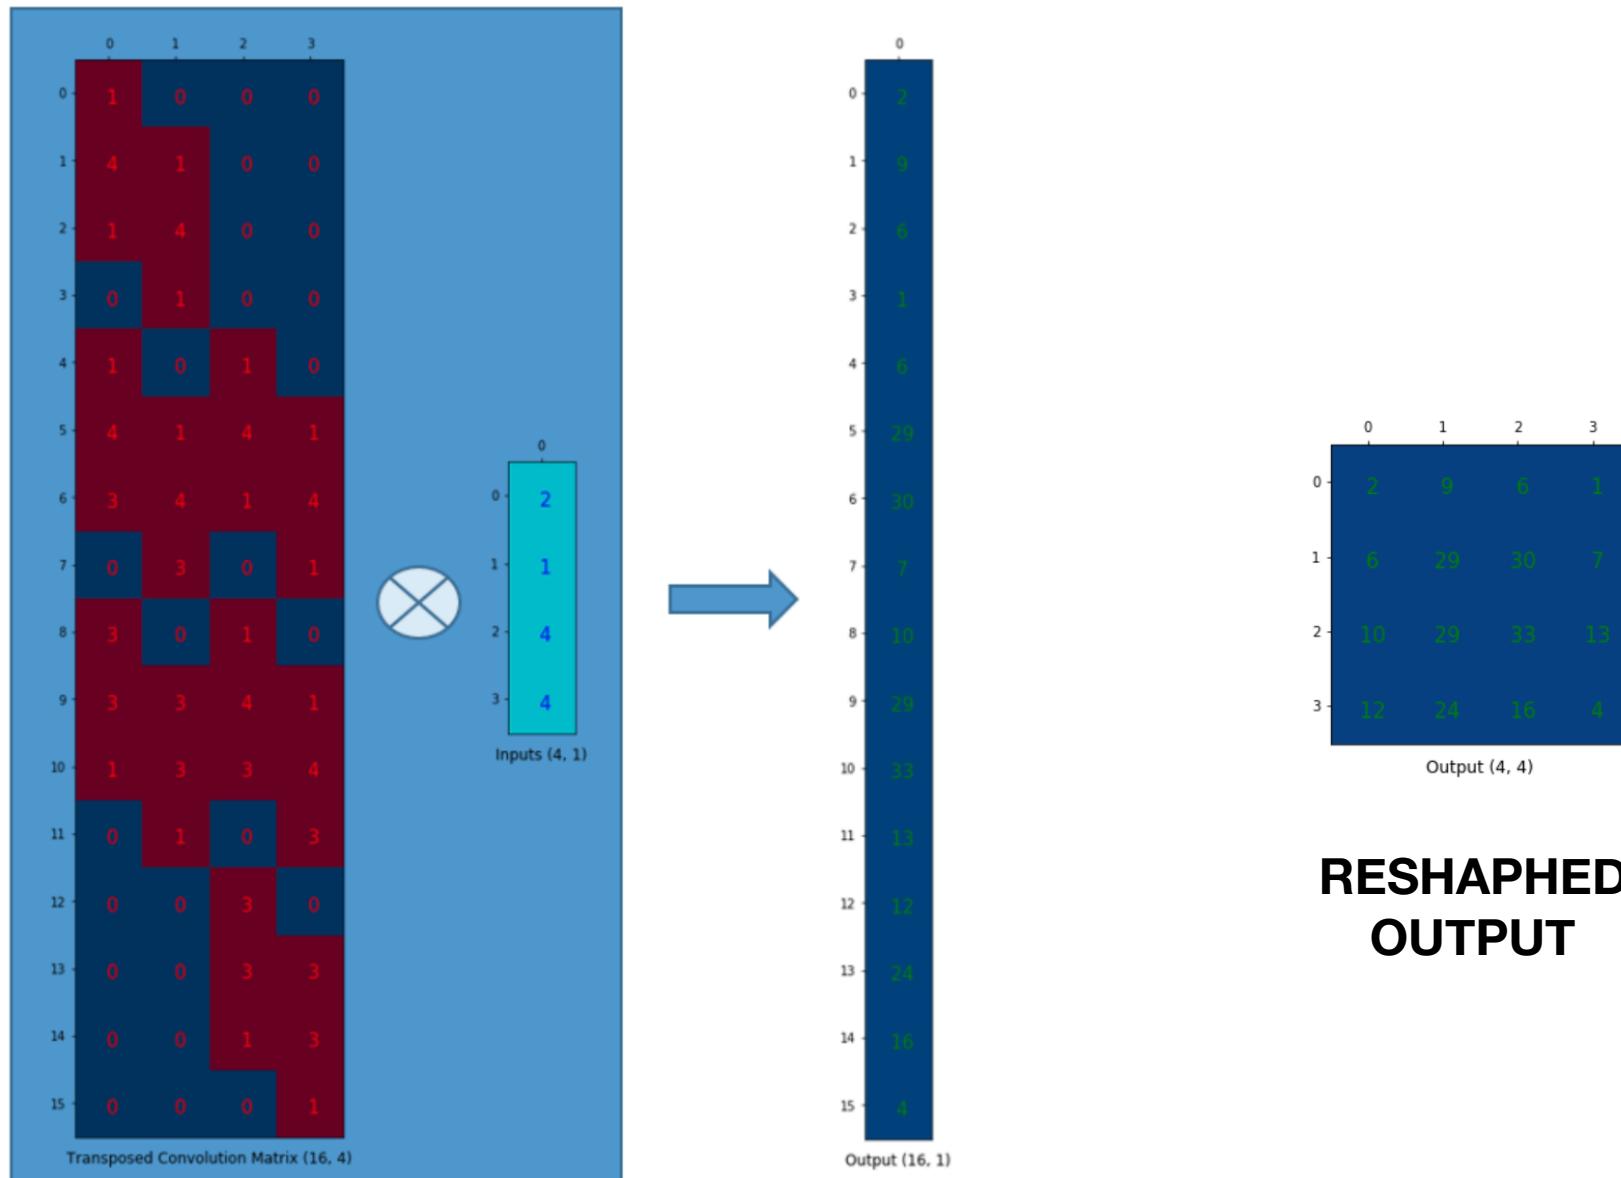
THE CONVOLUTION IS TRANSFORMED INTO A PRODUCT OF MATRICES



**RESHAPED
OUTPUT**

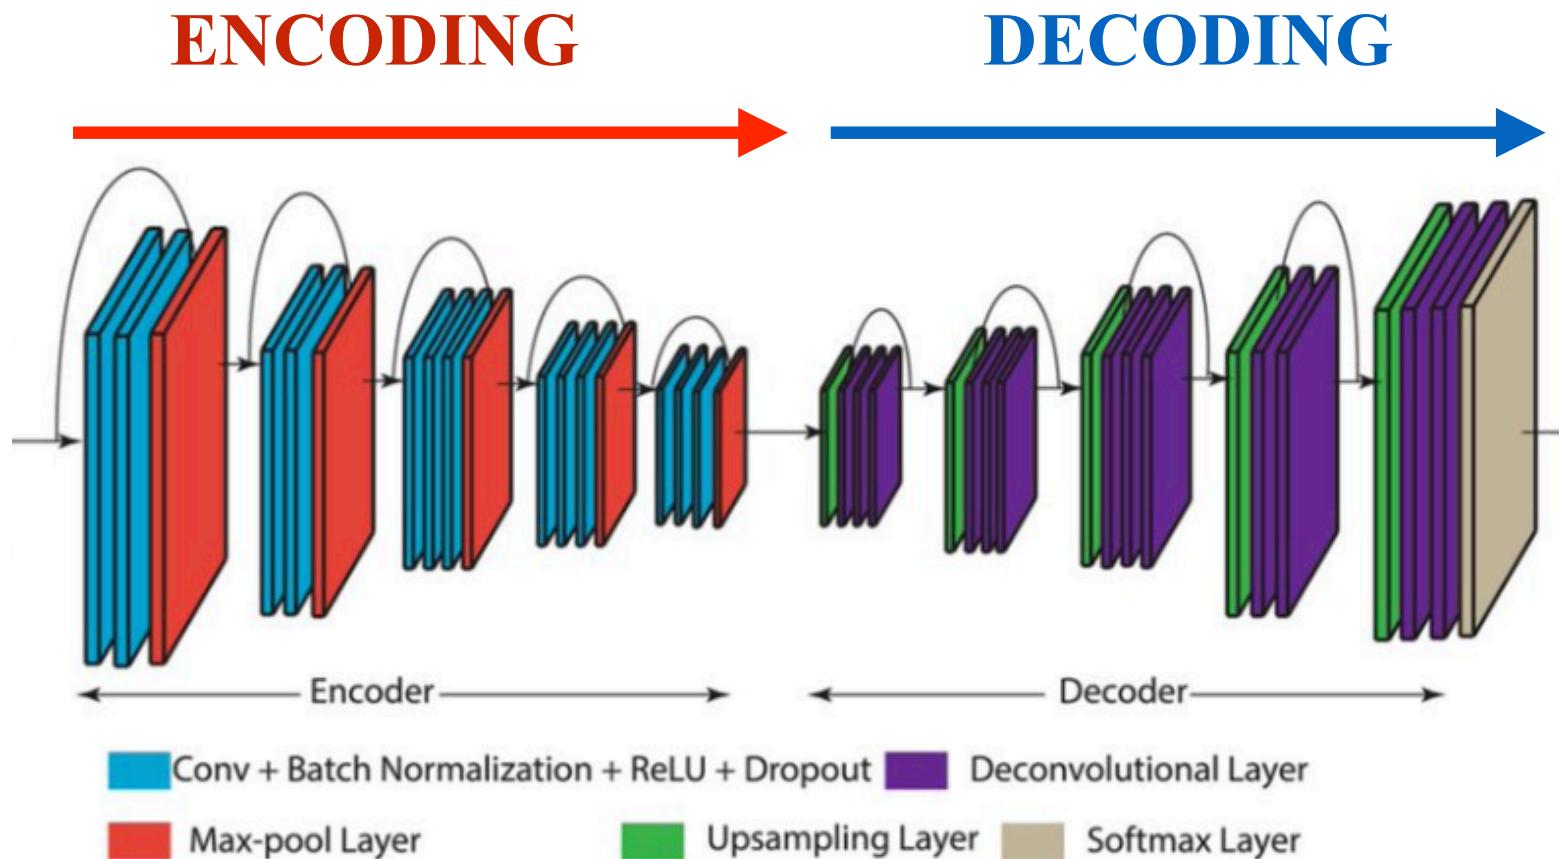
EXAMPLE TAKEN FROM HERE

THE TRANSPOSED CONVOLUTION IS THE INVERSE OPERATION



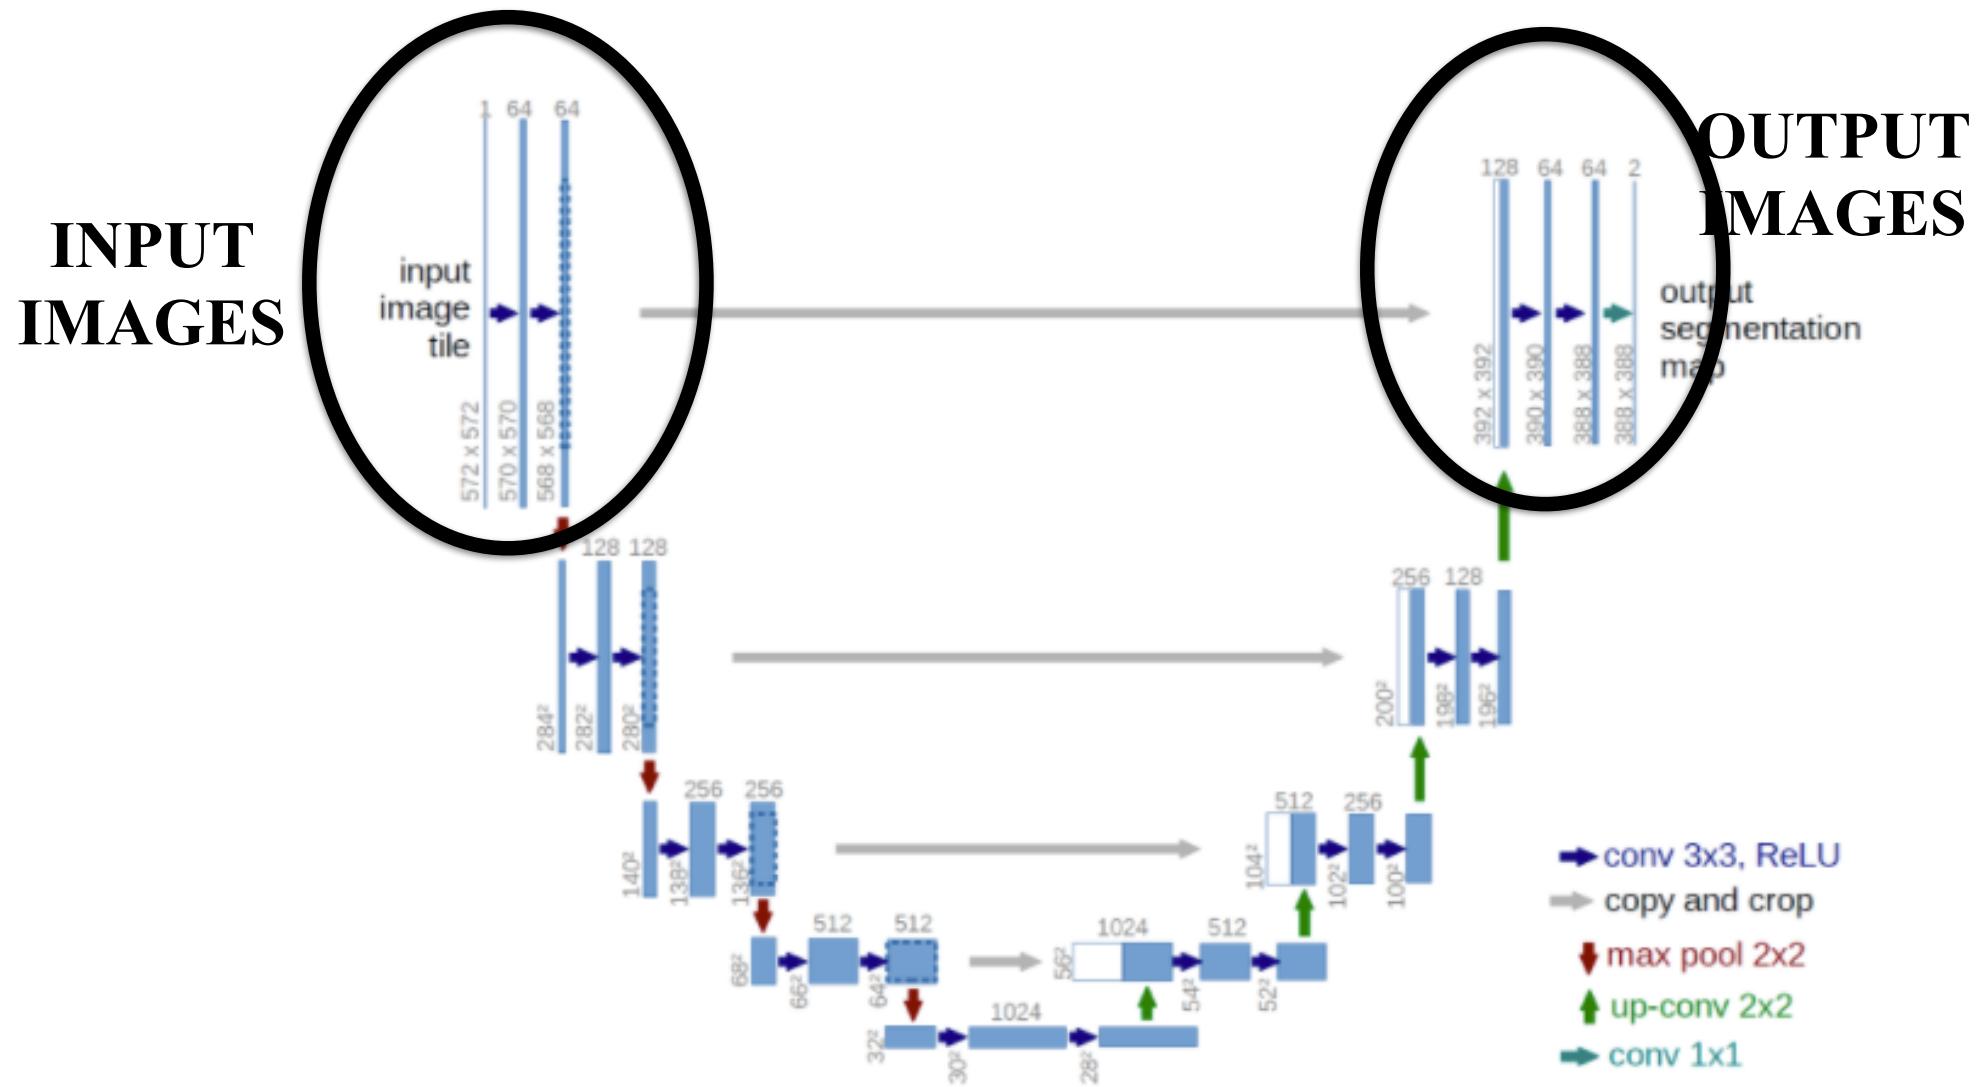
EXAMPLE TAKEN FROM HERE

ENCODER-DECODERS GO FROM IMAGE 2 IMAGE

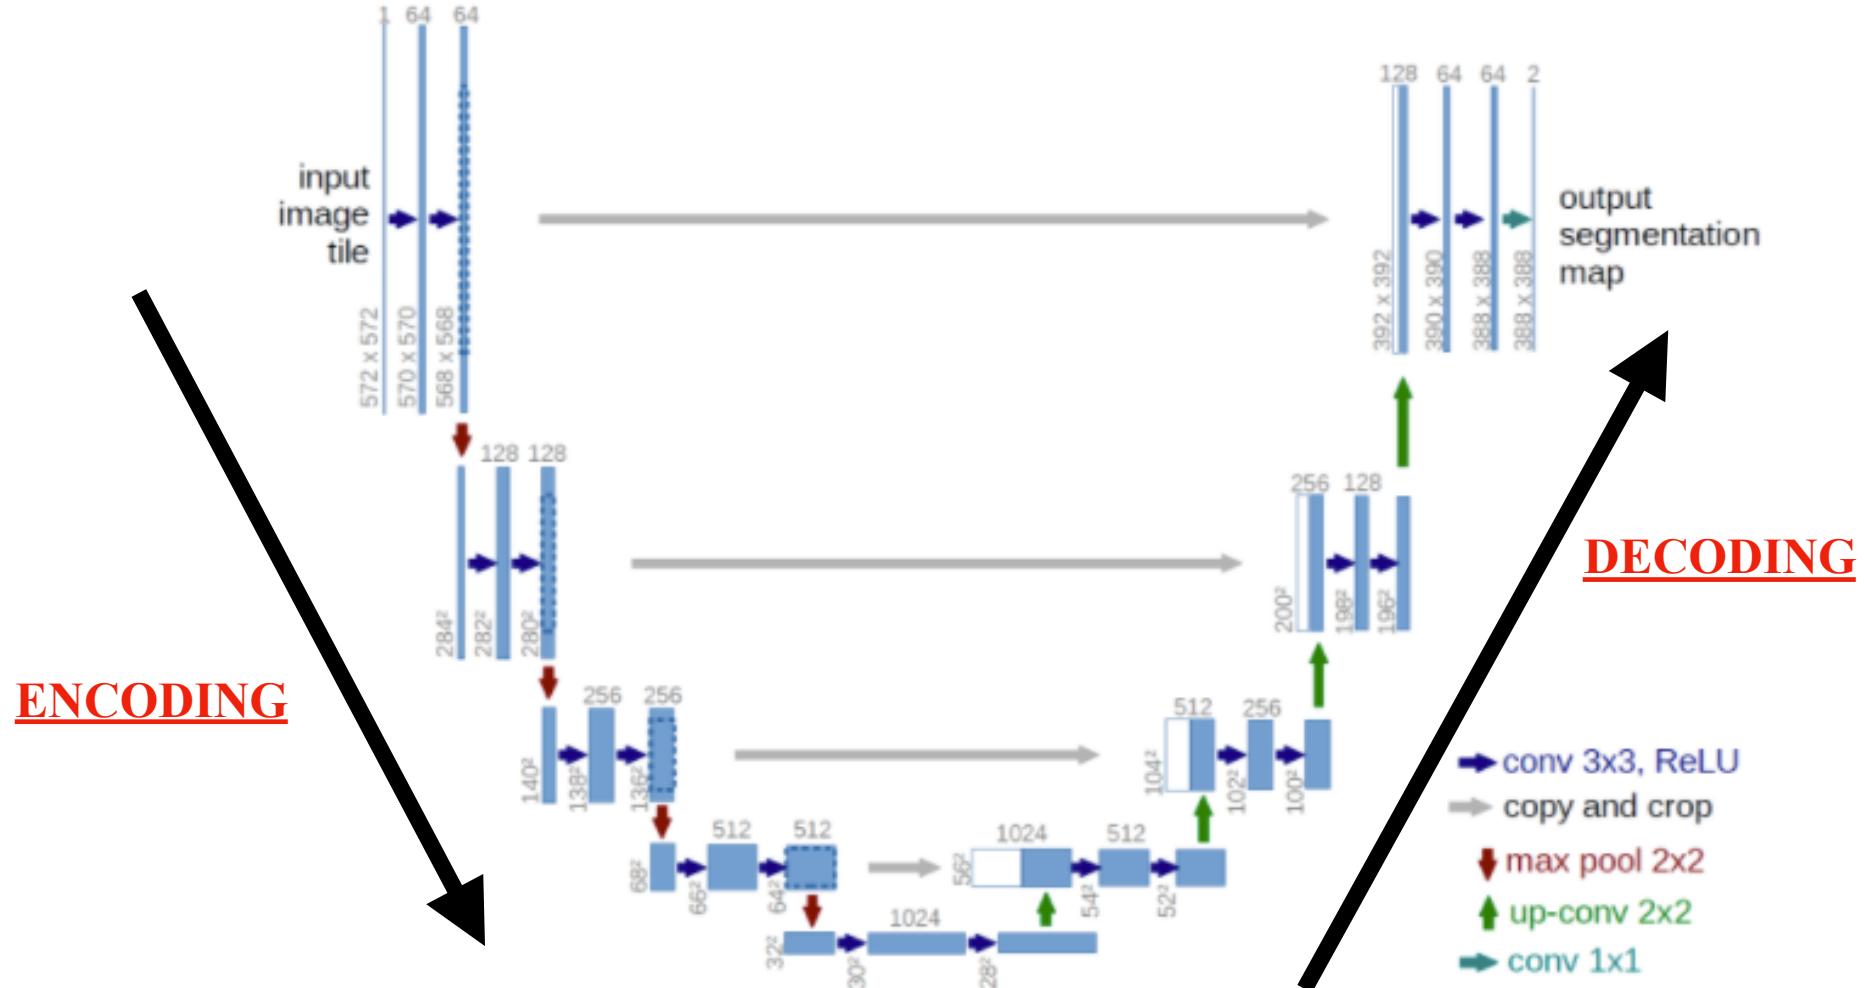


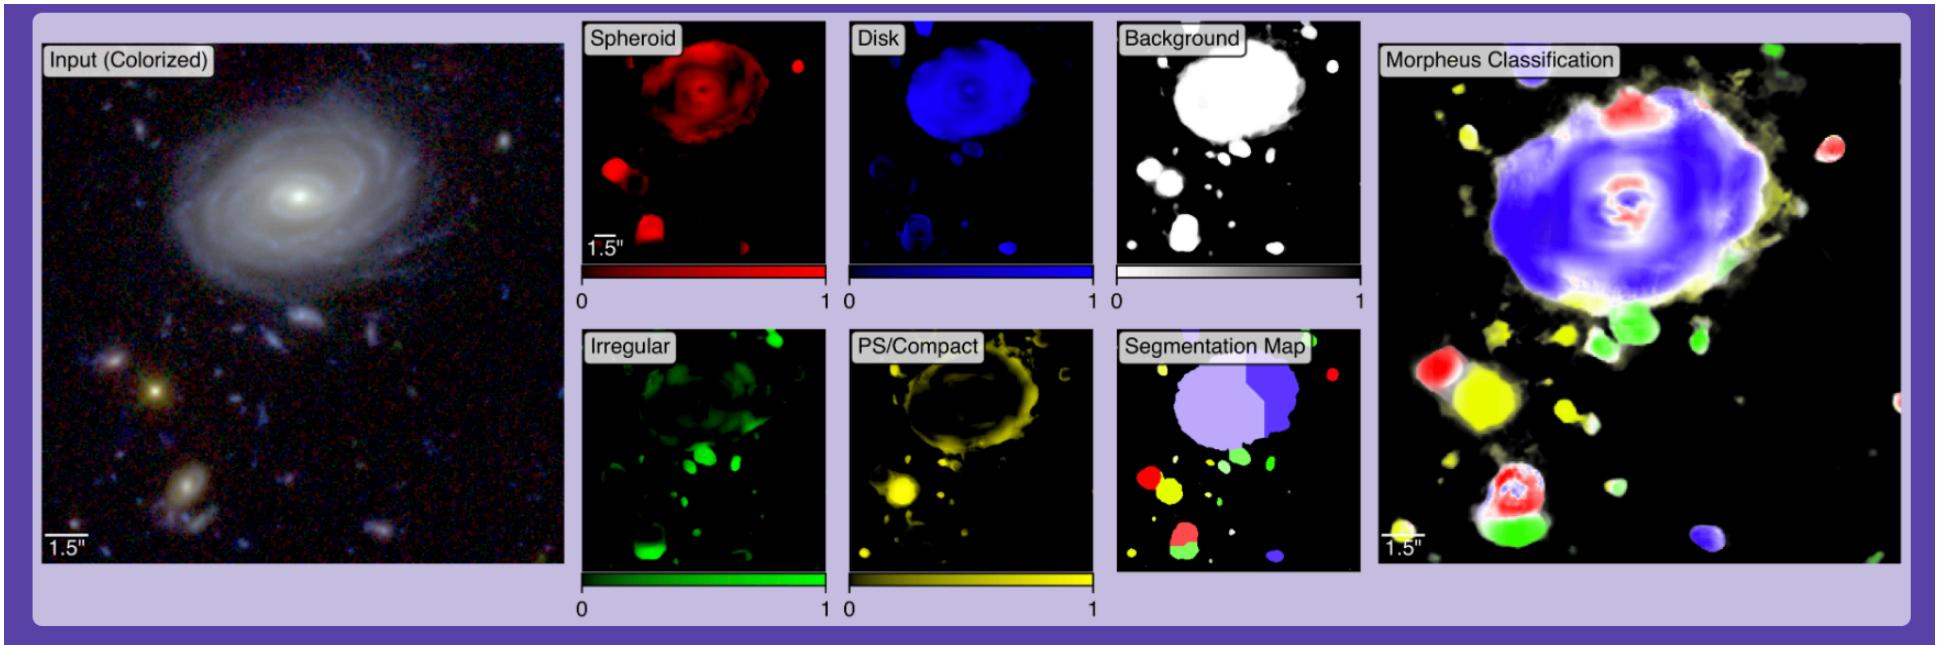
WE CALL THIS FULLY CONVOLUTIONAL NEURAL NETWORKS

ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



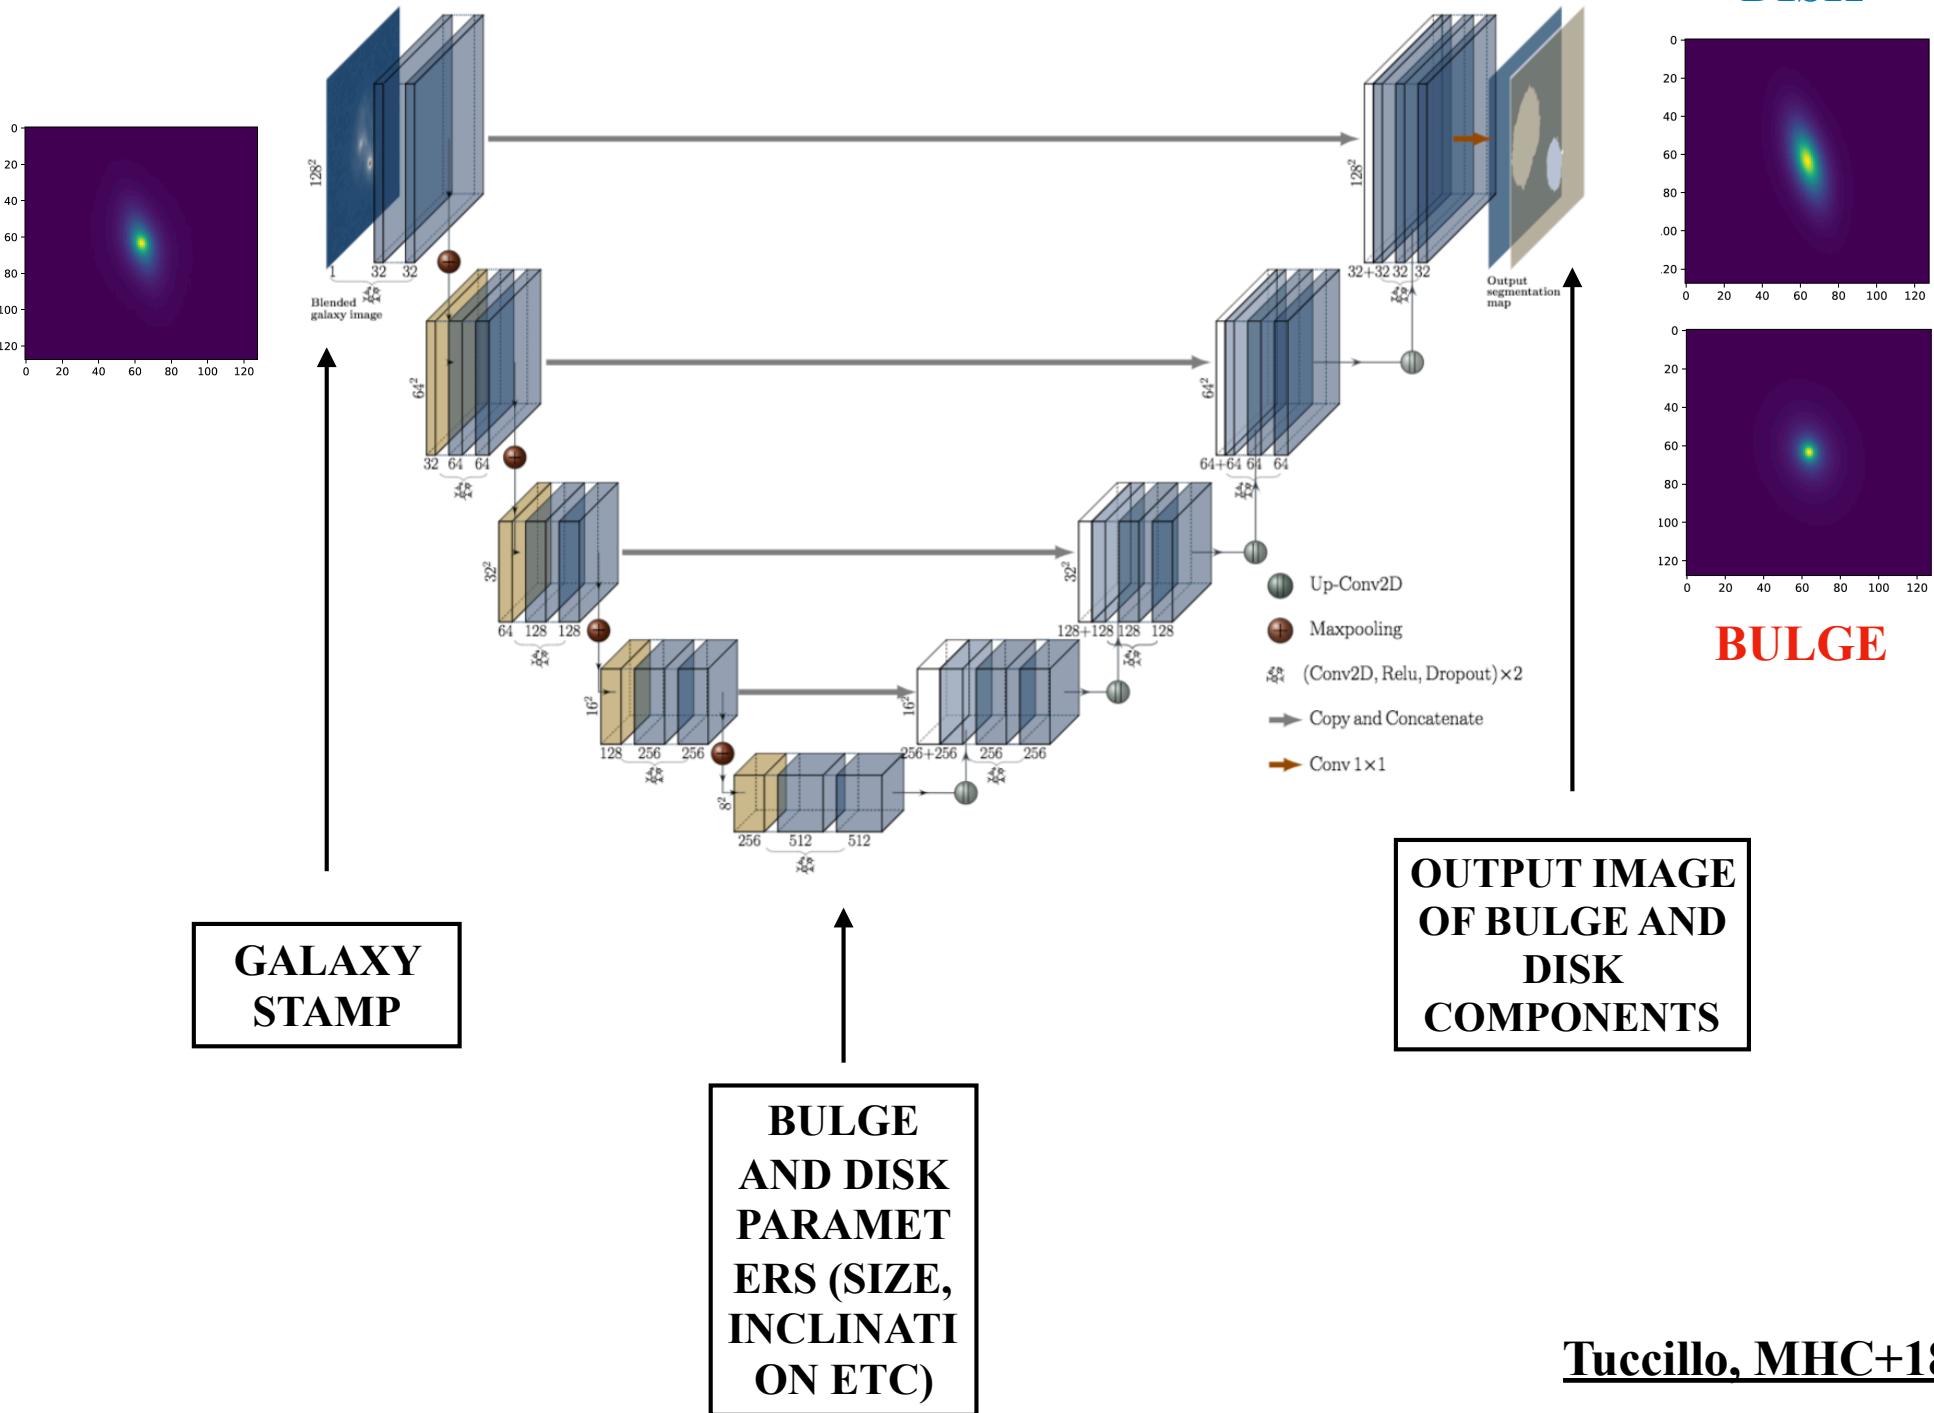
ENCODING-DECODING TO EXTRACT IMAGE FEATURES: THE U-NET





https://www.youtube.com/watch?v=hEL1h_dODkU#action=share

ALSO REGRESSION...

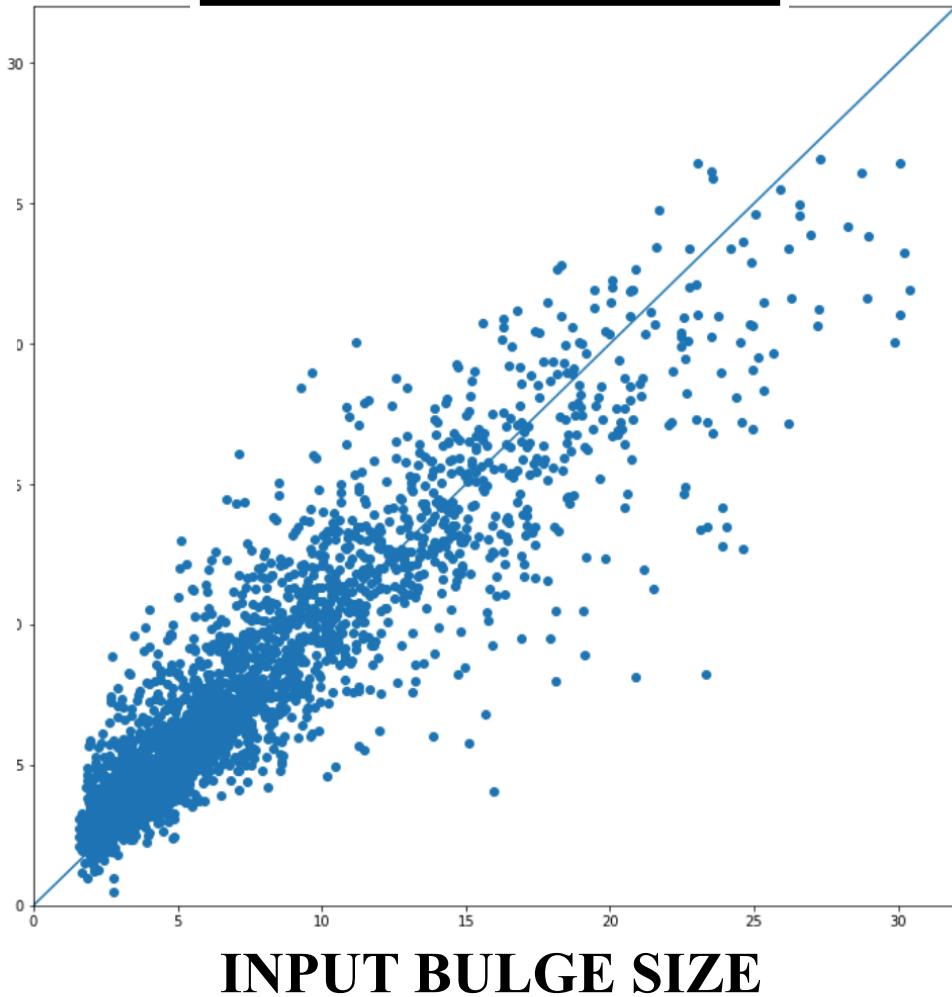


Tuccillo, MHC+18,19

PRELIMINARY!!

DEEP LEARNING

OUTPUT BULGE SIZE



GALFIT

OUTPUT BULGE SIZE

